

EBU7240

Computer Vision

Assignment Project Exam Help

- Introduction to Deep Learning -

<https://powcoder.com>

Add WeChat powcoder

Semester 1, 2021

Changjae Oh

Outline

- Machine learning basics++
- Introduction to deep learning
- Linear classifier

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

What is Machine Learning?

- Learning = Looking for a Function

Speech Recognition

$$f(\text{[audio waveform]}) = \text{"Hello"}$$

Handwritten Recognition

$$f(\text{[handwritten '2']}) = \text{"2"}$$

Weather forecast

$$f(\text{[sun icon] Thursday}) = \text{"[cloud icon] Saturday"}$$

Play video games

$$f(\text{[game state image]}) = \text{"move left"}$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Machine Learning: Basic Concept

- **Prediction task**
 - Regression: returns a specific value
 - Classification: returns a class label

Assignment Project Exam Help

Regression example


What is location when $x=10$?

<https://powcoder.com>


Add WeChat powcoder

Classification example

For a new image, what class does it belong to?



airplane
automobile
bird
cat
deer



Machine Learning: Basic Concept

- Training data

$\mathbb{X} = \{\mathbf{x}_1 = (2.0), \mathbf{x}_2 = (4.0), \mathbf{x}_3 = (6.0), \mathbf{x}_4 = (8.0)\}$

$\mathbb{Y} = \{y_1 = 3.0, y_2 = 4.0, y_3 = 5.0, y_4 = 6.0\}$

Assignment Project Exam Help


$y = wx + b$
<https://powcoder.com>

Optimal parameter $w=0.5$ $b=2.0$

Add WeChat powcoder

Goal

- Minimize errors for new samples (test set)
- Generalization
refers to high performance for test sets



Machine Learning: Basic Concept

- Multi-dimensional feature space

- d-dimensional data: $\mathbf{x}=(x_1, x_2, \dots, x_d)^T$ Note) d=784 in MNIST data

- Linear classifier for d-dimensional data

- 1-D linear classifier # of variables = $d + 1$

- Widely used in machine learning


$$y = \underline{w_1}x_1 + \underline{w_2}x_2 + \dots + \underline{w_d}x_d + \underline{b}$$

- 2-D linear classifier # of variables = $\frac{(d+1)(d+2)}{2}$


$$y = \underline{w_1}x_1^2 + \underline{w_2}x_2^2 + \dots + \underline{w_d}x_d^2 + \underline{w_{d+1}}x_1x_2 + \dots + w_{\frac{d(d+1)}{2}}x_{d-1}x_d + w_{\frac{d(d+1)}{2}+1}x_1 \dots + w_{\frac{d(d+1)}{2}+d}x_d + \underline{b}$$

Machine Learning: Basic Concept

- Feature space transformation
 - Map a linearly non-separable feature space into separable space



Original feature space



Transformed feature space,
which is linearly separable

Toy example

$\mathbf{x} = (x_1, x_2)^T \rightarrow \mathbf{x}' = \left(\frac{x_1}{2x_1x_2 + 0.5}, \frac{x_2}{2x_1x_2 + 0.5} \right)^T$

<https://powcoder.com>

Add WeChat powcoder

$$\mathbf{a} = (0,0)^T \rightarrow \mathbf{a}' = (0,0)^T$$

$$\mathbf{b} = (1,0)^T \rightarrow \mathbf{b}' = (2,0)^T$$

$$\mathbf{c} = (0,1)^T \rightarrow \mathbf{c}' = (0,2)^T$$

$$\mathbf{d} = (1,1)^T \rightarrow \mathbf{d}' = (0.4,0.4)^T$$

Machine Learning: Basic Concept

- Representation learning

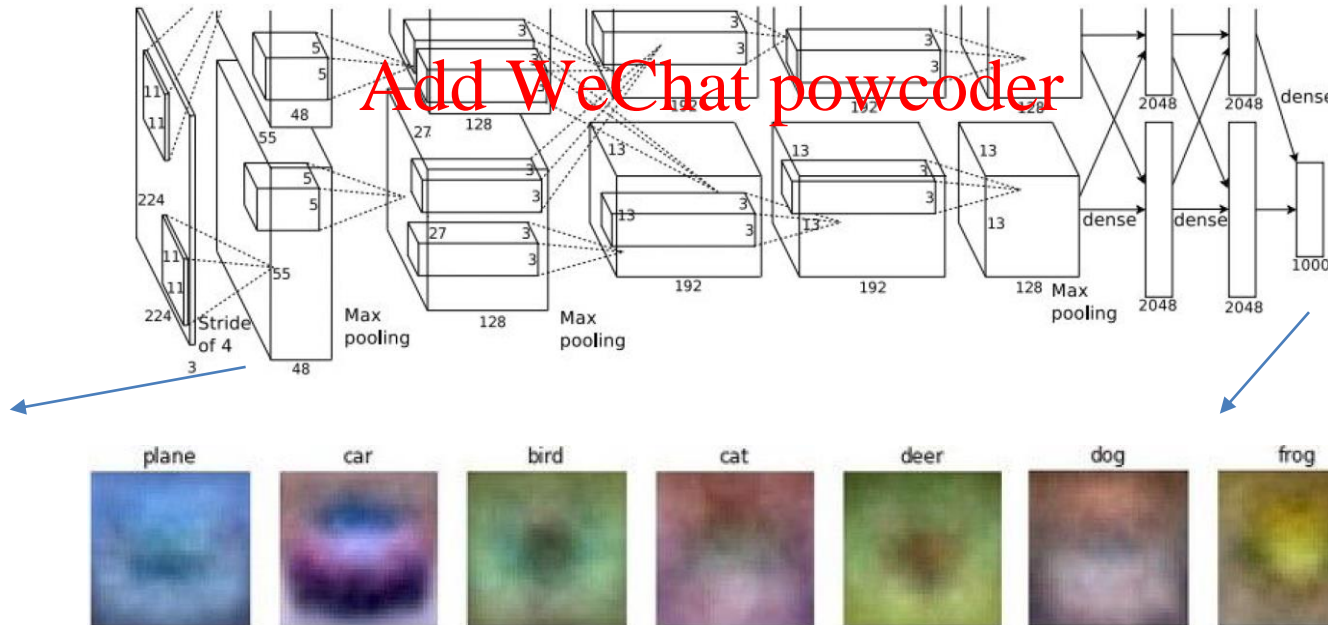

- Aims to find good feature space automatically
- Deep learning finds a hierarchical feature space by using neural networks with multiple hidden layers.

- The first hidden layer has low-level features (edge, corner points, etc.), and the right-hand side features advanced features (face, wheel, etc.)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Data for Machine Learning

- **The quality of the training data**

- To increase estimation accuracy, diverse and enough data should be collected for a given application.
- Ex) After learning from a database with a frontal face only, the recognition accuracy of side face will be degraded.


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

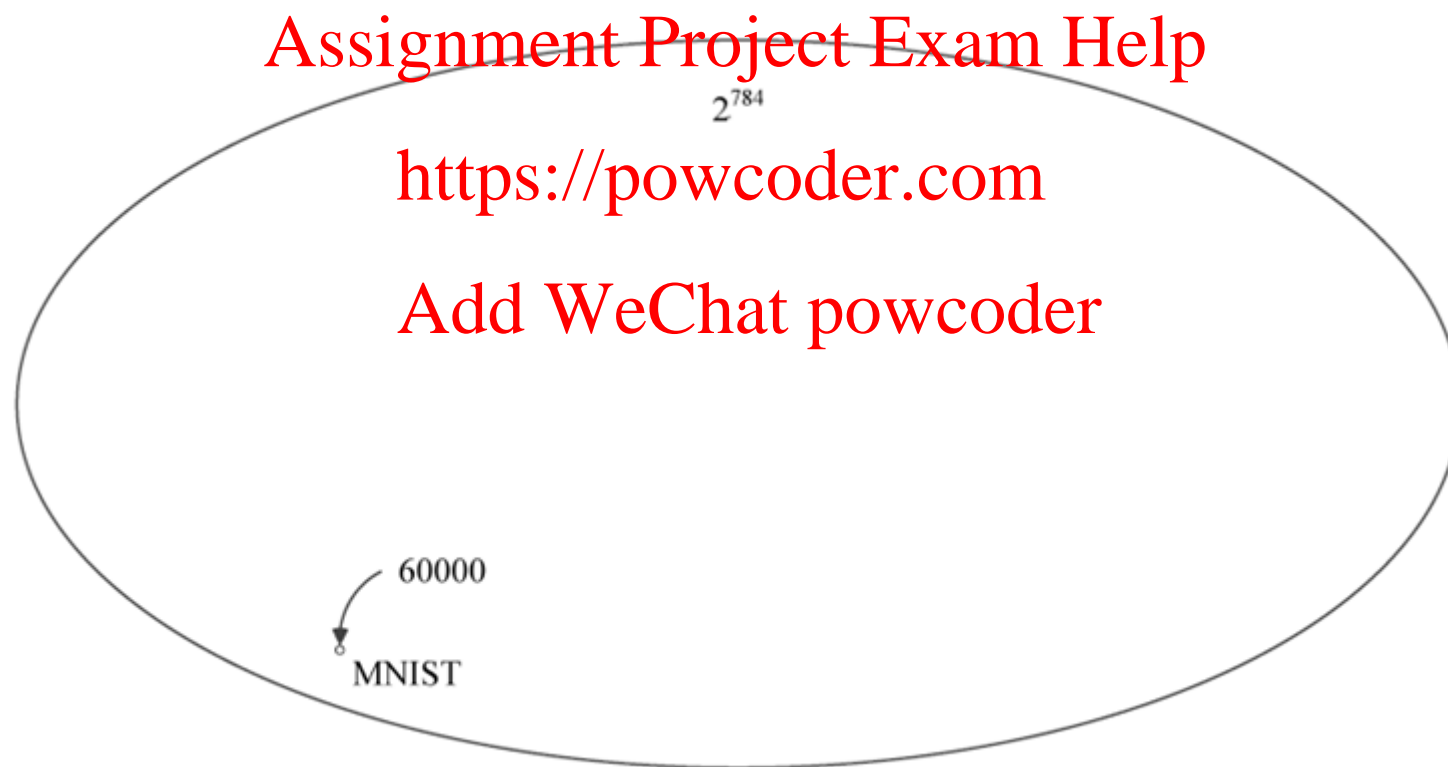
- **MNIST database**

- Handwritten numeric database
- Training data: 60,000
- Test data: 10,000



Data for Machine Learning

- Database size vs. training accuracy
 - Ex) MNIST: 28*28 binary image
 - The total number of possible samples is 2^{784} , but MNIST has 60,000 training images.



Data for Machine Learning

- How does a small database achieve high performance?
 - In a feature space, the actual data is generated in a very small subspace

 is unlikely to happen.

Assignment Project Exam Help

- Manifold assumption

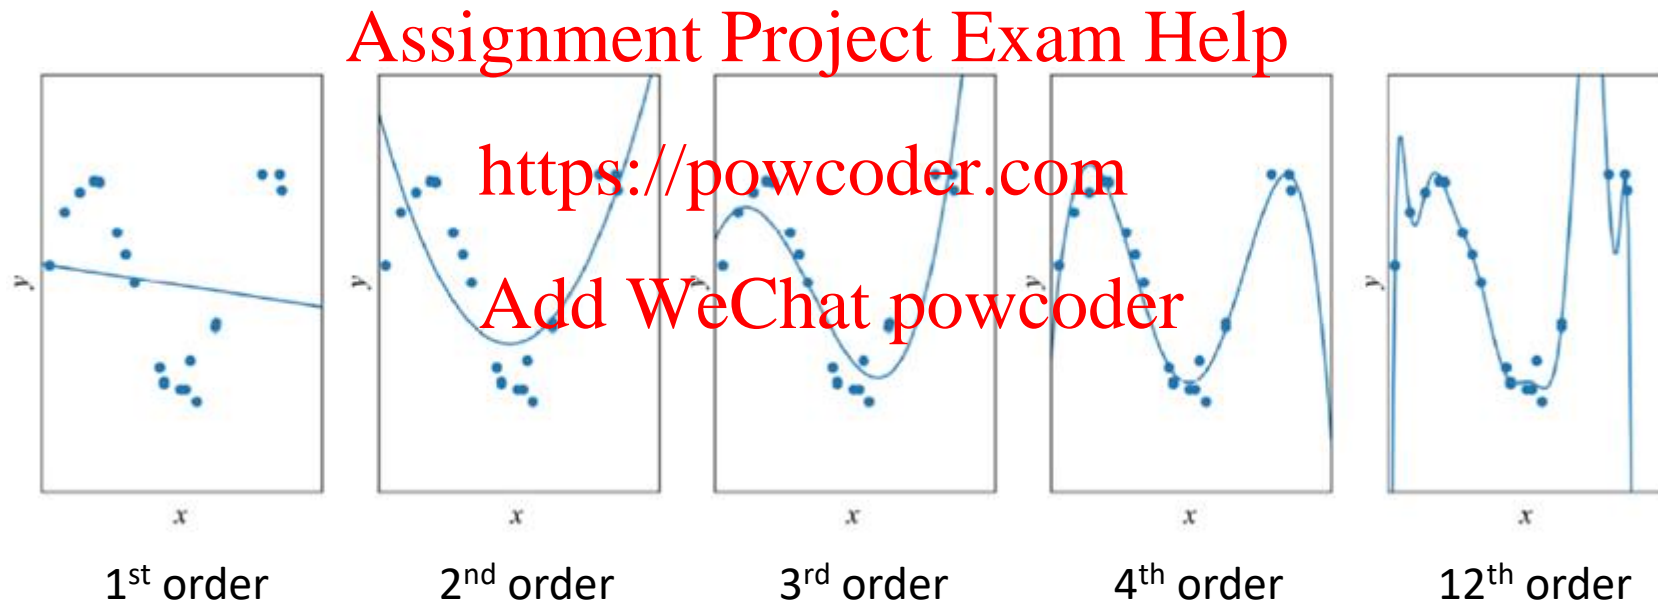
Smooth change according to certain rules like 

Add WeChat powcoder

Training Model: under-fitting vs. over-fitting

- Under-fitting

- Model capacity is too small to fit the data accordingly.
- Model with higher order can be used.



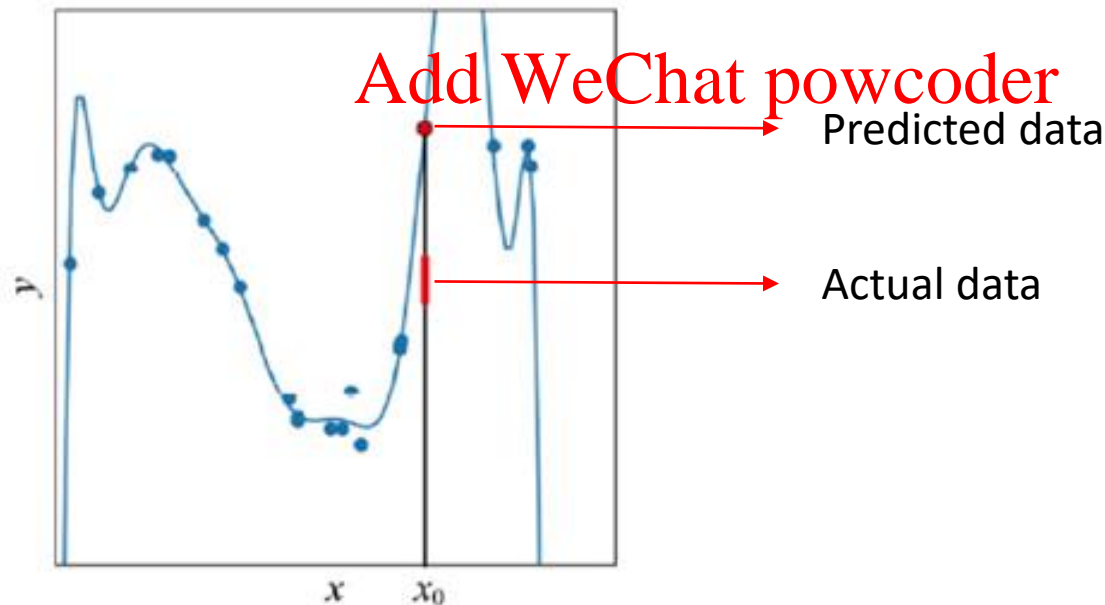
Training Model: under-fitting vs. over-fitting

- **Over-fitting**

- 12th order polynomial model approximates perfectly for the training set.
- But if you anticipate "new" data, there's a big problem.
- Since the model capacity is large, the training process also accepts data noise.
- The model with the appropriate capacity should be selected.

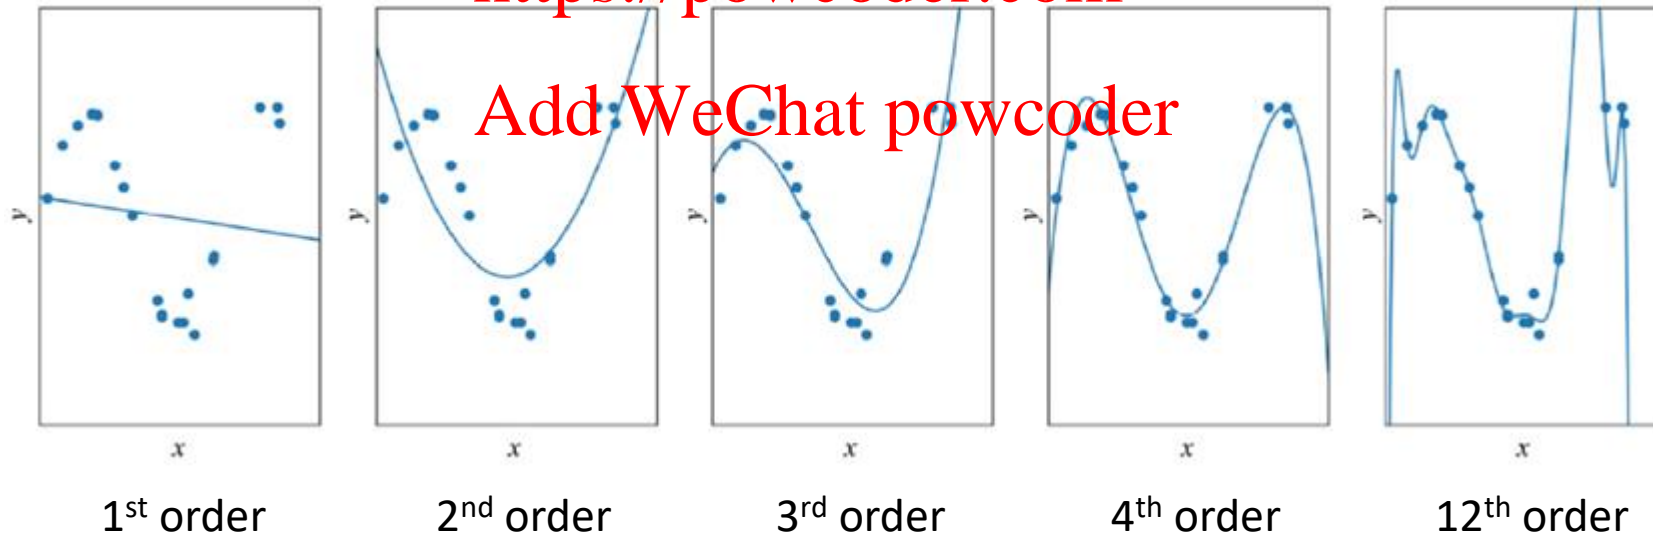
Assignment Project Exam Help

<https://powcoder.com>

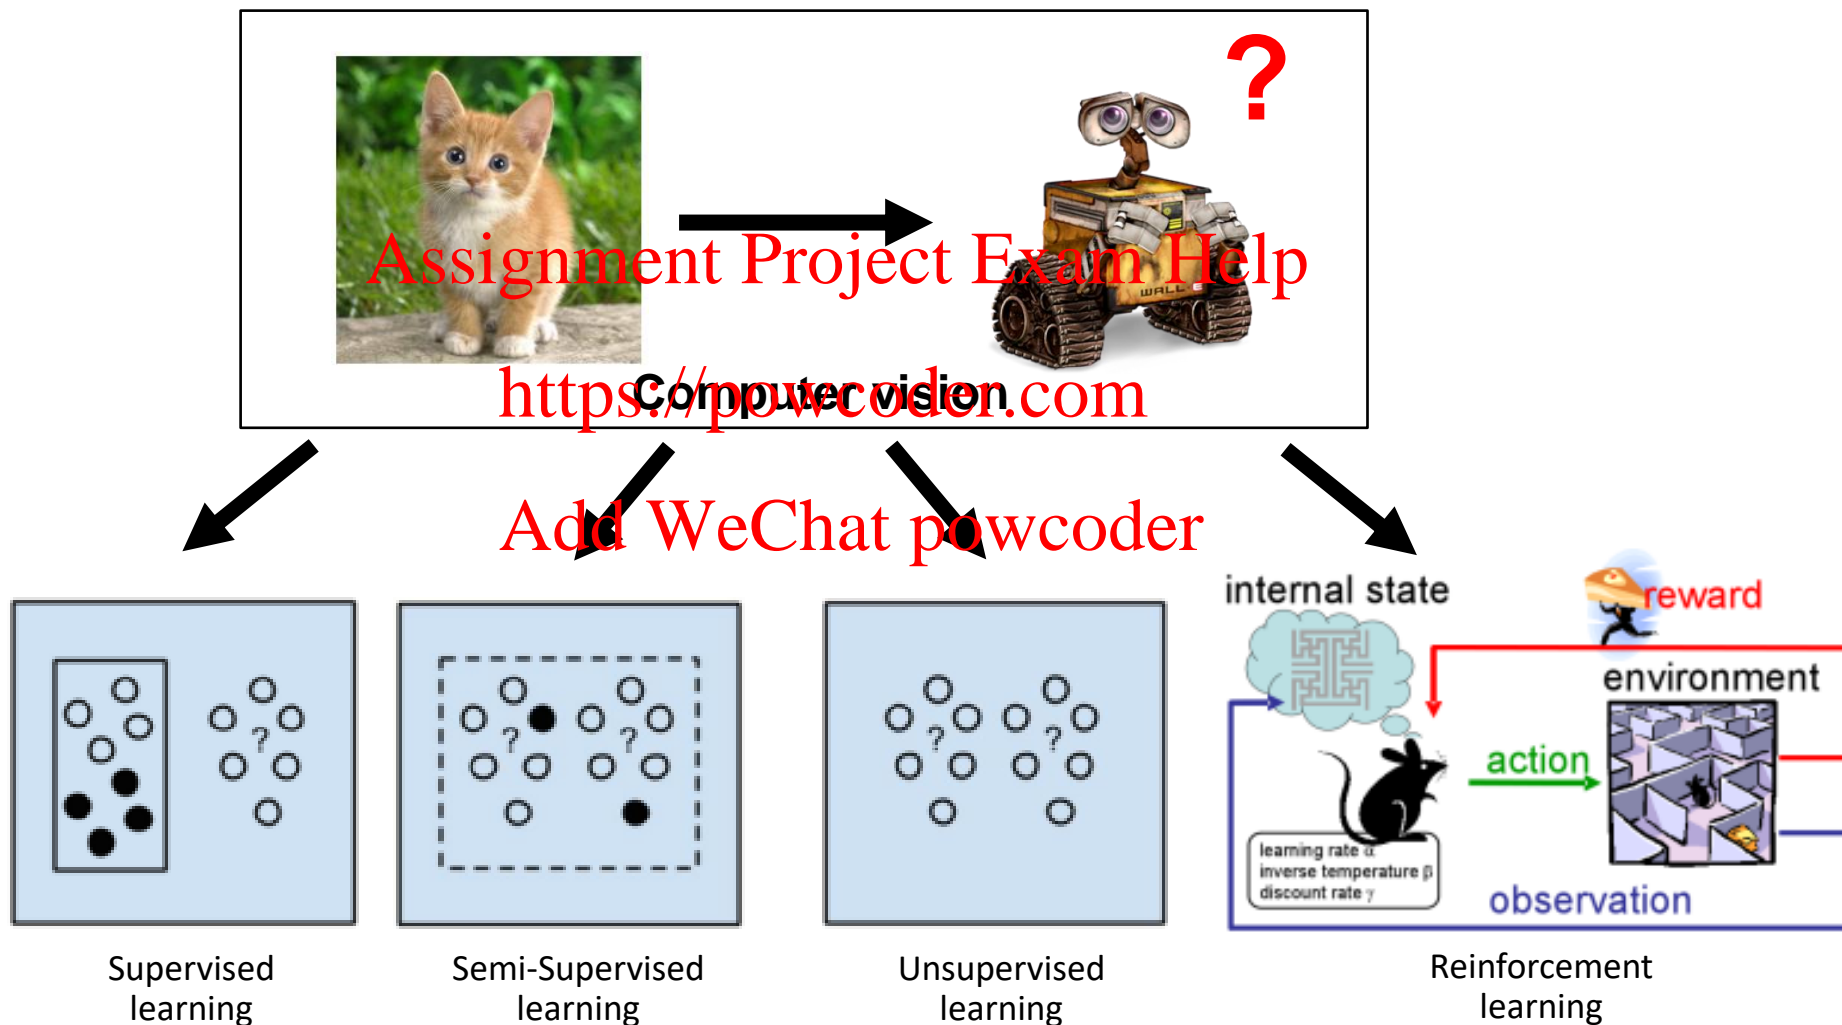


Training Model: under-fitting vs. over-fitting

- 1st and 2nd order model show poor performance for both the training and the test set.
- 12th order model shows high performance in training set, but low performance in test set. → low generalization ability
- 3rd and 4th order model are lower than the 12th order model for the training set, but the test set has high performance → higher generalization capability



Spectrum of supervision



Spectrum of supervision

- **Supervised learning**

- Both the feature vector \mathbb{X} and the output \mathbb{Y} are given.
- Regression and classification problem

- **Unsupervised learning**

- The feature vector \mathbb{X} is given, but the output \mathbb{Y} is not given.
- Ex) Clustering, density estimation, feature space conversion

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Spectrum of supervision

- **Reinforcement learning**

- The output is given, but it is different from supervised learning.
- Ex) Go
 - Once the game is over, you get a point (credit).
If you win, get 1, and -1 otherwise.
 - The credit should be distributed to each sample of the game.

- **Semi-supervised Learning**

- Some of data have both \mathbb{X} and \mathbb{Y} , but others have only \mathbb{X} .
- It is becoming important, since it is easy to collect \mathbb{X} , but \mathbb{Y} requires manual tasks.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Introduction to Deep Learning

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

From Wiki

- **Deep learning**

- is a branch of machine learning based on a set of algorithms that attempt to model high-level abstractions in data by using multiple processing layers, with complex structures or otherwise, composed of multiple non-linear transformations.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Deep learning success

- Image classification
- Machine translation
- Speech recognition
- Speech synthesis
- Game playing
- .. and many, many more

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



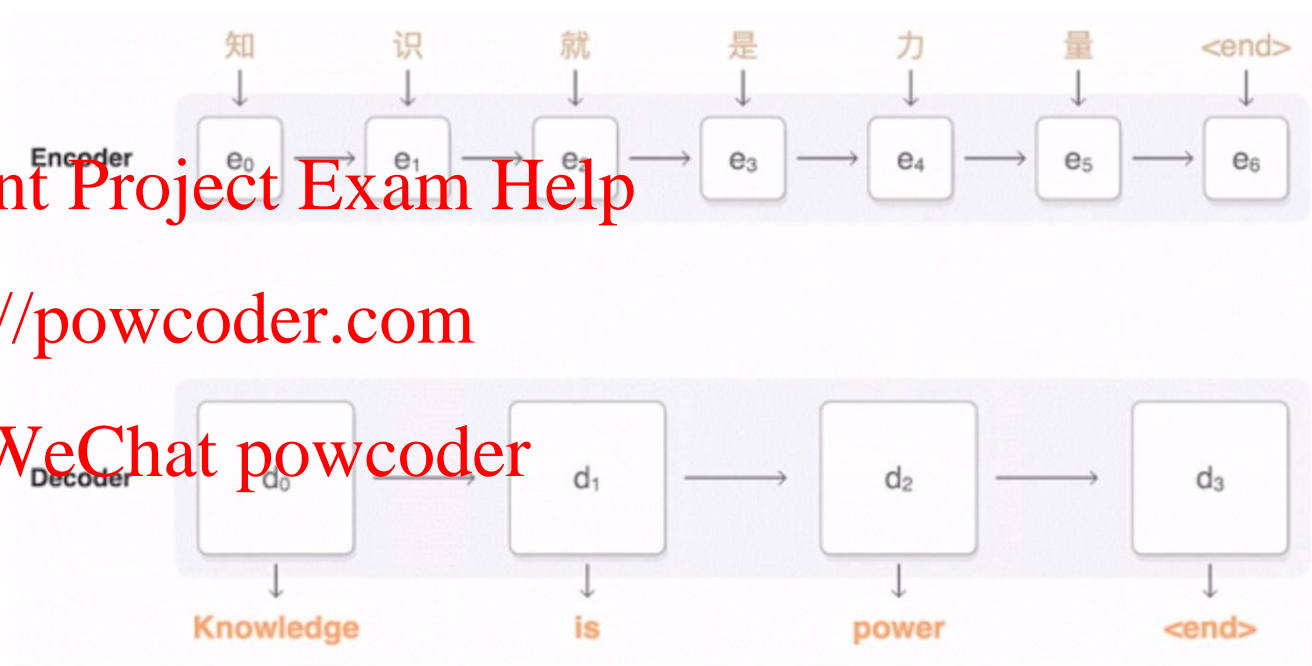
Deep learning success

- Image classification
- **Machine translation**
- Speech recognition
- Speech synthesis
- Game playing
- .. and many, many more

Assignment Project Exam Help


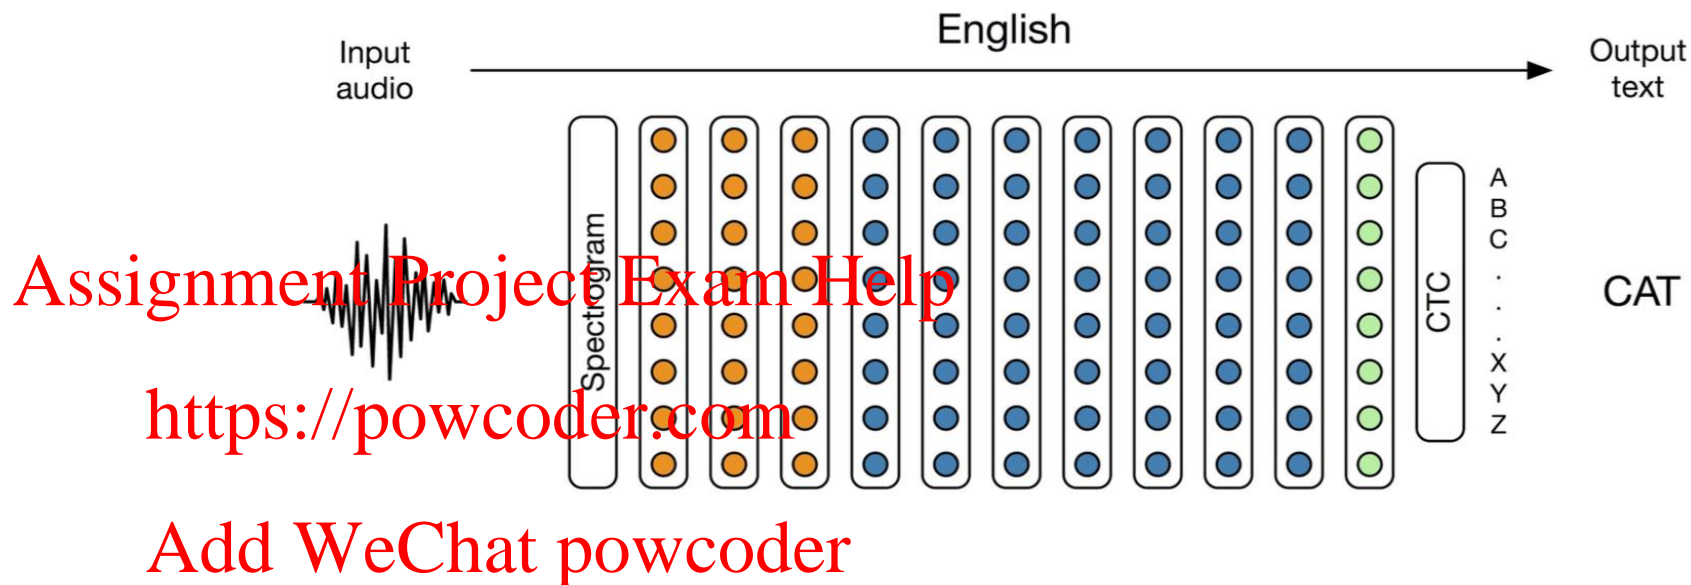
<https://powcoder.com>

Add WeChat powcoder



Deep learning success

- Image classification
- Machine translation
- **Speech recognition**
- Speech synthesis
- Game playing
- .. and many, many more




Deep learning success

- Image classification
 - Machine translation
 - Speech recognition
 - **Game playing**
-
- .. and many, many more

Assignment Project Exam Help

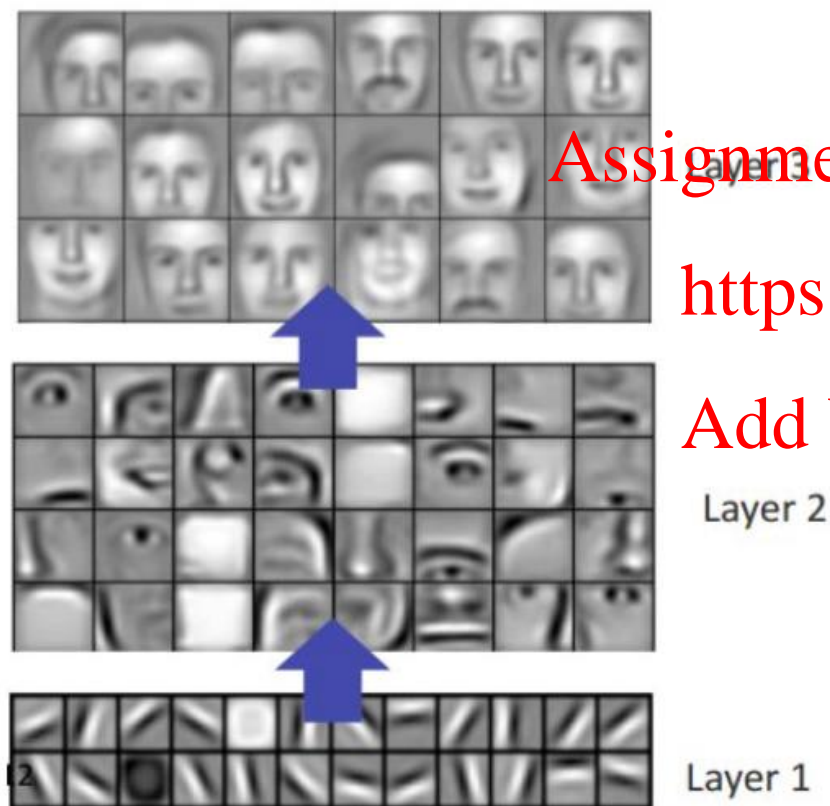
<https://powcoder.com>

Add WeChat powcoder



Why deep learning?

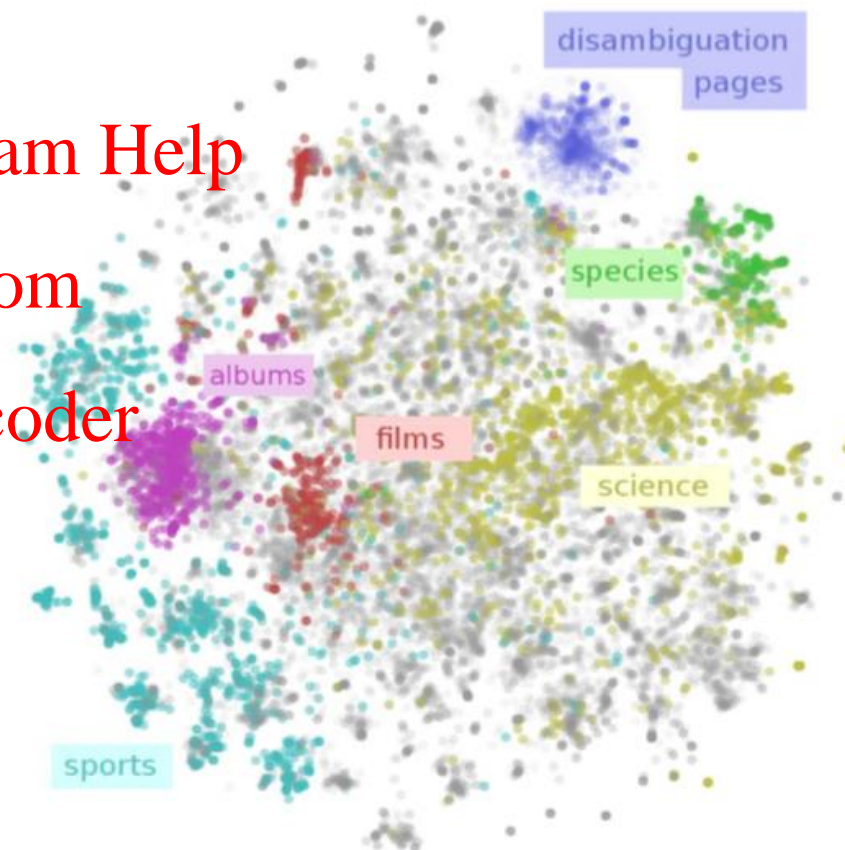
- Hand-crafted features vs. Learned features



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Why now?

- Large datasets
- GPU hardware advances + Price decreases
- Improved techniques (algorithm)

Assignment Project Exam Help

<https://powcoder.com>

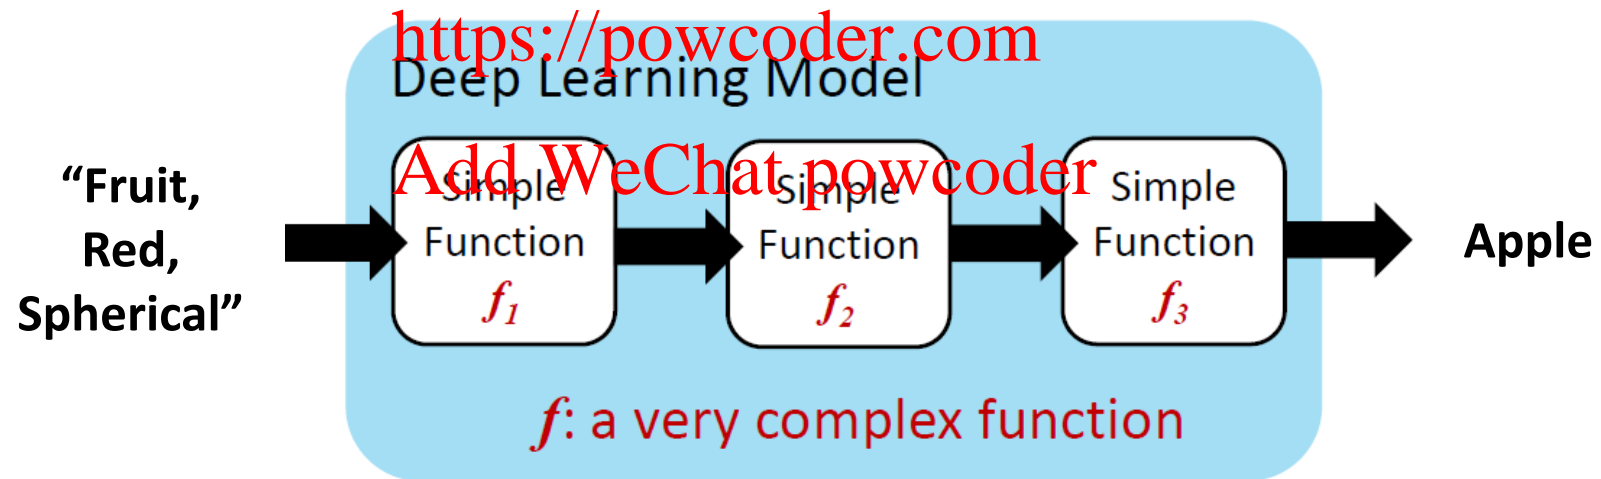
Add WeChat powcoder



What is Deep Learning?

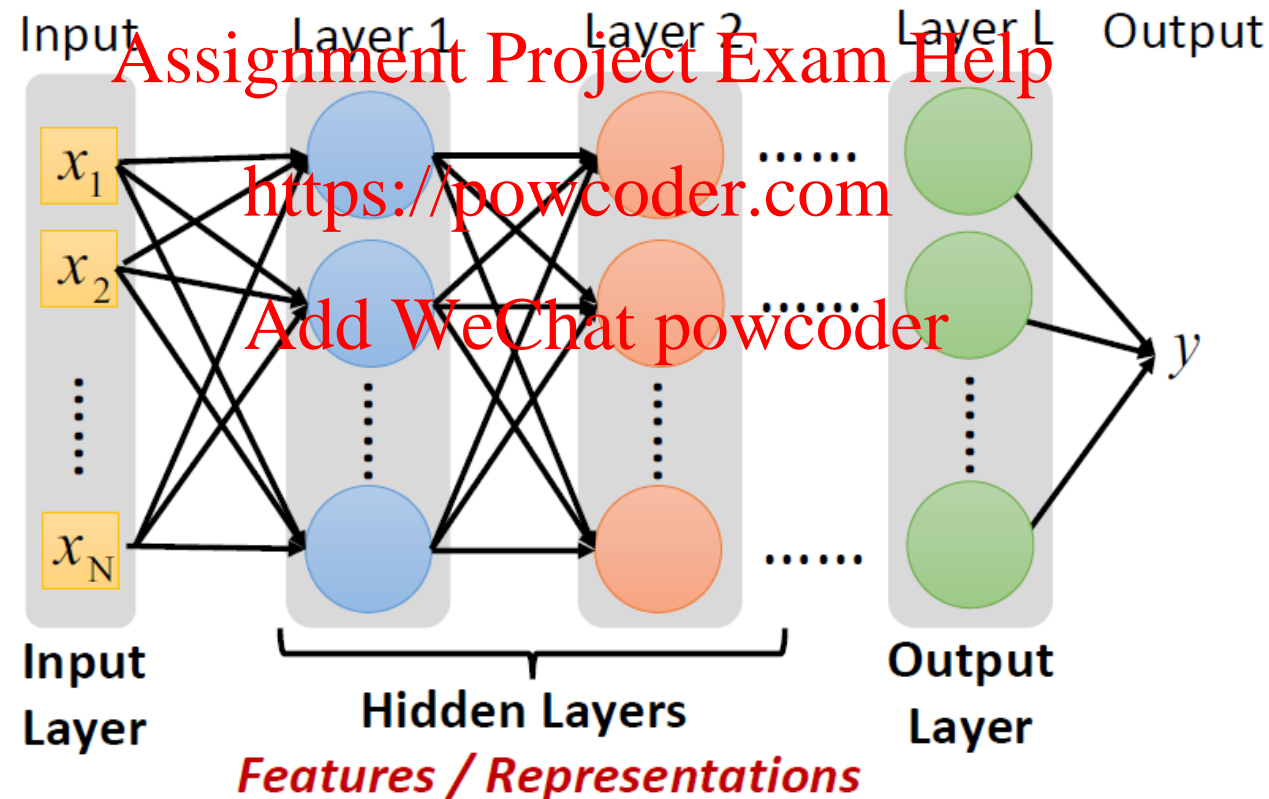
- **Stacked Functions Learned by Machine**
 - **End-to-end training**: what each function should do is learned automatically
 - Deep learning usually refers to **neural network** based model

Assignment Project Exam Help



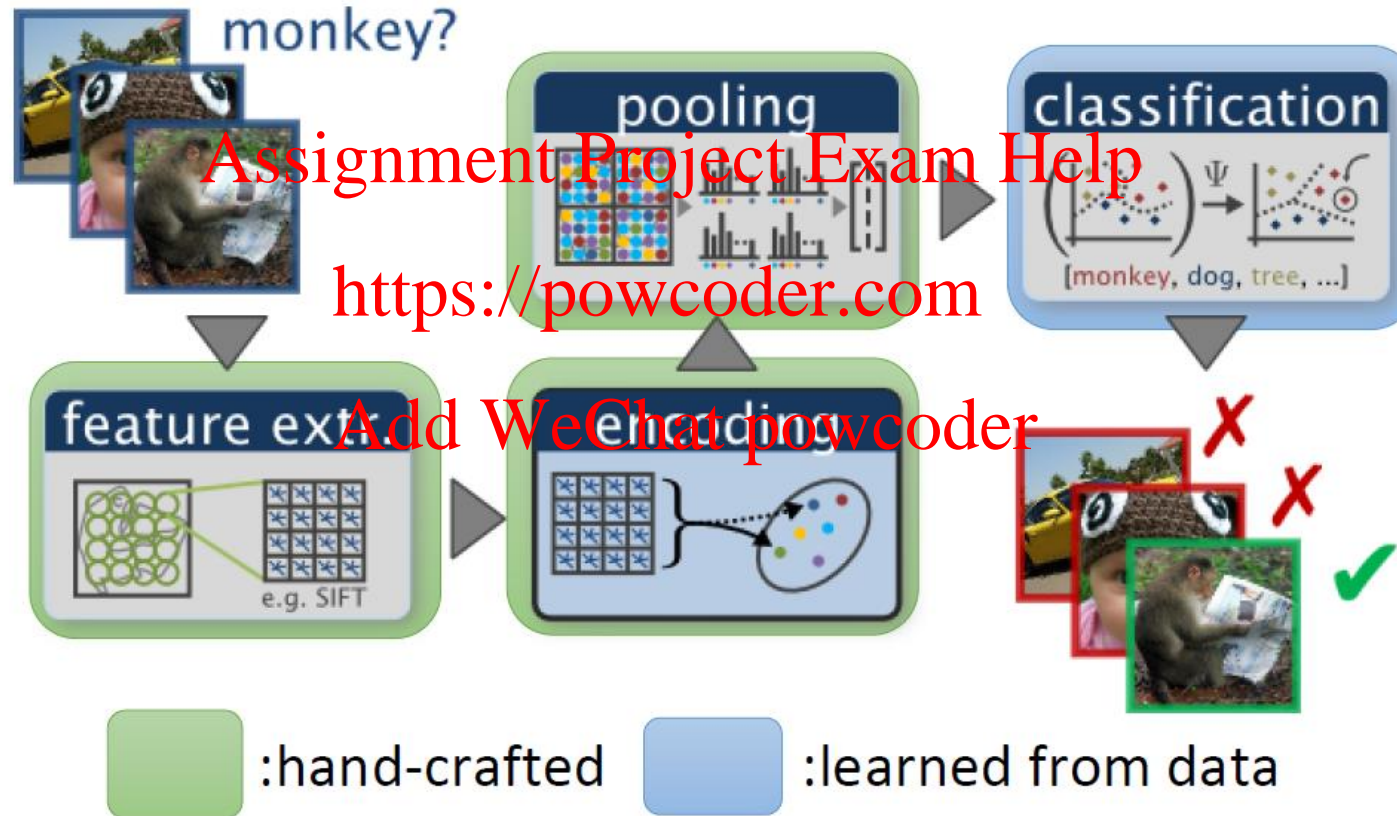
What is Deep Learning?

- **Stacked Functions Learned by Machine**
 - **Representation Learning:** learning features/representations
 - **Deep Learning:** learning (multi-level) features and an output



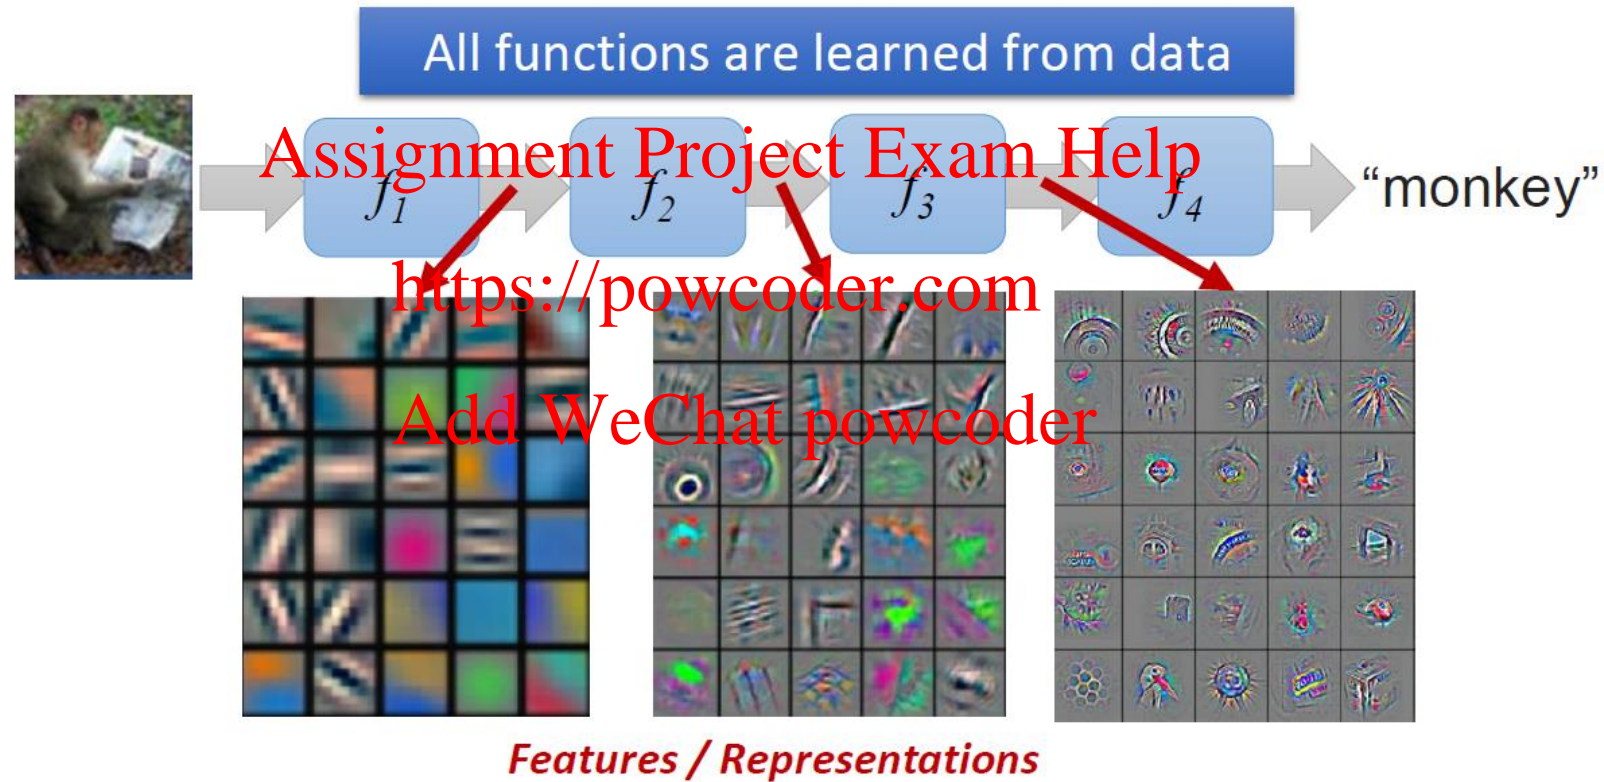
Machine Learning vs. Deep Learning

- Deep vs Shallow: Image Recognition
 - Shallow model using machine learning



Machine Learning vs. Deep Learning

- Deep vs Shallow: Image Recognition
 - Deep model using deep learning



Machine Learning vs. Deep Learning

- Machine Learning vs. Deep Learning

Machine Learning = Feature descriptor + Classifier

<https://powcoder.com>

Add WeChat powcoder

Feature: hand-crafted domain-specific knowledge
Describing your data with features that computer can understand
Ex) SIFT, Bag-of-Words (BoW), Histogram of Oriented Gradient (HOG)

Optimizing the classifier weights on features
Ex) Nearest Neighbor (NN), Support Vector Machine (SVM), Random Forest (RF)

Machine Learning vs. Deep Learning

- Machine Learning vs. Deep Learning


Deep Learning = Feature descriptor + Classifier

<https://powcoder.com>

Add WeChat powcoder

Feature: Representation learned by machine
Automatically learned internal knowledge

Optimizing the classifier weights on
features

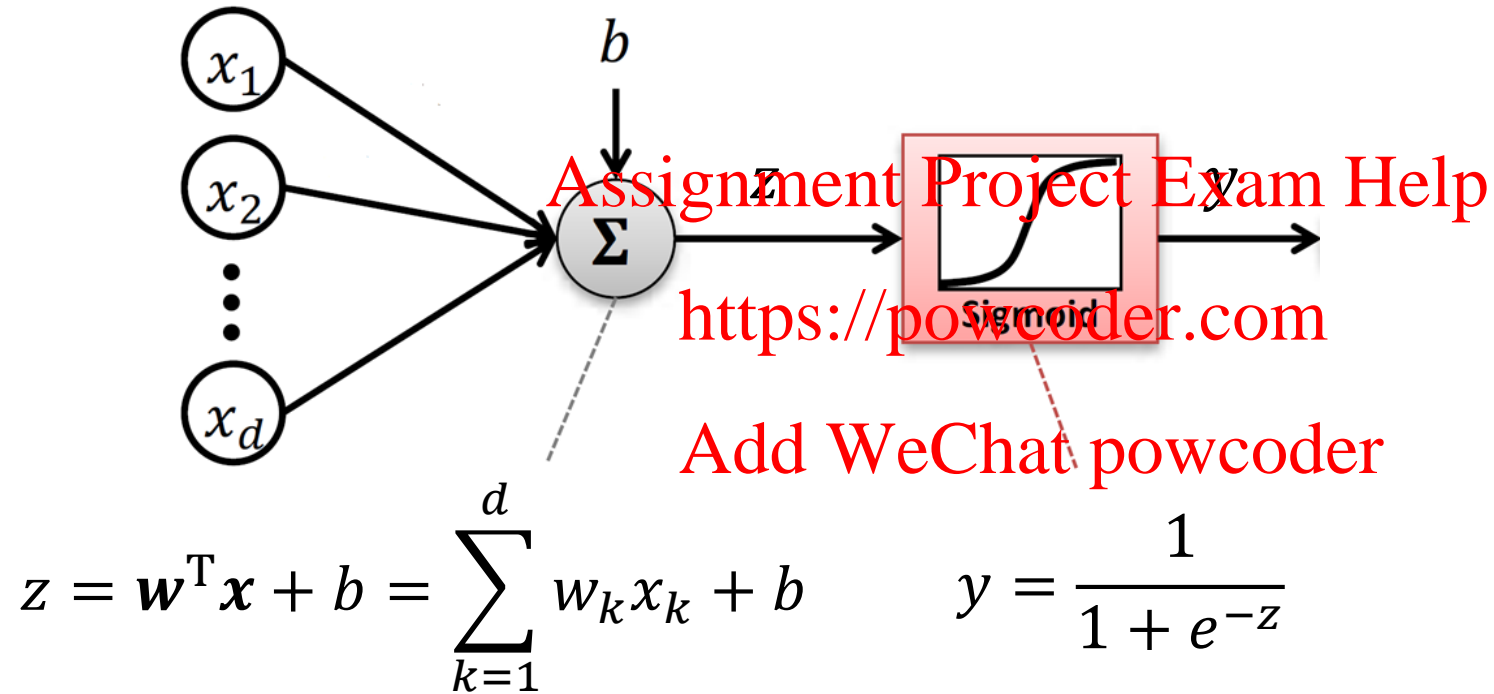


Neural network based model

A series of linear classifiers and non-linear activations + Loss function


Deep Learning

- A single neuron




Deep Learning

- A single layer with multiple neurons



$$z_1 = \mathbf{w}_1^T \mathbf{x} + b_1 = \sum_{k=1}^d w_{1k} x_k + b_1 \quad y_1 = \frac{1}{1 + e^{-z_1}}$$


Add WeChat powcoder



$$z_M = \mathbf{w}_M^T \mathbf{x} + b_M = \sum_{k=1}^d w_{Mk} x_k + b_M \quad y_M = \frac{1}{1 + e^{-z_n}}$$

Deep Learning

- **Deep Neural Network**
 - Cascading the neurons to form a neural network



Each layer consists of the linear classifier
and activation function

Linear classifier

Neural Network

Linear
classifiers



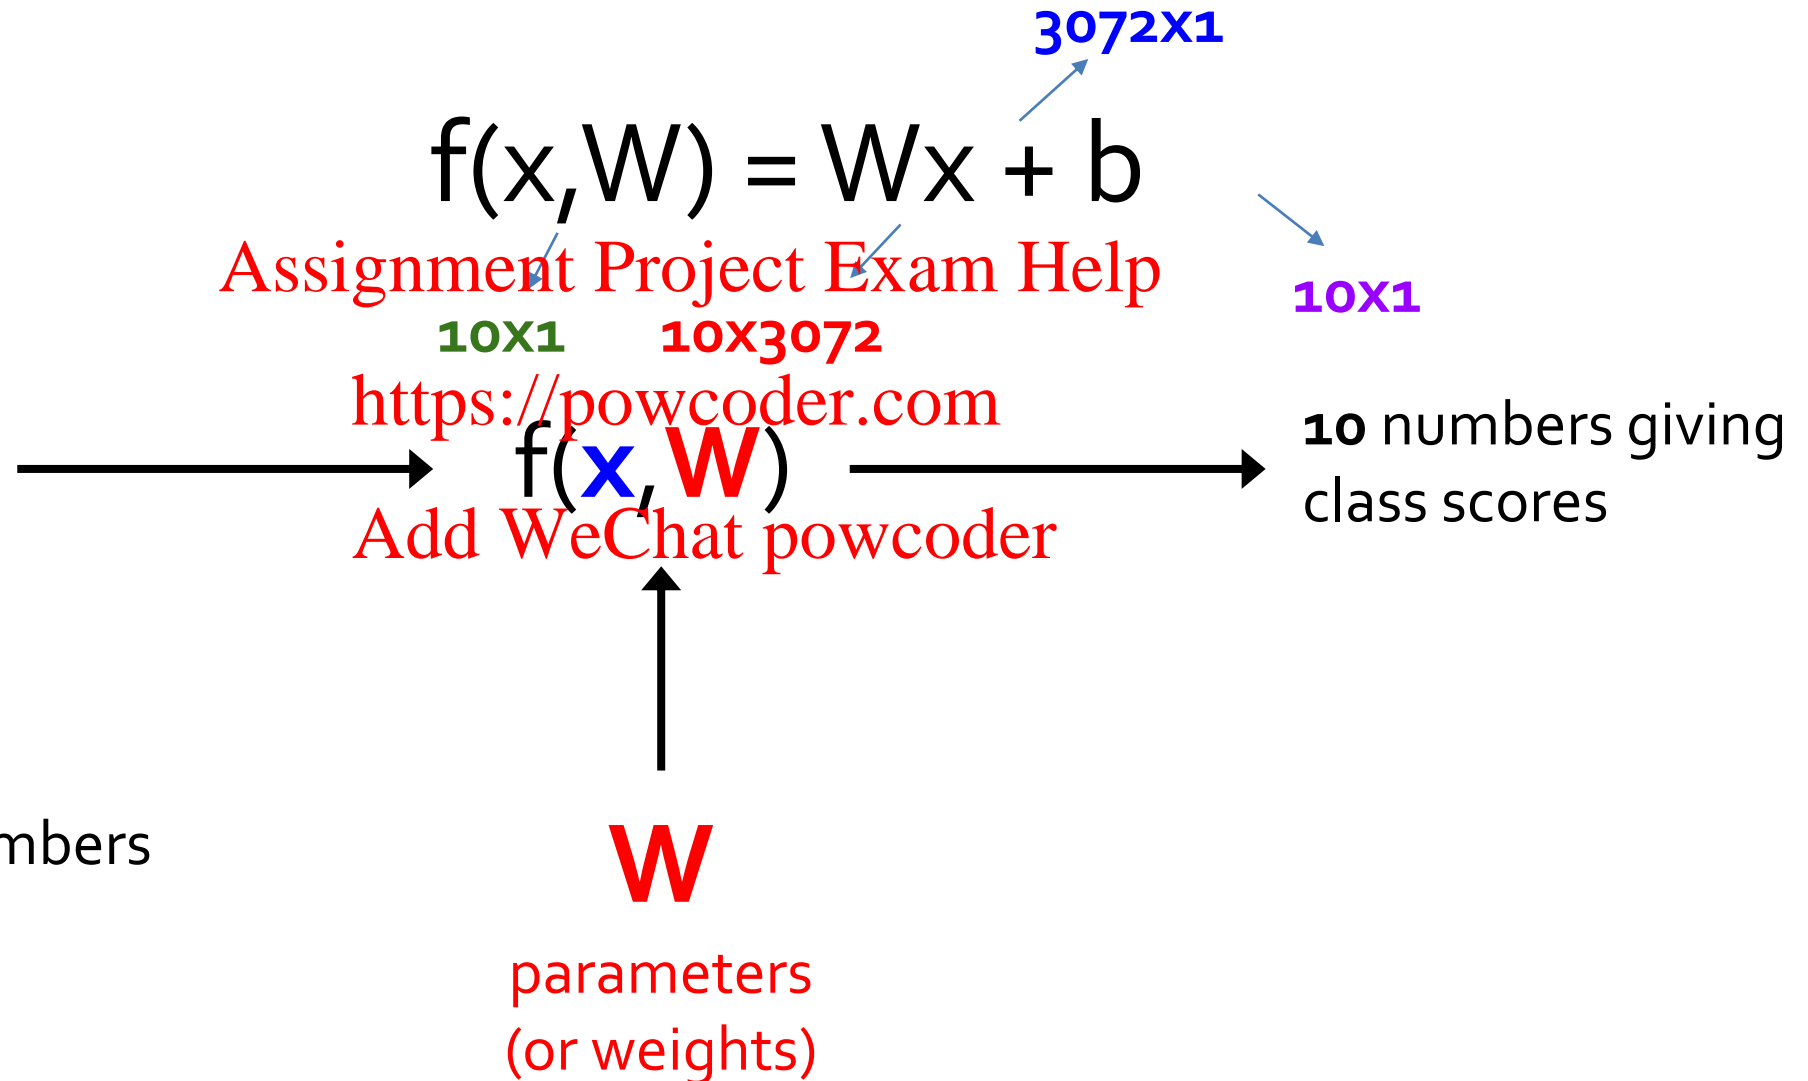
Parametric Approach

- (Review) Unit 3 ML basics and classification

Image



Array of $32 \times 32 \times 3$ numbers
(3072 numbers total)



Parametric Approach

- (Review) Unit 3 recognition

Example with an image with 4 pixels, and 3 classes (cat/dog/ship)

Stretch pixels into column



Input image

0.2	-0.5	0.1	2.0
1.5	1.3	2.1	0.0
0	0.25	0.2	-0.3

W

56
231
24
2

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

+

=

1.1
3.2
-1.2

-96.8
437.9
61.95

Cat score

Dog score

Ship score

What do we need now?

- Functions to measuring the error between the output of a classifier and the given target value.
 - Let's talk about designing error (a.k.a. loss) functions!

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

EBU7240

Computer Vision

Assignment Project Exam Help

<https://powcoder.com> -

Add WeChat powcoder

Semester 1, 2021

Changjae Oh

Loss function

- **Loss function**
 - quantifies our unhappiness with the scores across the training data.
- **Type of loss function**
 - Hinge loss
 - Cross-entropy loss
 - Log likelihood loss
 - Regression loss

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Loss Function: Hinge Loss

- Binary hinge loss (=binary SVM loss)

$$L_i = \max(0, 1 - y_i \cdot s) \quad s = \mathbf{w}^T \mathbf{x}_i + b$$

$y_i = \pm 1$ for positive/negative samples

Assignment Project Exam Help

- Hinge loss (=multiclass SVM loss)

– C : The number of class (> 2)


<https://powcoder.com>

Add WeChat powcoder

$$L_i = \sum_{j=1, j \neq y_i}^C \max(0, s_j - s_{y_i} + 1)$$

\mathbf{x}_i : input data (e.g. image)

y_i : class label (integer, $1 \leq y_i \leq C$)



$$\mathbf{s} = \mathbf{W}\mathbf{x}_i + \mathbf{b}$$

$$\mathbf{W} = \begin{pmatrix} \mathbf{w}_1^T \\ \mathbf{w}_2^T \\ \vdots \\ \mathbf{w}_C^T \end{pmatrix} \quad \mathbf{s} = \begin{pmatrix} s_1 \\ s_2 \\ \vdots \\ s_C \end{pmatrix}$$

Loss Function: Hinge Loss

Suppose: 3 training examples, 3 classes.

With some \mathbf{W} the scores $f(\mathbf{x}, \mathbf{W}) = \mathbf{W}\mathbf{x} + \mathbf{b}$ are



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

cat	3.2	1.3	2.2
car	5.1	4.9	2.5
frog	-1.7	2.0	-3.1

Given a dataset of examples

$$\{(\mathbf{x}_i, y_i)\}_{i=1}^N$$

\mathbf{x}_i : image

y_i : class label (integer)

Loss over the dataset is a sum of loss over examples:

$$L = \frac{1}{N} \sum_{i=1}^N L_i(f(\mathbf{x}_i, \mathbf{W}), y_i)$$

Loss Function: Hinge Loss

Suppose: 3 training examples, 3 classes.

With some W the scores $f(\mathbf{x}, \mathbf{W}) = \mathbf{W}\mathbf{x} + \mathbf{b}$ are

Multiclass SVM loss (=hinge loss)

$$L_i = \sum_{j \neq y_i} \begin{cases} 0 & \text{if } s_{y_i} \geq s_j + 1 \\ s_j - s_{y_i} + 1 & \text{otherwise} \end{cases}$$
$$= \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

cat	3.2	1.3	2.2
car	5.1	4.9	2.5
frog	-1.7	2.0	-3.1

where score vector $\mathbf{s} = f(\mathbf{x}_i, \mathbf{W})$

Loss Function: Hinge Loss

Suppose: 3 training examples, 3 classes.

With some W the scores $f(\mathbf{x}, \mathbf{W}) = \mathbf{W}\mathbf{x} + \mathbf{b}$ are

Multiclass SVM loss (=hinge loss)

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

cat	3.2	1.3	2.2
car	5.1	4.9	2.5
frog	-1.7	2.0	-3.1
Loss	2.9	0	12.9

$$\begin{aligned} &= \max(0, 5.1 - 3.2 + 1) + \max(0, -1.7 - 3.2 + 1) \\ &= \max(0, 2.9) + \max(0, -3.9) \\ &= 2.9 + 0 = 2.9 \end{aligned}$$

$$\begin{aligned} &= \max(0, 1.3 - 4.9 + 1) + \max(0, 2.0 - 4.9 + 1) \\ &= \max(0, -2.6) + \max(0, -1.9) \\ &= 0 + 0 = 0 \end{aligned}$$

$$\begin{aligned} &= \max(0, 2.2 - (-3.1) + 1) + \max(0, 2.5 - (-3.1) + 1) \\ &= \max(0, 6.3) + \max(0, 6.6) \\ &= 6.3 + 6.6 \end{aligned}$$

Loss over full dataset is average

$$L = (2.9 + 0 + 12.9)/3 = 5.27$$

Loss Function: Log Likelihood Loss

- Log likelihood loss

$$L_i = -\log p_j \text{ where } j \text{ satisfies } z_{ij} = 1$$

$$\mathbf{p} = \begin{pmatrix} p_1 \\ p_2 \\ \vdots \\ p_C \end{pmatrix} \quad \begin{array}{l} \text{probability for } i^{th} \text{ image} \\ \text{(It is assumed to be } \textcolor{red}{normalized}, \text{ i.e. } |\mathbf{p}| = 1.) \end{array}$$

Assignment Project Exam Help

z_i : class label for i^{th} image
($C \times 1$ vector, $z_{ij} = 1$ when $j = y_i$ and 0 otherwise)

<https://powcoder.com>
 y_i : class label (integer, $1 \leq y_i \leq C$)

Example

Suppose i^{th} image belongs to class 2 and $C = 10$.

$$\mathbf{z}_i = \begin{pmatrix} 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{pmatrix}$$

$$\mathbf{p} = \begin{pmatrix} 0.1 \\ 0.7 \\ 0 \\ \vdots \\ 0.2 \end{pmatrix}$$



$$L_i = -\log 0.7$$

Add WeChat powcoder

Loss Function: Cross-entropy Loss

- Cross-entropy loss

$$L_i = - \sum_{j=1}^C (z_{ij} \log p_j + (1 - z_{ij}) \log(1 - p_j))$$

$$\mathbf{p} = \begin{pmatrix} p_1 \\ p_2 \\ \vdots \\ p_C \end{pmatrix} \quad \begin{array}{l} \text{probability for } i^{th} \text{ image} \\ \text{(It is assumed to be } \textcolor{red}{normalized}, \text{ i.e. } |\mathbf{p}| = 1.) \end{array}$$

Assignment Project Exam Help
 \mathbf{z}_i : class label for i^{th} image
($C \times 1$ vector, $z_{ij} = 1$ when $j = y_i$ and 0 otherwise)

<https://powcoder.com>
 y_i : class label (integer, $1 \leq y_i \leq C$)

Example

Suppose i^{th} image belongs to class 2 and $C = 10$.

$$\mathbf{z}_i = \begin{pmatrix} 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{pmatrix} \quad \mathbf{p} = \begin{pmatrix} 0.1 \\ 0.7 \\ 0 \\ \vdots \\ 0.2 \end{pmatrix} \quad \Rightarrow \quad L_i = -\log(1 - 0.1) - \log 0.7 - \log(1 - 0.2)$$

Add WeChat powcoder

Softmax Activation Function

- Softmax activation function

scores = unnormalized log probabilities of the classes.

→ Probability can be computed using scores as below.

Probability of class label being k for an image x_i

Assignment Project Exam Help

<https://powcoder.com>

$$p_k = \frac{e^{s_k}}{\sum_{j=1}^C e^{s_j}}$$

Softmax activation function



unnormalized probabilities

Add WeChat powcoder

x_i : image

y_i : class label (integer, $1 \leq y_i \leq C$)

cat
car
frog

cat	3.2
car	5.1
frog	-1.7

exp

cat	24.5
car	164.0
frog	0.18

normalize

cat	0.13
car	0.87
frog	0.00

unnormalized log probabilities

probabilities

$$s = Wx_i + b$$

$$W = \begin{pmatrix} w_1^T \\ w_2^T \\ \vdots \\ w_C^T \end{pmatrix}$$

Softmax + Log Likelihood Loss



$$L_i = -\log \left(\frac{e^{s_{y_i}}}{\sum_{j=1}^C e^{s_j}} \right)$$

Assignment Project Exam Help

unnormalized probabilities

<https://powcoder.com>

Add WeChat powcoder

cat
car
frog

3.2
5.1
-1.7

exp

24.5
164.0
0.18

normalize

0.13
0.87
0.00

$$L_i = -\log(0.13) = 0.89$$

unnormalized log probabilities

probabilities

Softmax + Log likelihood loss:
is often called 'softmax classifier'

Softmax + Cross-entropy Loss



$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_{j=1}^C e^{s_j}}\right) - \sum_{k=1, k \neq y_i}^C \log\left(1 - \frac{e^{s_k}}{\sum_{j=1}^C e^{s_j}}\right)$$

Assignment Project Exam Help

unnormalized probabilities

<https://powcoder.com>

Add WeChat powcoder

cat
car
frog

3.2
5.1
-1.7

exp

24.5
164.0
0.18

normalize

0.13
0.87
0.00

$$L_i = -\log(0.13) - \log(1 - 0.87) - \log(1 - 0.0) = 0.89 * 2$$

unnormalized log probabilities

probabilities

Loss Function: Regression Loss

- Regression loss
 - Using L1 or L2 norms
 - Widely used in pixel-level prediction (e.g. image denoising)

$$L_i = |\mathbf{y}_i - \mathbf{s}_i|$$

Assignment Project Exam Help

$$L_i = (\mathbf{y}_i - \mathbf{s}_i)^2$$

<https://powcoder.com>

Add WeChat powcoder

$$\mathbf{y}_i = \begin{pmatrix} 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{pmatrix} \quad \mathbf{s}_i = \begin{pmatrix} 0.1 \\ 0.7 \\ 0 \\ \vdots \\ 0.2 \end{pmatrix} \quad \Rightarrow \quad L_i = |\mathbf{y}_i - \mathbf{s}_i| = |0 - 0.1| + |1 - 0.7| + |0 - 0.2|$$

Regularization

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1) \quad \mathbf{s} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

Suppose that we found a \mathbf{W} such that $L = 0$. Is this \mathbf{W} unique?

No! $2\mathbf{W}$ is also has $L = 0$!



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Before:

$$\begin{aligned} &= \max(0, 1.3 - 4.9 + 1) + \max(0, 2.0 - 4.9 + 1) \\ &= \max(0, -2.6) + \max(0, -1.9) \\ &= 0 + 0 = 0 \end{aligned}$$

cat	3.2	1.3	2.2
car	5.1	4.9	2.5
frog	-1.7	2.0	-3.1

0

With \mathbf{W} twice as large:

$$\begin{aligned} &= \max(0, 2.6 - 9.8 + 1) + \max(0, 4.0 - 9.8 + 1) \\ &= \max(0, -6.2) + \max(0, -4.8) \\ &= 0 + 0 = 0 \end{aligned}$$

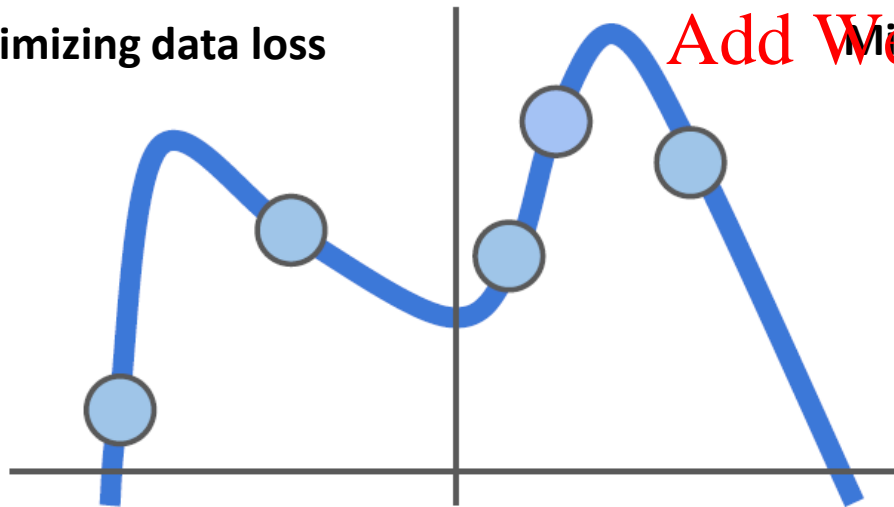
Regularization

$$L = \frac{1}{N} \sum_{i=1}^N L_i(f(\mathbf{x}_i, \mathbf{W}), y_i) + \lambda R(\mathbf{W})$$

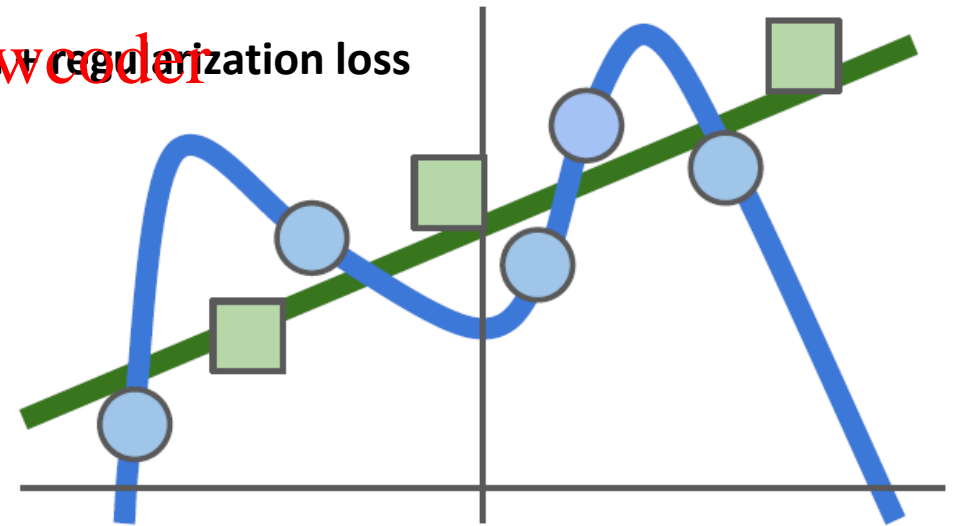
Data loss: Model predictions should match training data

Regularization: Model should be “simple” to avoid *overfitting*, so it works on test data

Minimizing data loss



Minimizing data + regularization loss



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Regularization

$$L = \frac{1}{N} \sum_{i=1}^N L_i(f(\mathbf{x}_i, \mathbf{W}), y_i) + \lambda R(\mathbf{W})$$

λ : regularization strength
(hyperparameter)

Assignment Project Exam Help

- L2 regularization: $R(\mathbf{W}) = \sum_{k,l} W_{k,l}^2$
- L1 regularization: $R(\mathbf{W}) = \sum_{k,l} |W_{k,l}|$
- Elastic net (L1 + L2): $R(\mathbf{W}) = \sum_{k,l} \beta W_{k,l}^2 + |W_{k,l}|$
- Max norm regularization: $|\mathbf{w}_j^T| < c$ **for all** j
- Dropout (will see later)
- Batch normalization, stochastic depth (will see later)

$$\mathbf{W} = \begin{pmatrix} \mathbf{w}_1^T \\ \mathbf{w}_2^T \\ \vdots \\ \mathbf{w}_C^T \end{pmatrix}$$


Optimization: Gradient Descent

- Gradient Descent

- The simplest approach to minimizing a loss function

$$\mathbf{w}^{T+1} = \mathbf{w}^T - \alpha \frac{\partial L}{\partial \mathbf{w}^T}$$

- α : step size (a.k.a. learning rate)



$$y = f(x) = x^2 - 4x + 3$$

$$y' = f'(x) = 2x - 4$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


```
# Vanilla Gradient Descent
```

```
while True:
```

```
    weights_grad = evaluate_gradient(loss_fun, data, weights)
```

```
    weights += - step_size * weights_grad # perform parameter update
```


Optimization: Gradient Descent



Optimization: Stochastic Gradient Descent (SGD)

$$L = \frac{1}{N} \sum_{i=1}^N L_i(f(\mathbf{x}_i, \mathbf{W}), y_i) + \lambda R(\mathbf{W})$$

Full sum is too expensive when N is large!

$$\frac{\partial L}{\partial \mathbf{W}} = \frac{1}{N} \sum_{i=1}^N \frac{\partial L_i(f(\mathbf{x}_i, \mathbf{W}), y_i)}{\partial \mathbf{W}} + \frac{\partial R(\mathbf{W})}{\partial \mathbf{W}}$$

Assignment Project Exam Help

<https://powcoder.com>

Instead, approximating sum using a minibatch of 32 / 64 / 128 / 256 examples is common

Add WeChat powcoder

```
# Vanilla Minibatch Gradient Descent
```

```
while True:
```

```
    data_batch = sample_training_data(data, 256) # sample 256 examples
```

```
    weights_grad = evaluate_gradient(loss_fun, data_batch, weights)
```

```
    weights += - step_size * weights_grad # perform parameter update
```

EBU7240

Computer Vision

Assignment Project Exam Help

- <https://powcoder.com> -
Backpropagation-

Add WeChat powcoder

Semester 1, 2021

Changjae Oh

Backpropagation

- A widely used algorithm for training feedforward neural networks.
- A way of computing gradients of expressions through recursive application of chain rule.
 - Backpropagation computes the gradient of the loss function with respect to the weights of the network (model) for a single input–output example.

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder

- **Gradient Descent**
 - The simplest approach to minimizing a loss function

$$\mathbf{w}^{T+1} = \mathbf{w}^T - \alpha \frac{\partial L}{\partial \mathbf{w}^T}$$

- α : step size (a.k.a. learning rate)



Derivative


- Optimization using derivative

- 1st order derivative

$$f'(x) = \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x}$$

- $f'(x)$: The slope of the function, indicating the direction in which the value increases
→ The minima of the objective function may exist in the direction of $-f'(x)$.

- Gradient descent algorithm: $-f'(x)$



$$y = f(x) = x^2 - 4x + 3$$

$$y' = f'(x) = 2x - 4$$

Add WeChat powcoder

$$\mathbf{w}^{T+1} = \mathbf{w}^T - \alpha \frac{\partial L}{\partial \mathbf{w}^T}$$

Derivative

- Partial derivative

- Derivatives of functions with multiple variables
- Gradient: the vector of the partial derivative

$$\text{Ex) } \nabla f, \frac{\partial f}{\partial \mathbf{x}}, \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2} \right)^T$$

Assignment Project Exam Help

$$f(\mathbf{x}) = f(x_1, x_2) = \left(4 - 2.1x_1^2 + \frac{x_1^4}{3} \right) x_1^2 + x_1 x_2 + (-4 + 4x_2^2) x_2^2$$

<https://powcoder.com>

$$\nabla f = f'(\mathbf{x}) = \frac{\partial f}{\partial \mathbf{x}} = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2} \right)^T = (2x_1^5 - 8.4x_1^3 + 8x_1 + x_2, 16x_2^3 - 8x_2 + x_1)^T$$

Add WeChat powcoder

Derivative

- **Jacobian matrix**

- 1st order partial derivative matrix for $\mathbf{f}: \mathbb{R}^d \mapsto \mathbb{R}^m$

$$\mathbf{J} = \begin{pmatrix} \frac{\partial f_1}{\partial x_1} & \frac{\partial f_1}{\partial x_2} & \dots & \frac{\partial f_1}{\partial x_d} \\ \frac{\partial f_2}{\partial x_1} & \frac{\partial f_2}{\partial x_2} & \dots & \frac{\partial f_2}{\partial x_d} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \frac{\partial f_m}{\partial x_2} & \dots & \frac{\partial f_m}{\partial x_d} \end{pmatrix}$$

Ex) $\mathbf{f}: \mathbb{R}^2 \mapsto \mathbb{R}^3$ $\mathbf{f}(\mathbf{x}) = (2x_1 + x_2^2, -x_1^2 + 3x_2, 4x_1x_2)^T$

$$\mathbf{J} = \begin{pmatrix} 2 & 2x_2 \\ -2x_1 & 3 \\ 4x_2 & 4x_1 \end{pmatrix} \quad \mathbf{J}|_{(2,1)^T} = \begin{pmatrix} 2 & 2 \\ -4 & 3 \\ 4 & 8 \end{pmatrix}$$

Assignment Project Exam Help

<https://powcoder.com>

- **Hessian matrix**

- 2nd order partial derivative matrix

$$\mathbf{H} = \begin{pmatrix} \frac{\partial^2 f}{\partial x_1 x_1} & \frac{\partial^2 f}{\partial x_1 x_2} & \dots & \frac{\partial^2 f}{\partial x_1 x_n} \\ \frac{\partial^2 f}{\partial x_2 x_1} & \frac{\partial^2 f}{\partial x_2 x_2} & \dots & \frac{\partial^2 f}{\partial x_2 x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n x_1} & \frac{\partial^2 f}{\partial x_n x_2} & \dots & \frac{\partial^2 f}{\partial x_n x_n} \end{pmatrix}$$

Ex) $f(\mathbf{x}) = f(x_1, x_2)$
 $= \left(4 - 2.1x_1^2 + \frac{x_1^4}{3}\right)x_1^2 + x_1x_2 + (-4 + 4x_2^2)x_2^2$

$$\mathbf{H} = \begin{pmatrix} 10x_1^4 - 25.2x_1^2 + 8 & 1 \\ 1 & 48x_2^2 - 8 \end{pmatrix}$$

$$\mathbf{H}|_{(0,1)^T} = \begin{pmatrix} 8 & 1 \\ 1 & 40 \end{pmatrix}$$

Add WeChat powcoder

Derivative


- Chain rule

$$\begin{aligned} f(x) &= g(h(x)) && \longrightarrow && f'(x) = g'(h(x))h'(x) \\ f(x) &= g(h(i(x))) && && f'(x) = g'(h(i(x)))h'(i(x))i'(x) \end{aligned}$$

Ex) $f(x) = 3(2x^2 - 1)^2 - 2(2x^2 - 1) + 5$ Assignment Project Exam Help $h(x) = 2x^2 - 1$

$\longrightarrow f'(x) = \underbrace{(3 * 2(2x^2 - 1) - 2)}_{g'(h(x))} \underbrace{(2 * 2x)}_{h'(x)} = 48x^3 - 32x$ https://powcoder.com

Add WeChat powcoder




Why are we talking about derivatives?

- **Gradient Descent**

- The simplest approach to minimizing a loss function

$$\mathbf{w}^{T+1} = \mathbf{w}^T - \alpha \frac{\partial L}{\partial \mathbf{w}^T}$$

- α : step size (a.k.a. learning rate)



$$y = f(x) = x^2 - 4x + 3$$

$$y' = f'(x) = 2x - 4$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

```
# Vanilla Gradient Descent
```

```
while True:
```

```
    weights_grad = evaluate_gradient(loss_fun, data, weights)
```

```
    weights += - step_size * weights_grad # perform parameter update
```

Derivative

- Example) Applying chain rule to single-layer perceptron

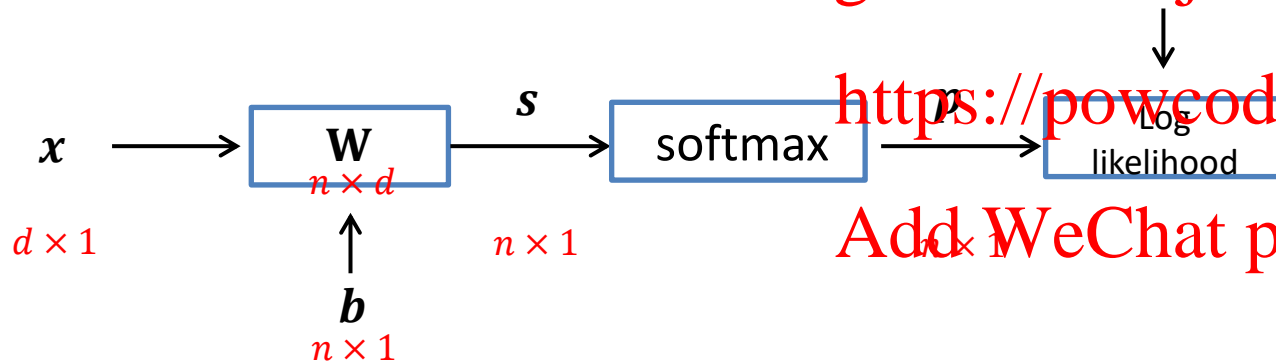
- Example of composite function

- Back-propagation: use the chain rule to compute $\frac{\partial L}{\partial \mathbf{w}}$ and $\frac{\partial L}{\partial \mathbf{b}}$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



$$\frac{\partial L}{\partial \mathbf{p}} \Rightarrow \frac{\partial L}{\partial \mathbf{s}} = \frac{\partial \mathbf{p}}{\partial \mathbf{s}} \frac{\partial L}{\partial \mathbf{p}} \Rightarrow \frac{\partial L}{\partial \mathbf{w}_j} = \frac{\partial \mathbf{s}}{\partial \mathbf{w}_j} \frac{\partial L}{\partial \mathbf{s}} \quad \frac{\partial L}{\partial \mathbf{w}} = \left(\frac{\partial L}{\partial \mathbf{w}_1} \quad \frac{\partial L}{\partial \mathbf{w}_2} \quad \dots \quad \frac{\partial L}{\partial \mathbf{w}_n} \right)^T$$

$$\frac{\partial L}{\partial \mathbf{b}} = \frac{\partial \mathbf{s}}{\partial \mathbf{b}} \frac{\partial L}{\partial \mathbf{s}} = \frac{\partial L}{\partial \mathbf{s}}$$

Analytic Gradient: Linear Equation

$$\mathbf{s} = \mathbf{W}\mathbf{x} + \mathbf{b} \longleftrightarrow \begin{aligned} s_1 &= \mathbf{w}_1^T \mathbf{x} + b_1 \\ s_2 &= \mathbf{w}_2^T \mathbf{x} + b_2 \\ &\vdots \\ s_n &= \mathbf{w}_n^T \mathbf{x} + b_n \end{aligned}$$

$$\mathbf{s} = \begin{pmatrix} s_1 \\ s_2 \\ \vdots \\ s_n \end{pmatrix} \quad \mathbf{W} = \begin{pmatrix} \mathbf{w}_1^T \\ \mathbf{w}_2^T \\ \vdots \\ \mathbf{w}_n^T \end{pmatrix} = \begin{pmatrix} w_{11} & w_{12} & \cdots & w_{1d} \\ w_{21} & w_{22} & \cdots & w_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n1} & w_{n2} & \cdots & w_{nd} \end{pmatrix} \quad \mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{pmatrix} \quad \mathbf{b} = \begin{pmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{pmatrix}$$

Assignment Project Exam Help

$$\frac{\partial s_1}{\partial \mathbf{w}_1} = \mathbf{x}$$


<https://powcoder.com>

$$\frac{\partial s_2}{\partial \mathbf{w}_1} = \mathbf{0}$$

Add WeChat powcoder

$$\vdots$$

$$\frac{\partial s_n}{\partial \mathbf{w}_1} = \mathbf{0}$$



$$\frac{\partial \mathbf{s}}{\partial \mathbf{w}_1} = [\mathbf{x} \ \mathbf{0} \ \mathbf{0} \ \cdots \ \mathbf{0}] \in \mathbb{R}^{d \times n}$$

$$\frac{\partial \mathbf{s}}{\partial \mathbf{b}} = \begin{pmatrix} 1 & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & 1 \end{pmatrix} = \mathbf{I} \in \mathbb{R}^{n \times n}$$

$$\frac{\partial \mathbf{s}}{\partial \mathbf{w}_j} = [\mathbf{0} \ \mathbf{0} \ \mathbf{x} \ \cdots \ \mathbf{0}] \in \mathbb{R}^{d \times n}$$

jth column

Analytic Gradient: Linear Equation

$$\mathbf{s} = \mathbf{W}\mathbf{x} + \mathbf{b} \quad \longleftrightarrow \quad \begin{aligned} s_1 &= \mathbf{w}_1^T \mathbf{x} + b_1 \\ s_2 &= \mathbf{w}_2^T \mathbf{x} + b_2 \\ &\vdots \\ s_n &= \mathbf{w}_n^T \mathbf{x} + b_n \end{aligned}$$
$$\mathbf{s} = \begin{pmatrix} s_1 \\ s_2 \\ \vdots \\ s_n \end{pmatrix} \quad \mathbf{W} = \begin{pmatrix} \mathbf{w}_1^T \\ \mathbf{w}_2^T \\ \vdots \\ \mathbf{w}_n^T \end{pmatrix} = \begin{pmatrix} w_{11} & w_{12} & \cdots & w_{1d} \\ w_{21} & w_{22} & \cdots & w_{2d} \\ & \vdots & & \\ w_{n1} & w_{n2} & \cdots & w_{nd} \end{pmatrix} \quad \mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{pmatrix} \quad \mathbf{b} = \begin{pmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{pmatrix}$$

Assignment Project Exam Help

$$\frac{\partial s_1}{\partial \mathbf{x}} = \mathbf{w}_1$$


$$\frac{\partial s_2}{\partial \mathbf{x}} = \mathbf{w}_2$$

\vdots

$$\frac{\partial s_n}{\partial \mathbf{x}} = \mathbf{w}_n$$

<https://powcoder.com>

Add WeChat powcoder



$$\frac{\partial \mathbf{s}}{\partial \mathbf{x}} = [\mathbf{w}_1 \ \mathbf{w}_2 \ \cdots \ \mathbf{w}_n] = \mathbf{W}^T \in \mathbb{R}^{d \times n}$$

Analytic Gradient: Sigmoid Function

- Sigmoid function

For a scalar x

$$\sigma(x) = \frac{1}{1 + e^{-x}} \rightarrow \frac{\partial \sigma(x)}{\partial x} = \frac{e^{-x}}{(1 + e^{-x})^2} = \frac{1 + e^{-x} - 1}{1 + e^{-x}} \frac{1}{1 + e^{-x}} = (1 - \sigma(x))\sigma(x)$$

Similarly, for a vector $\mathbf{s} \in \mathbb{R}^{n \times 1}$

$$\mathbf{p} = \sigma(\mathbf{s}) = \frac{1}{1 + e^{-s}} \rightarrow \frac{\partial \mathbf{p}}{\partial \mathbf{s}} = \text{diag}((1 - \sigma(s_j))\sigma(s_j)) = \begin{bmatrix} (1 - \sigma(s_1))\sigma(s_1) & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & (1 - \sigma(s_n))\sigma(s_n) \end{bmatrix}$$

for $j = 1, \dots, n$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Analytic Gradient: Softmax Activation Function

- Softmax function

$$p_k = \frac{e^{s_k}}{\sum_{j=1}^n e^{s_j}} \quad \Rightarrow \quad \mathbf{p} = \frac{e^{\mathbf{s}}}{\sum_{j=1}^n e^{s_j}} \quad \text{in vector form}$$

score function

$$\mathbf{s} = \begin{pmatrix} s_1 \\ s_2 \\ \vdots \\ s_n \end{pmatrix}$$

probability

$$\mathbf{p} = \begin{pmatrix} p_1 \\ p_2 \\ \vdots \\ p_n \end{pmatrix}$$

<https://powcoder.com>

- 1st order derivative of softmax function

Add WeChat powcoder

$$\frac{\partial \mathbf{p}}{\partial \mathbf{s}} = \frac{\text{diag}(e^{\mathbf{s}}) \cdot \sum e^{s_j} - e^{\mathbf{s}}(e^{\mathbf{s}})^T}{(\sum e^{s_j})^2} = \frac{1}{(\sum e^{s_j})^2} \left\{ \begin{pmatrix} e^{s_1} \sum e^{s_j} & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & e^{s_n} \sum e^{s_j} \end{pmatrix} - \begin{pmatrix} e^{s_1} e^{s_1} & \dots & e^{s_1} e^{s_n} \\ \vdots & \ddots & \vdots \\ e^{s_n} e^{s_1} & \dots & e^{s_n} e^{s_n} \end{pmatrix} \right\}$$

Analytic Gradient: Softmax Activation Function

$$\mathbf{D} = \frac{\partial \mathbf{p}}{\partial \mathbf{s}} = \frac{\text{diag}(e^{\mathbf{s}}) \cdot \sum e^{s_j} - e^{\mathbf{s}}(e^{\mathbf{s}})^T}{(\sum e^{s_j})^2} = \frac{1}{(\sum e^{s_j})^2} \left\{ \begin{pmatrix} e^{s_1} \sum e^{s_j} & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & e^{s_n} \sum e^{s_j} \end{pmatrix} - \begin{pmatrix} e^{s_1} e^{s_1} & \dots & e^{s_1} e^{s_n} \\ \vdots & \ddots & \vdots \\ e^{s_n} e^{s_1} & \dots & e^{s_n} e^{s_n} \end{pmatrix} \right\}$$

Assignment Project Exam Help

For $a = b$

$$\frac{e^{s_a} (\sum e^{s_j} - e^{s_a})}{(\sum e^{s_j})^2} = p_a(1 - p_a)$$

For $a \neq b$

<https://powcoder.com>

$$-\frac{e^{s_a} e^{s_b}}{(\sum e^{s_j})^2} = -p_a p_b$$

Add WeChat powcoder

$$D_{ab} = p_a(\delta_{ab} - p_b)$$

$$\delta_{ab} = \begin{cases} 1 & a = b \\ 0 & \text{otherwise} \end{cases}$$

Analytic Gradient: Hinge Loss

For simplicity of notation, the index of training image i is omitted here

- 1st order derivative of binary hinge loss

$$L = \max(0, 1 - y \cdot s)$$


$$s = \mathbf{w}^T \mathbf{x} + b$$

$y = \pm 1$ for positive/negative samples

$$= \begin{cases} 1 - y \cdot s & \text{if } 1 - y \cdot s > 0 \\ 0 & \text{otherwise} \end{cases}$$

<https://powcoder.com>

⇒ $\frac{\partial L}{\partial s} = \begin{cases} -y & \text{if } 1 - y \cdot s > 0 \\ 0 & \text{otherwise} \end{cases}$



Analytic Gradient: Hinge Loss

For simplicity of notation, the index of training image i is omitted here

- 1st order derivative of hinge loss

$$L = \sum_{j=1, j \neq y}^n \max(0, s_j - s_y + 1)$$

Assignment Project Exam Help

x : image

y : class label (integer, $1 \leq y \leq n$)

<https://powcoder.com>

Add WeChat powcoder

$$\mathbf{s} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$$\mathbf{W} = \begin{pmatrix} \mathbf{w}_1^T \\ \mathbf{w}_2^T \\ \vdots \\ \mathbf{w}_n^T \end{pmatrix} \quad \mathbf{s} = \begin{pmatrix} s_1 \\ s_2 \\ \vdots \\ s_n \end{pmatrix}$$



$$\frac{\partial L}{\partial s_y} = - \sum_{j=1, j \neq y}^n 1(s_j - s_y + 1 > 0) \quad \text{for } j = y$$

$$1(F) = \begin{cases} 1 & \text{if } F \text{ is true} \\ 0 & \text{otherwise} \end{cases}$$

$$\frac{\partial L}{\partial s_j} = 1(s_j - s_y + 1 > 0) \quad \text{for } j \neq y$$

Analytic Gradient: Log Likelihood Loss

For simplicity of notation, the index of training image i is omitted here

$$L = -\log p_y \text{ where } y \text{ satisfies } z_y = 1$$
$$\mathbf{p} = \begin{pmatrix} p_1 \\ p_2 \\ \vdots \\ p_n \end{pmatrix} \begin{array}{l} \text{probability for } i^{th} \text{ image} \\ \text{(It is assumed to be } \textcolor{red}{normalized}, \text{ i.e. } |\mathbf{p}| = 1.) \end{array}$$

$$\Rightarrow \frac{\partial L}{\partial \mathbf{p}} = - \begin{bmatrix} 0 \\ \vdots \\ 1/p_y \\ \vdots \\ 0 \end{bmatrix}$$

y^{th} row

y : class label for i^{th} image ($1 \leq y \leq n$)

\mathbf{z} : class probability for i^{th} image
 $\mathbf{z} = (z_1, z_2, \dots, z_n)^T, z_y = 1 \text{ and } z_{k \neq y} = 0$

<https://powcoder.com>

Add WeChat powcoder

Example

Suppose i^{th} image belongs to class 2 and $n = 10$. $\rightarrow y = 2$

$$\mathbf{z} = \begin{pmatrix} 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{pmatrix}$$

$$\mathbf{p} = \begin{pmatrix} 0.1 \\ 0.7 \\ 0 \\ \vdots \\ 0.2 \end{pmatrix}$$



$$\frac{\partial L}{\partial \mathbf{p}} = - \begin{bmatrix} 0 \\ 1/0.7 \\ \vdots \\ 0 \end{bmatrix}$$

Analytic Gradient: Cross-entropy Loss

For simplicity of notation, the index of training image i is omitted here

$$L = - \sum_{j=1}^n (z_j \log p_j + (1 - z_j) \log(1 - p_j)) \quad \mathbf{p} = \begin{pmatrix} p_1 \\ p_2 \\ \vdots \\ p_n \end{pmatrix} \begin{array}{l} \text{probability for } i^{th} \text{ image} \\ \text{(It is assumed to be } \textcolor{red}{\textit{normalized}}, \text{ i.e. } |\mathbf{p}| = 1.) \end{array}$$

$$\Rightarrow \frac{\partial L}{\partial \mathbf{p}} = - \begin{bmatrix} 1/(1 - p_1) \\ \vdots \\ 1/p_y \\ \vdots \\ 1/(1 - p_n) \end{bmatrix}$$

$y^{th} \text{ row}$

y : class label for i^{th} image ($1 \leq y \leq n$)

\mathbf{z} : class probability for i^{th} image

$\mathbf{z} = (z_1, z_2, \dots, z_n)^T, z_y = 1 \text{ and } z_{k \neq y} = 0$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Example

Suppose i^{th} image belongs to class 2 and $n = 10$. $\rightarrow y = 2$

$$\mathbf{z} = \begin{pmatrix} 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{pmatrix}$$

$$\mathbf{p} = \begin{pmatrix} 0.1 \\ 0.7 \\ 0 \\ \vdots \\ 0.2 \end{pmatrix}$$



$$\frac{\partial L}{\partial \mathbf{p}} = - \begin{bmatrix} 1/(1 - 0.1) \\ 1/0.7 \\ 0 \\ \vdots \\ 1/(1 - 0.2) \end{bmatrix}$$

Analytic Gradient: Regression Loss

For simplicity of notation, the index of training image i is omitted here

- Regression loss

$$L = (\mathbf{y} - \mathbf{s})^2 \quad \Rightarrow \quad \frac{\partial L}{\partial \mathbf{s}} = -2(\mathbf{y} - \mathbf{s}) \quad \mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix} \quad \mathbf{s} = \begin{pmatrix} s_1 \\ s_2 \\ \vdots \\ s_n \end{pmatrix}$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


Why are we talking about derivatives?

- **Gradient Descent**

- The simplest approach to minimizing a loss function

$$\mathbf{w}^{T+1} = \mathbf{w}^T - \alpha \frac{\partial L}{\partial \mathbf{w}^T}$$

- α : step size (a.k.a. learning rate)



$$y = f(x) = x^2 - 4x + 3$$

$$y' = f'(x) = 2x - 4$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

```
# Vanilla Gradient Descent
```

```
while True:
```

```
    weights_grad = evaluate_gradient(loss_fun, data, weights)
```

```
    weights += - step_size * weights_grad # perform parameter update
```