

EBU7240

Computer Vision

- Multi-layer Perception (MLP)-

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Semester 1, 2021

Changjae Oh

Neural networks



(Before) Linear score function: $f = Wx$

(Now) 2-layer Neural Network: $f = W_2 \max(0, W_1 x)$

3-layer Neural Network: $f = W_3 \max(0, W_2 \max(0, W_1 x))$

<https://powcoder.com>

Add WeChat powcoder



Activation functions

- Adding non-linearities into neural networks, allowing the neural networks to learn powerful operations.
- A crucial component of deep learning
 - If the activation functions were to be removed from a feedforward neural network, the entire network could be re-factored to a simple linear operation or matrix transformation on its input
 - It would no longer be capable of performing complex tasks such as image recognition.

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder

Activation functions


Sigmoid

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$




Leaky ReLU

$$\max(0.1x, x)$$



tanh

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



Maxout


$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

<https://powcoder.com>

Add WeChat powcoder


ReLU

$$\max(0, x)$$




ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$




Neural networks: Architectures

“2-layer Neural Net”, or
“1-hidden-layer Neural Net”



“3-layer Neural Net”, or
“2-hidden-layer Neural Net”



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

“Fully-connected” layers

Derivative of Neural Net using Chain Rules

- **Example**

- 1-layer Neural Net (L2 regression loss)
- 2-layer Neural Net (L2 regression loss)


Assignment Project Exam Help

- 1-layer Neural Net (Softmax classifier)

- 2-layer Neural Net (Softmax classifier)

<https://powcoder.com>
Add WeChat powcoder

1. 1-layer Neural Net (L2 regression loss)



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



1. Linear score $\mathbf{s} = \mathbf{W}\mathbf{x} + \mathbf{b} \longleftrightarrow s_j = \mathbf{w}_j^T \mathbf{x} + b_j$

2. Activation function $\mathbf{p} = \sigma(\mathbf{s}) = \frac{1}{1 + e^{-\mathbf{s}}}$

3. Loss $L = (\mathbf{z} - \mathbf{p})^2$

$$\mathbf{s} = \begin{pmatrix} s_1 \\ s_2 \\ \vdots \\ s_n \end{pmatrix} \quad \mathbf{W} = \begin{pmatrix} \mathbf{w}_1^T \\ \mathbf{w}_2^T \\ \vdots \\ \mathbf{w}_n^T \end{pmatrix} = \begin{pmatrix} w_{11} & w_{12} & \cdots & w_{1d} \\ w_{21} & w_{22} & \cdots & w_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n1} & w_{n2} & \cdots & w_{nd} \end{pmatrix} \quad \mathbf{b} = \begin{pmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{pmatrix} \quad \mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{pmatrix}$$

1. 1-layer Neural Net (L2 regression loss)




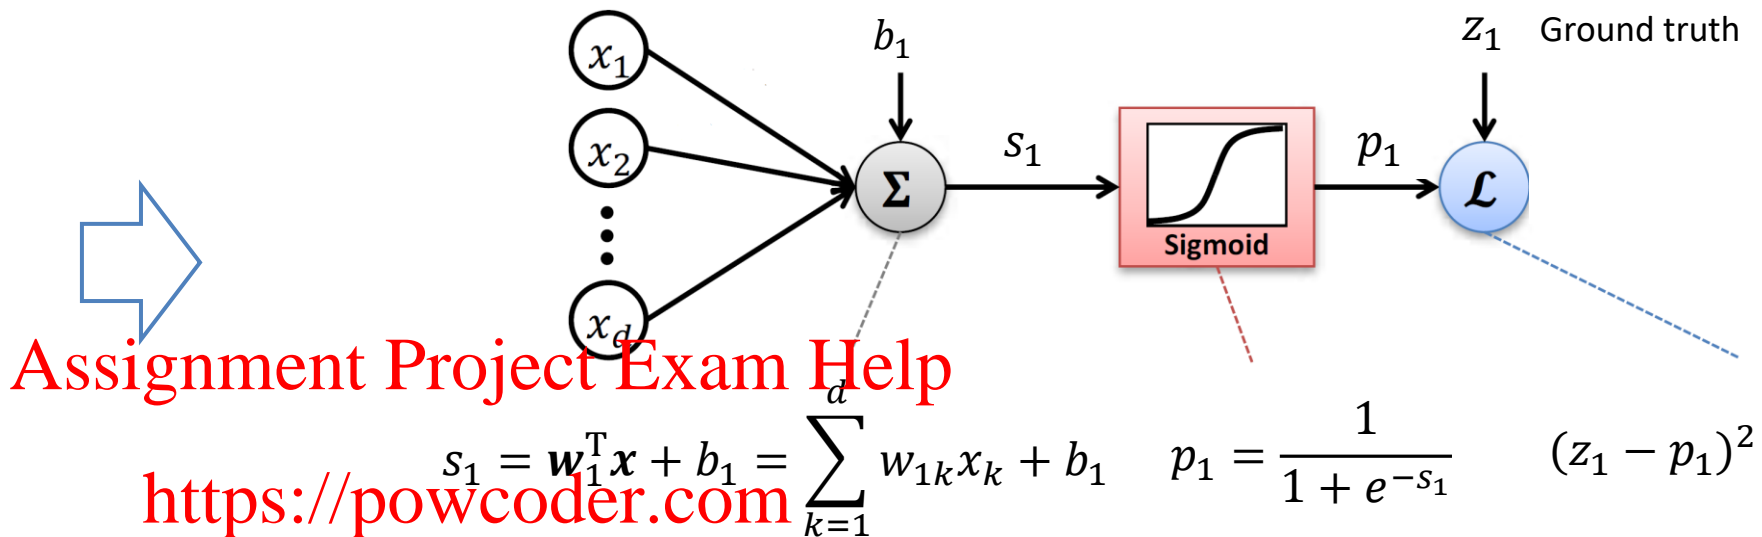

1. Linear score $\mathbf{s} = \mathbf{W}\mathbf{x} + \mathbf{b} \longleftrightarrow s_j = \mathbf{w}_j^T \mathbf{x} + b_j$

2. Activation function $\mathbf{p} = \sigma(\mathbf{s}) = \frac{1}{1 + e^{-\mathbf{s}}}$

3. Loss $L = (\mathbf{z} - \mathbf{p})^2$

$$\mathbf{s} = \begin{pmatrix} s_1 \\ s_2 \\ \vdots \\ s_n \end{pmatrix} \quad \mathbf{W} = \begin{pmatrix} \mathbf{w}_1^T \\ \mathbf{w}_2^T \\ \vdots \\ \mathbf{w}_n^T \end{pmatrix} = \begin{pmatrix} w_{11} & w_{12} & \cdots & w_{1d} \\ w_{21} & w_{22} & \cdots & w_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n1} & w_{n2} & \cdots & w_{nd} \end{pmatrix} \quad \mathbf{b} = \begin{pmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{pmatrix} \quad \mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{pmatrix}$$


1. 1-layer Neural Net (L2 regression loss)




1. Linear score $\mathbf{s} = \mathbf{W}\mathbf{x} + \mathbf{b} \longleftrightarrow s_j = \mathbf{w}_j^T \mathbf{x} + b_j$
2. Activation function $\mathbf{p} = \sigma(\mathbf{s}) = \frac{1}{1 + e^{-\mathbf{s}}}$
3. Loss $L = (\mathbf{z} - \mathbf{p})^2$

$$\mathbf{s} = \begin{pmatrix} s_1 \\ s_2 \\ \vdots \\ s_n \end{pmatrix} \quad \mathbf{W} = \begin{pmatrix} \mathbf{w}_1^T \\ \mathbf{w}_2^T \\ \vdots \\ \mathbf{w}_n^T \end{pmatrix} = \begin{pmatrix} w_{11} & w_{12} & \cdots & w_{1d} \\ w_{21} & w_{22} & \cdots & w_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n1} & w_{n2} & \cdots & w_{nd} \end{pmatrix} \quad \mathbf{b} = \begin{pmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{pmatrix} \quad \mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{pmatrix}$$

1. 1-layer Neural Net (L2 regression loss)



In a vector form



Assignment Project Exam Help

<https://powcoder.com>

We need to compute gradients of $\mathbf{W}, \mathbf{b}, \mathbf{s}, \mathbf{p}$ with respect to the loss function L .

1. Linear score $\mathbf{s} = \mathbf{W}\mathbf{x} + \mathbf{b} \longleftrightarrow s_j = \mathbf{w}_j^T \mathbf{x} + b_j$

2. Activation function $\mathbf{p} = \sigma(\mathbf{s}) = \frac{1}{1 + e^{-\mathbf{s}}}$

3. Loss $L = (\mathbf{z} - \mathbf{p})^2$

$$\mathbf{s} = \begin{pmatrix} s_1 \\ s_2 \\ \vdots \\ s_n \end{pmatrix} \quad \mathbf{W} = \begin{pmatrix} \mathbf{w}_1^T \\ \mathbf{w}_2^T \\ \vdots \\ \mathbf{w}_n^T \end{pmatrix} = \begin{pmatrix} w_{11} & w_{12} & \cdots & w_{1d} \\ w_{21} & w_{22} & \cdots & w_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ w_{n1} & w_{n2} & \cdots & w_{nd} \end{pmatrix} \quad \mathbf{b} = \begin{pmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{pmatrix} \quad \mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{pmatrix}$$

1. 1-layer Neural Net (L2 regression loss)

- Gradient Descent**

- The simplest approach to minimizing a loss function

$$\mathbf{W}^{T+1} = \mathbf{W}^T - \alpha \frac{\partial L}{\partial \mathbf{W}^T}$$

- α : step size (a.k.a. learning rate)

$$\frac{\partial L}{\partial \mathbf{p}} = -2(\mathbf{z} - \mathbf{p})$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

$$\frac{\partial L}{\partial \mathbf{p}} = -2 \begin{bmatrix} (1 - \sigma(s_1))\sigma(s_1)(z_1 - p_1) \\ (1 - \sigma(s_2))\sigma(s_2)(z_2 - p_2) \\ \vdots \\ (1 - \sigma(s_n))\sigma(s_n)(z_n - p_n) \end{bmatrix}$$

\otimes : element-wise multiplication

$$\frac{\partial L}{\partial \mathbf{s}} = \frac{\partial \mathbf{p}}{\partial \mathbf{s}} \frac{\partial L}{\partial \mathbf{p}} = \text{diag}((1 - \sigma(s_j))\sigma(s_j)) \frac{\partial L}{\partial \mathbf{p}} = (1 - \sigma(\mathbf{s})) \otimes \sigma(\mathbf{s}) \otimes \frac{\partial L}{\partial \mathbf{p}}$$


$$\frac{\partial L}{\partial \mathbf{w}_j} = \frac{\partial \mathbf{s}}{\partial \mathbf{w}_j} \frac{\partial L}{\partial \mathbf{s}} = \mathbf{X}_j \frac{\partial L}{\partial \mathbf{s}} = [\mathbf{0} \ \mathbf{0} \ \overset{\text{jth column}}{\mathbf{x}} \ \cdots \ \mathbf{0}] \frac{\partial L}{\partial \mathbf{s}} = \left(\frac{\partial L}{\partial \mathbf{s}} \right)_j \mathbf{x}$$

$(\mathbf{a})_j$: jth element at vector \mathbf{a}

$$\frac{\partial L}{\partial \mathbf{W}} = \left(\frac{\partial L}{\partial \mathbf{w}_1} \quad \frac{\partial L}{\partial \mathbf{w}_2} \quad \cdots \quad \frac{\partial L}{\partial \mathbf{w}_n} \right)^T = \frac{\partial L}{\partial \mathbf{s}} \mathbf{x}^T$$


$$\frac{\partial L}{\partial \mathbf{b}} = \frac{\partial \mathbf{s}}{\partial \mathbf{b}} \frac{\partial L}{\partial \mathbf{s}} = \frac{\partial L}{\partial \mathbf{s}}$$

In a vector form



1. 1-layer Neural Net (L2 regression loss)

In a vector form



Summary

$$\frac{\partial L}{\partial p} = -2(z - p)$$

$$\frac{\partial L}{\partial s} = \frac{\partial p}{\partial s} \frac{\partial L}{\partial p} = (1 - \sigma(s)) \otimes \sigma(s) \otimes \frac{\partial L}{\partial p}$$

$$\frac{\partial L}{\partial W} = \frac{\partial L}{\partial s} x^T$$

$$\frac{\partial L}{\partial b} = \frac{\partial L}{\partial s}$$

Assignment Project Exam Help

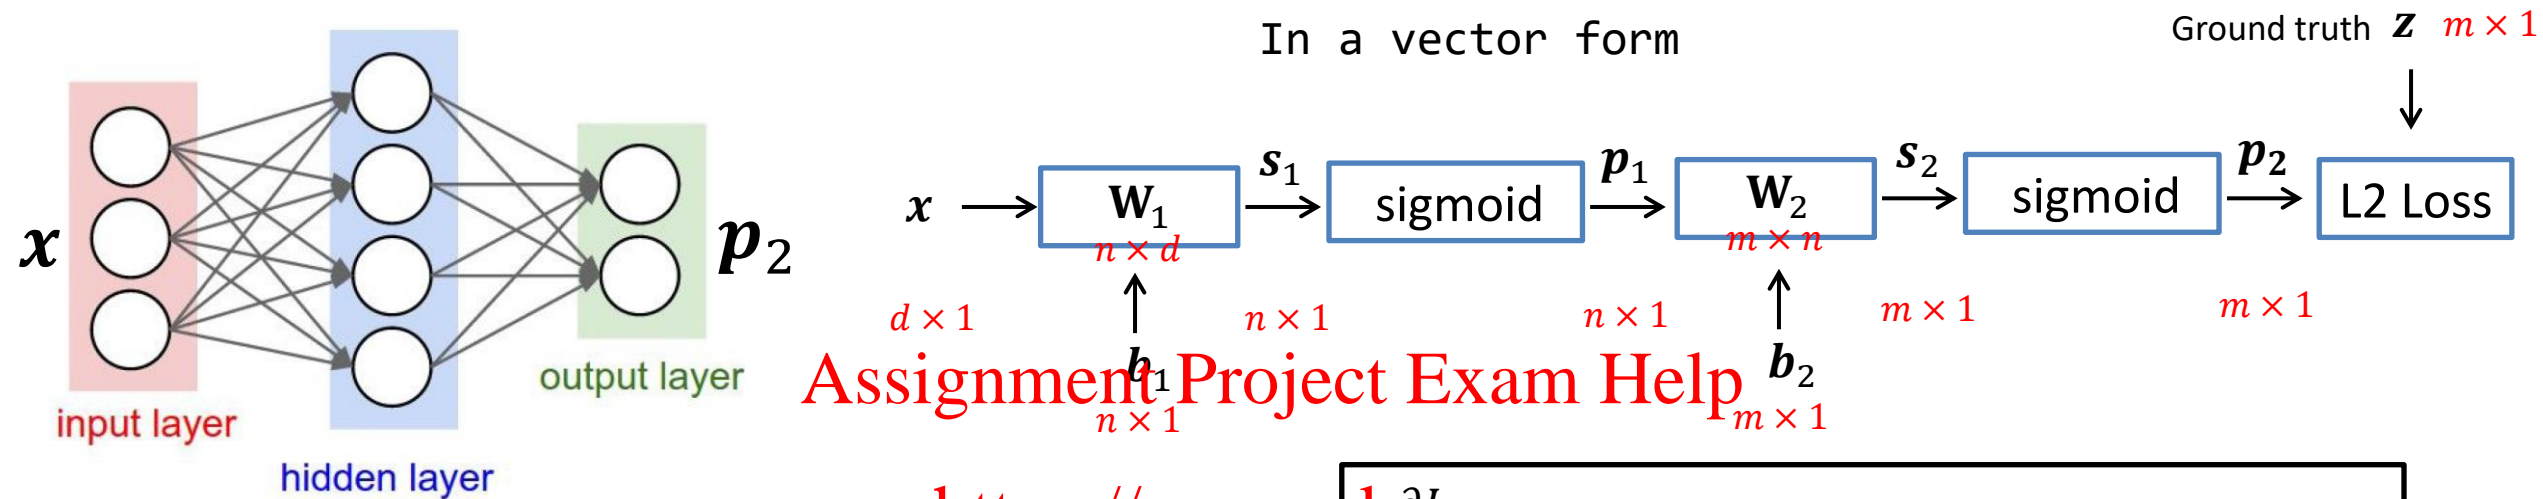
<https://powcoder.com>

Add WeChat powcoder

Note that the following derivative can also be computed, but here x is an input data that is fixed during training. Thus, it is not necessary to compute its derivative.

$$\frac{\partial L}{\partial x} = \frac{\partial s}{\partial x} \frac{\partial L}{\partial s} = W^T \frac{\partial L}{\partial s}$$

2. 2-layer Neural Net (L2 regression loss)



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

$$\frac{\partial L}{\partial s_1} = \frac{\partial p_1}{\partial s_1} \frac{\partial L}{\partial p_1} = \text{diag}((1 - \sigma(s_{1,j}))\sigma(s_{1,j})) \frac{\partial L}{\partial p_1}$$

$$\frac{\partial L}{\partial W_1} = \frac{\partial L}{\partial s_1} x^T \quad \frac{\partial L}{\partial b_1} = \frac{\partial L}{\partial s_1}$$

$$\frac{\partial L}{\partial p_2} = -2(z - p_2)$$

$$\frac{\partial L}{\partial s_2} = \frac{\partial p_2}{\partial s_2} \frac{\partial L}{\partial p_2} = \text{diag}((1 - \sigma(s_{2,j}))\sigma(s_{2,j})) \frac{\partial L}{\partial p_2}$$

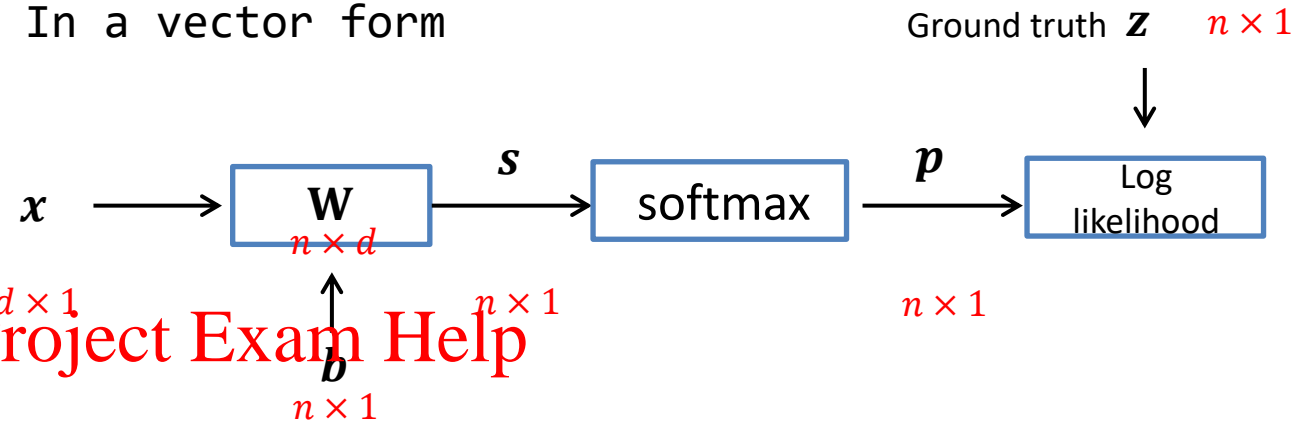
$$\frac{\partial L}{\partial W_2} = \frac{\partial L}{\partial s_2} p_1^T \quad \frac{\partial L}{\partial b_2} = \frac{\partial L}{\partial s_2}$$

$$\frac{\partial L}{\partial p_1} = \frac{\partial s_2}{\partial p_1} \frac{\partial L}{\partial s_2} = W_2^T \frac{\partial L}{\partial s_2}$$

3. 1-layer Neural Net (Softmax classifier)

$$\frac{\partial L}{\partial \mathbf{p}} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ -1/p_y \\ 0 \end{bmatrix} \quad \text{y}^{\text{th}} \text{ row}$$

In a vector form



$$\frac{\partial L}{\partial \mathbf{s}} = \frac{\partial \mathbf{p}}{\partial \mathbf{s}} \frac{\partial L}{\partial \mathbf{p}} = \mathbf{D} \frac{\partial L}{\partial \mathbf{p}} = -\frac{1}{p_y} \begin{bmatrix} D_{1y} \\ D_{2y} \\ \vdots \\ D_{ny} \end{bmatrix}$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat: powcoder

$$\frac{\partial L}{\partial \mathbf{w}_j} = \frac{\partial \mathbf{s}}{\partial \mathbf{w}_j} \frac{\partial L}{\partial \mathbf{s}} = \mathbf{x}_j \frac{\partial L}{\partial \mathbf{s}} = [\mathbf{0} \ \mathbf{0} \ \mathbf{x} \ \cdots \ \mathbf{0}] \frac{\partial L}{\partial \mathbf{s}} = \left(\frac{\partial L}{\partial \mathbf{s}} \right)_j \mathbf{x}$$

jth column

(\mathbf{a})_j: jth element at vector \mathbf{a}



$$\frac{\partial L}{\partial \mathbf{W}} = \left(\frac{\partial L}{\partial \mathbf{w}_1} \quad \frac{\partial L}{\partial \mathbf{w}_2} \quad \cdots \quad \frac{\partial L}{\partial \mathbf{w}_n} \right)^T = \frac{\partial L}{\partial \mathbf{s}} \mathbf{x}^T$$

$$\frac{\partial L}{\partial \mathbf{b}} = \frac{\partial \mathbf{s}}{\partial \mathbf{b}} \frac{\partial L}{\partial \mathbf{s}} = \frac{\partial L}{\partial \mathbf{s}}$$

3. 1-layer Neural Net (Softmax classifier)

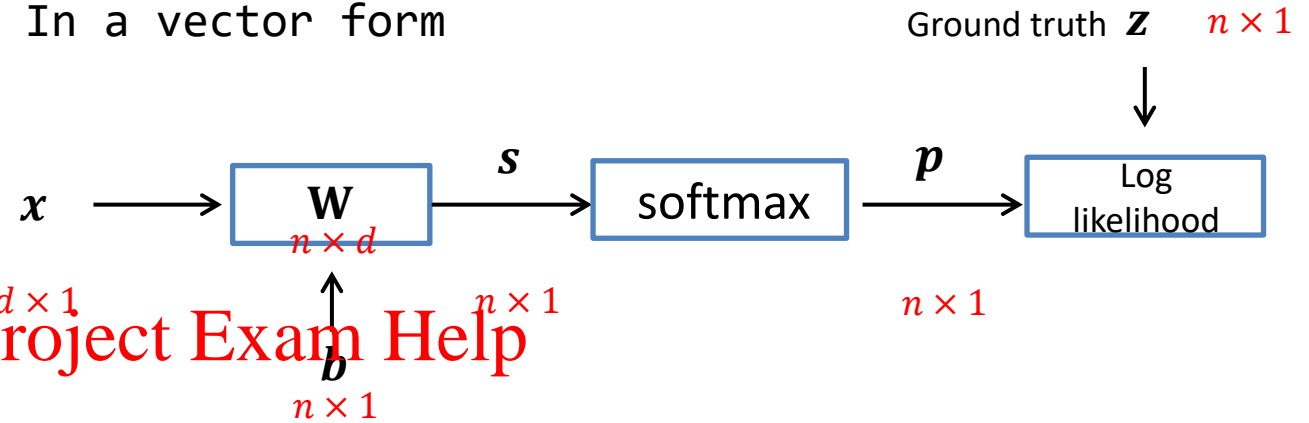
Summary

$$\frac{\partial L}{\partial \mathbf{p}} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ -1/p_y \\ 0 \end{bmatrix} \quad \text{y}^{\text{th}} \text{ row}$$

$$\frac{\partial L}{\partial \mathbf{s}} = \frac{\partial \mathbf{p}}{\partial \mathbf{s}} \frac{\partial L}{\partial \mathbf{p}} = \mathbf{D} \frac{\partial L}{\partial \mathbf{p}} = -\frac{1}{p_y} \begin{bmatrix} D_{1y} \\ D_{2y} \\ \vdots \\ D_{ny} \end{bmatrix} = \mathbf{p} - \mathbf{z}$$

$$\frac{\partial L}{\partial \mathbf{W}} = \frac{\partial L}{\partial \mathbf{s}} \mathbf{x}^T \quad \frac{\partial L}{\partial \mathbf{b}} = \frac{\partial L}{\partial \mathbf{s}}$$

In a vector form



Assignment Project Exam Help

<https://powcoder.com>

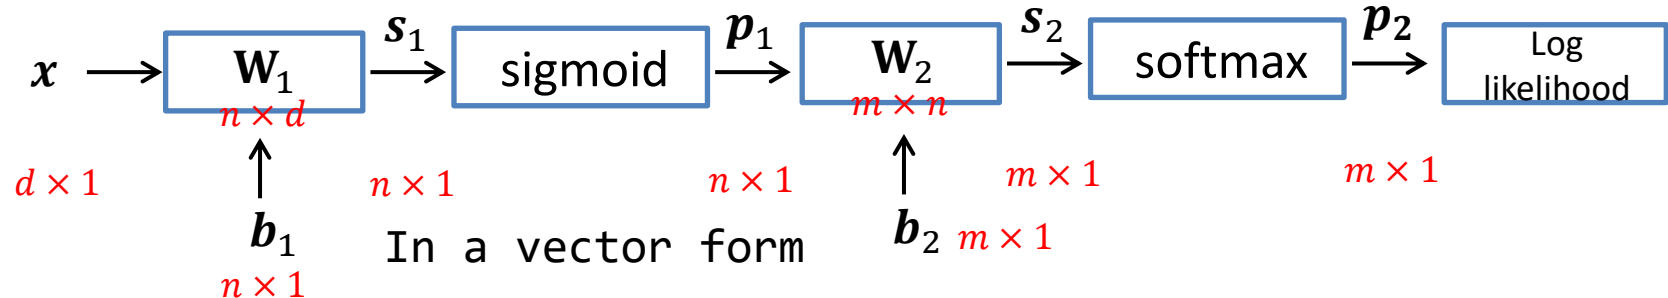

Add WeChat powcoder

Note that the following derivative can also be computed, but here \mathbf{x} is an input data that is fixed during training. Thus, it is not necessary to compute its derivative.

$$\frac{\partial L}{\partial \mathbf{x}} = \frac{\partial \mathbf{s}}{\partial \mathbf{x}} \frac{\partial L}{\partial \mathbf{s}} = \mathbf{W}^T \frac{\partial L}{\partial \mathbf{s}}$$

4. 2-layer Neural Net (Softmax classifier)

Ground truth \mathbf{z} $m \times 1$



Assignment Project Exam Help

<https://powcoder.com>

$$\frac{\partial L}{\partial \mathbf{p}_2} = \begin{bmatrix} 0 \\ 0 \\ -1/p_y \\ \vdots \\ 0 \end{bmatrix}$$

$y^{\text{th}} \text{ row}$

Add WeChat powcoder

$$\frac{\partial L}{\partial \mathbf{s}_1} = \frac{\partial \mathbf{p}_1}{\partial \mathbf{s}_1} \frac{\partial L}{\partial \mathbf{p}_1} = \text{diag}((1 - \sigma(s_{1,j}))\sigma(s_{1,j})) \frac{\partial L}{\partial \mathbf{p}_1}$$

$$\frac{\partial L}{\partial \mathbf{W}_1} = \frac{\partial L}{\partial \mathbf{s}_1} \mathbf{x}^T \quad \frac{\partial L}{\partial \mathbf{b}_1} = \frac{\partial L}{\partial \mathbf{s}_1}$$

$$\frac{\partial L}{\partial \mathbf{s}_2} = \frac{\partial \mathbf{p}_2}{\partial \mathbf{s}_2} \frac{\partial L}{\partial \mathbf{p}_2} = \mathbf{D} \frac{\partial L}{\partial \mathbf{p}_2}$$

$$D_{ab} = p_a(\delta_{ab} - p_b)$$

$$\delta_{ab} = \begin{cases} 1 & a = b \\ 0 & \text{otherwise} \end{cases}$$

$$\frac{\partial L}{\partial \mathbf{W}_2} = \frac{\partial L}{\partial \mathbf{s}_2} \mathbf{p}_1^T$$

$$\frac{\partial L}{\partial \mathbf{p}_1} = \frac{\partial \mathbf{s}_2}{\partial \mathbf{p}_1} \frac{\partial L}{\partial \mathbf{s}_2} = \mathbf{W}_2^T \frac{\partial L}{\partial \mathbf{s}_2}$$

$$\frac{\partial L}{\partial \mathbf{b}_2} = \frac{\partial L}{\partial \mathbf{s}_2}$$

Full implementation of training a 2-layer Neural Network

```
1 import numpy as np
2 from numpy.random import randn
3
4 N, D_in, H, D_out = 64, 1000, 100, 10
5 x, y = randn(N, D_in), randn(N, D_out)
6 w1, w2 = randn(D_in, H), randn(H, D_out)
7
8 for t in range(2000):
9     h = 1 / (1 + np.exp(-x.dot(w1)))
10    y_pred = h.dot(w2)
11    loss = np.square(y_pred - y).sum()
12    print(t, loss)
13
14    grad_y_pred = 2.0 * (y_pred - y)
15    grad_w2 = h.T.dot(grad_y_pred)
16    grad_h = grad_y_pred.dot(w2.T)
17    grad_w1 = x.T.dot(grad_h * h * (1 - h))
18
19    w1 -= 1e-4 * grad_w1
20    w2 -= 1e-4 * grad_w2
```

N: batch size

D_in: input feature size



H: input feature size of the second layer

D_out: output feature size

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Neural networks: Pros and cons

- **Pros**

- Flexible and general function approximation framework
- Can build extremely powerful models by adding more layers

- **Cons**

- Hard to analyze theoretically (e.g., training is prone to local optima)
- Huge amount of training data, computing power may be required to get good performance
- The space of implementation choices is huge (network architectures, parameters)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Summary

- We arrange neurons into fully-connected layers
- The layer allows us to use efficient vectorized code (e.g. matrix multiplication)

- Using back-propagation

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

EBU7240

Computer Vision

Assignment Project Exam Help

- Convolutional Neural Networks -

<https://powcoder.com>


Add WeChat powcoder

Semester 1, 2021

Changjae Oh

CNN Introduction

- Image Recognition
 - Recognizing the object class in the image



<http://www.cs.toronto.edu/~fritz/absps/imagenet.pdf>

CNN (= ConvNet)

- **ConvNet**

- is a sequence of layers
- Every layer of a ConvNet transforms one volume of activations to another through a differentiable function.

Assignment Project Exam Help

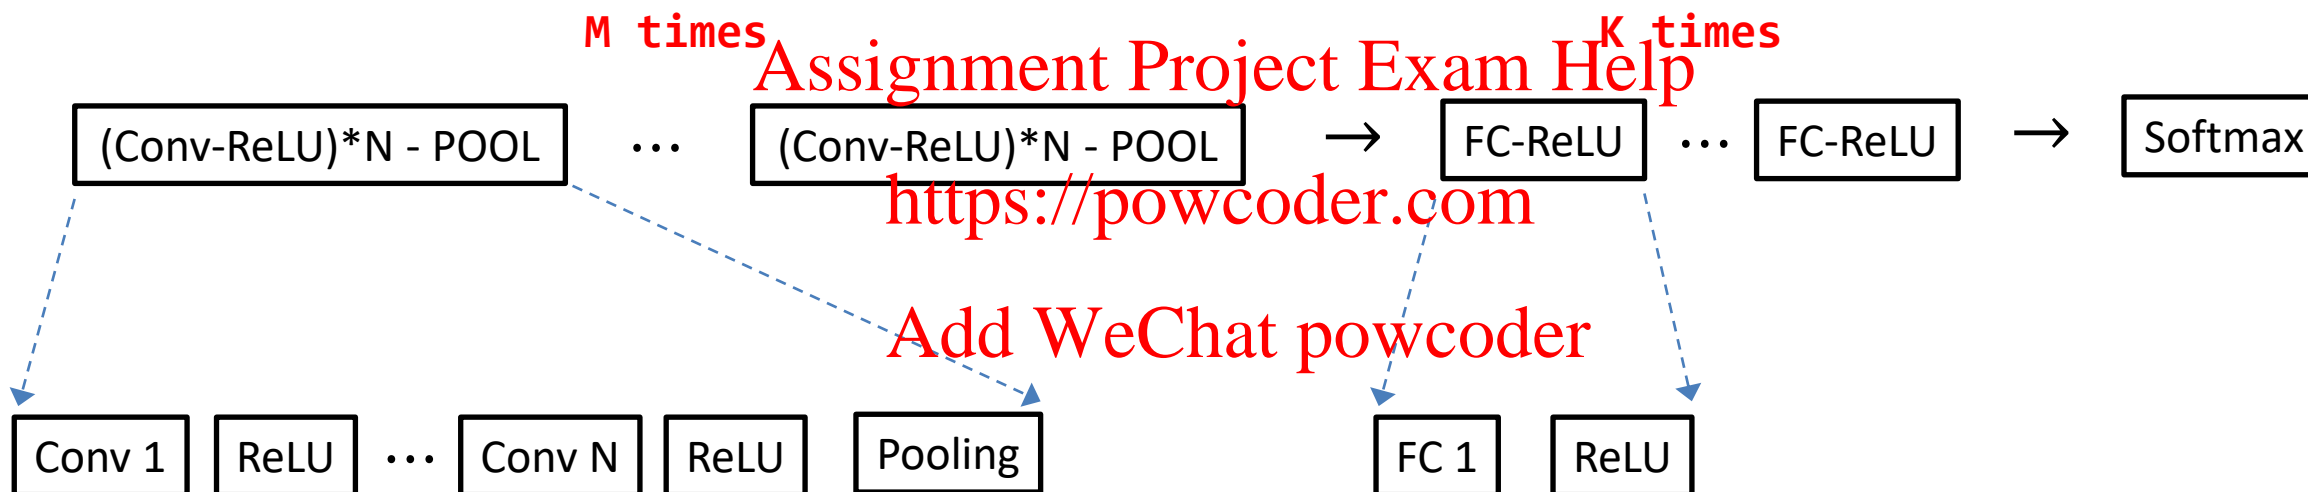
- **Convolutional Layer**: computes the output of neurons that are connected to local regions in the input
- **ReLU (nonlinear) layer**: activates relevant responses
- **Pooling Layer**: performs a downsampling operation along the spatial dimensions
- **Fully-Connected Layer**: each neuron in this layer will be connected to all the numbers in the previous volume

<https://powcoder.com>

Add WeChat powcoder

Typical architectures of ConvNet

$[(\text{Conv-ReLU})^*N - \text{POOL}]^*M - (\text{FC-RELU})^*K - \text{Softmax}$

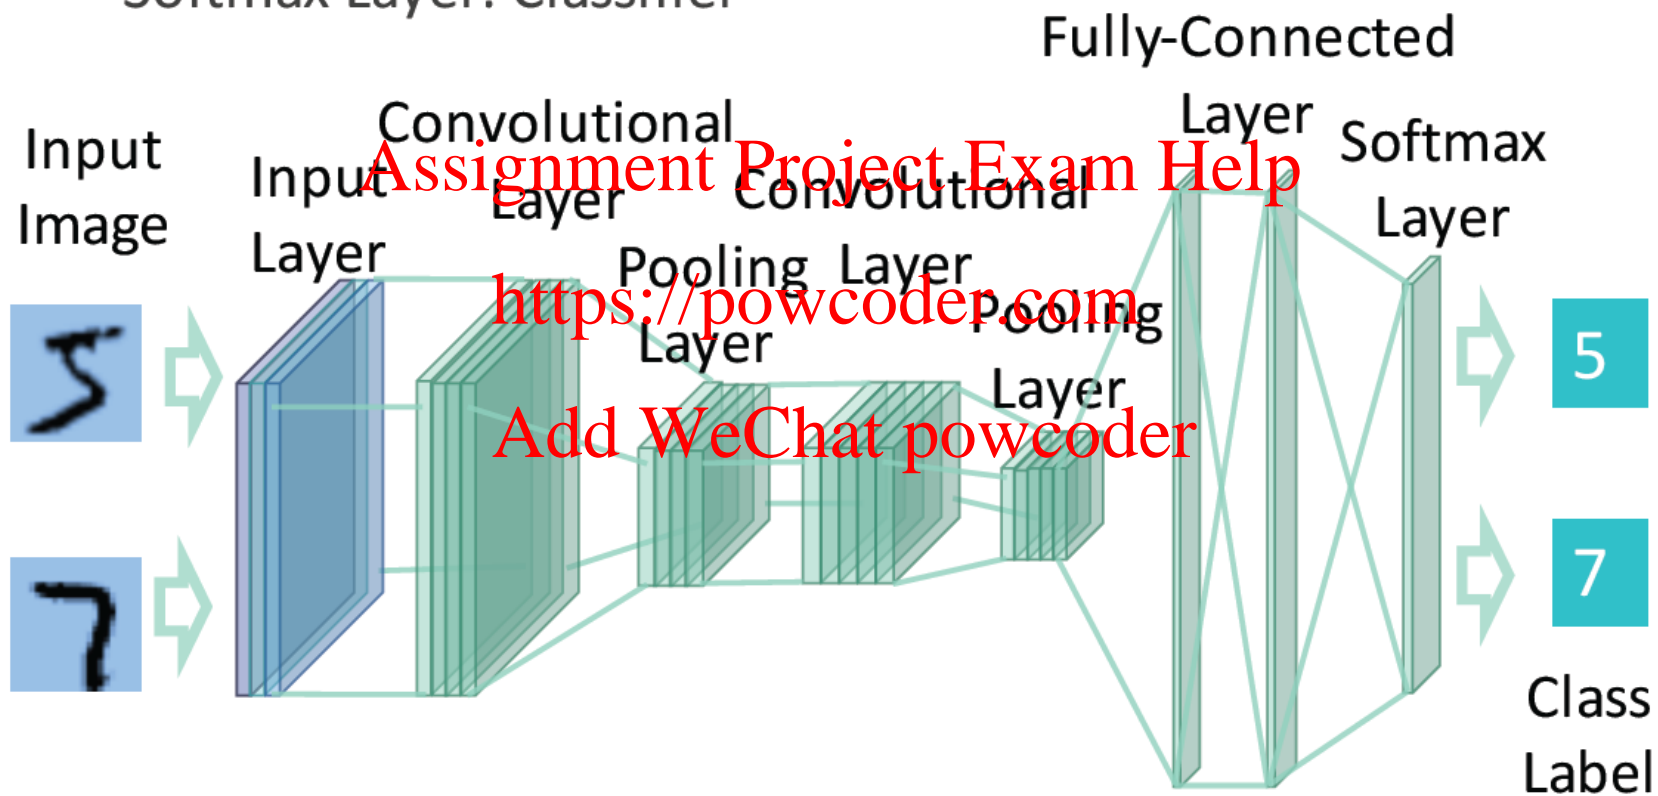


N is usually up to ~5, M is large, $0 \leq K \leq 2$

but some advances such as ResNet/GoogLeNet challenge this paradigm

Typical architectures of ConvNet

- Fully-Connected Layers : Global feature extraction
- Softmax Layer: Classifier



Assignment Project Exam Help


Convolutional Layer

<https://powcoder.com>

Add WeChat powcoder

Convolutional Layer

To preserve spatial structure, use an original 2D image




Assignment Project Exam Help

<https://powcoder.com>

5x5x3 filter

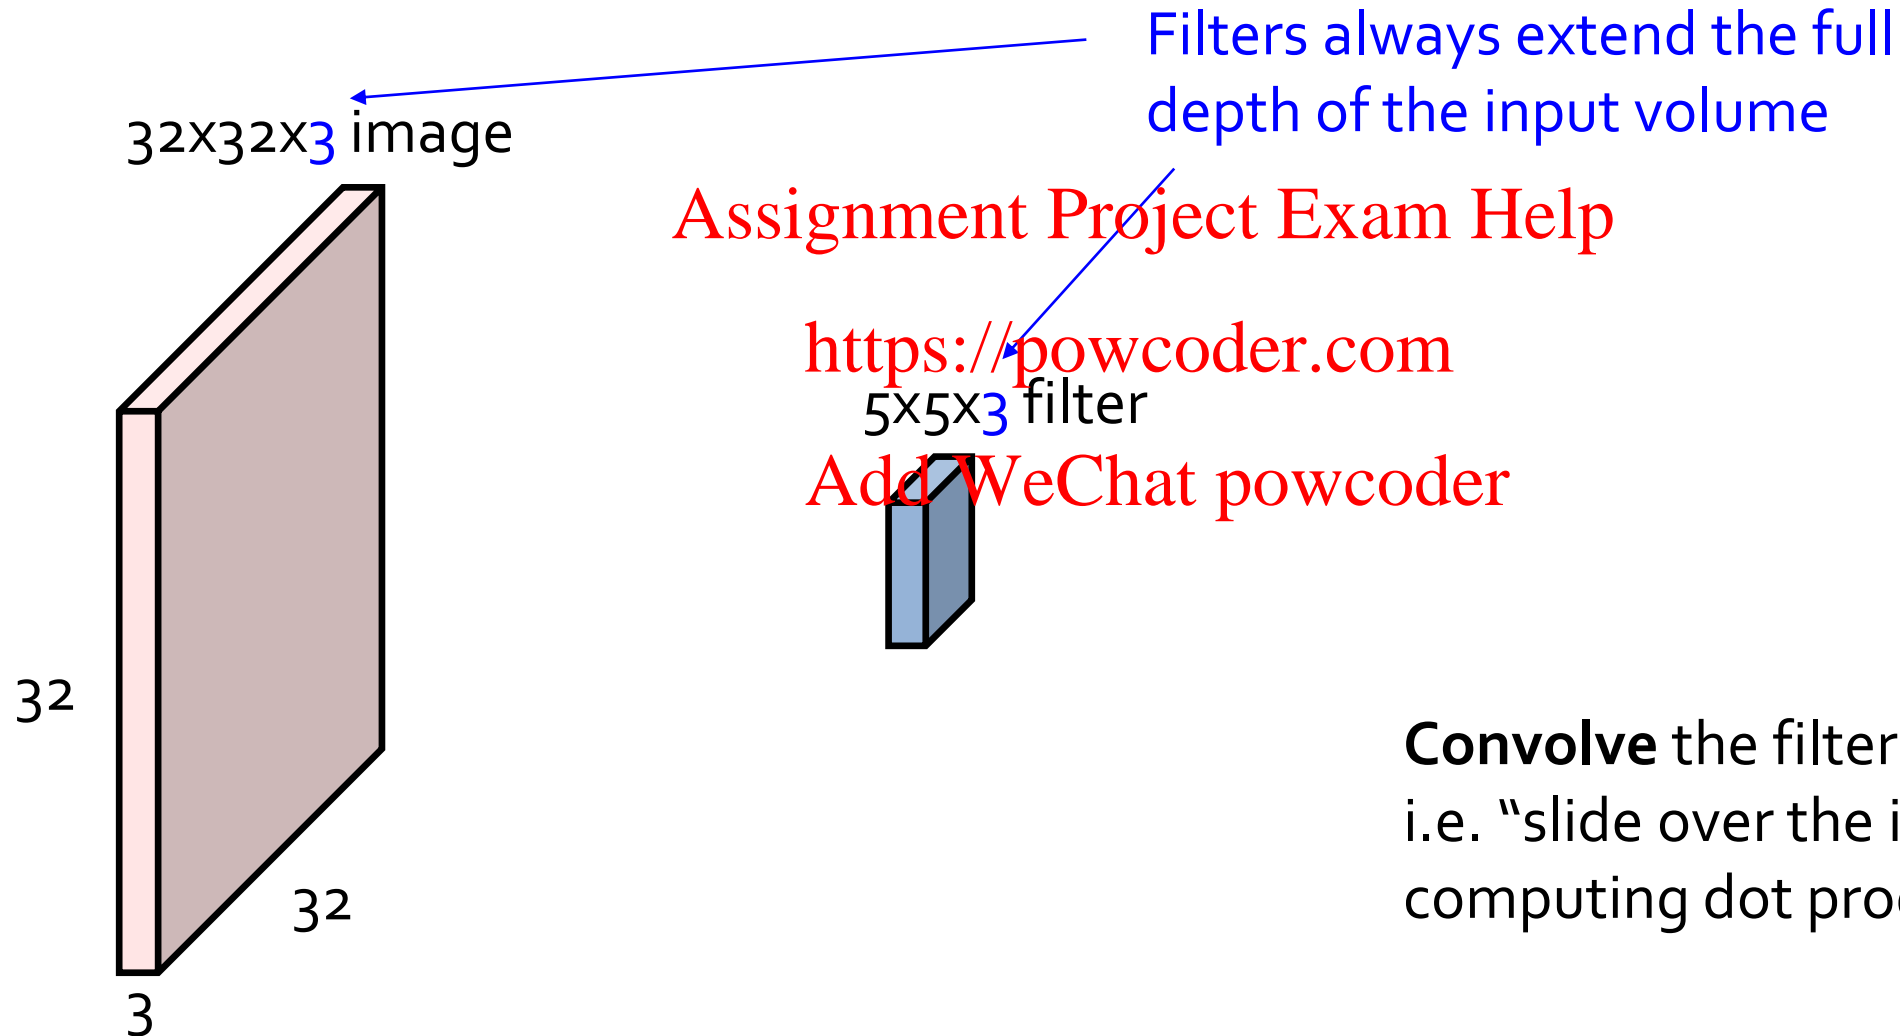
Add WeChat powcoder



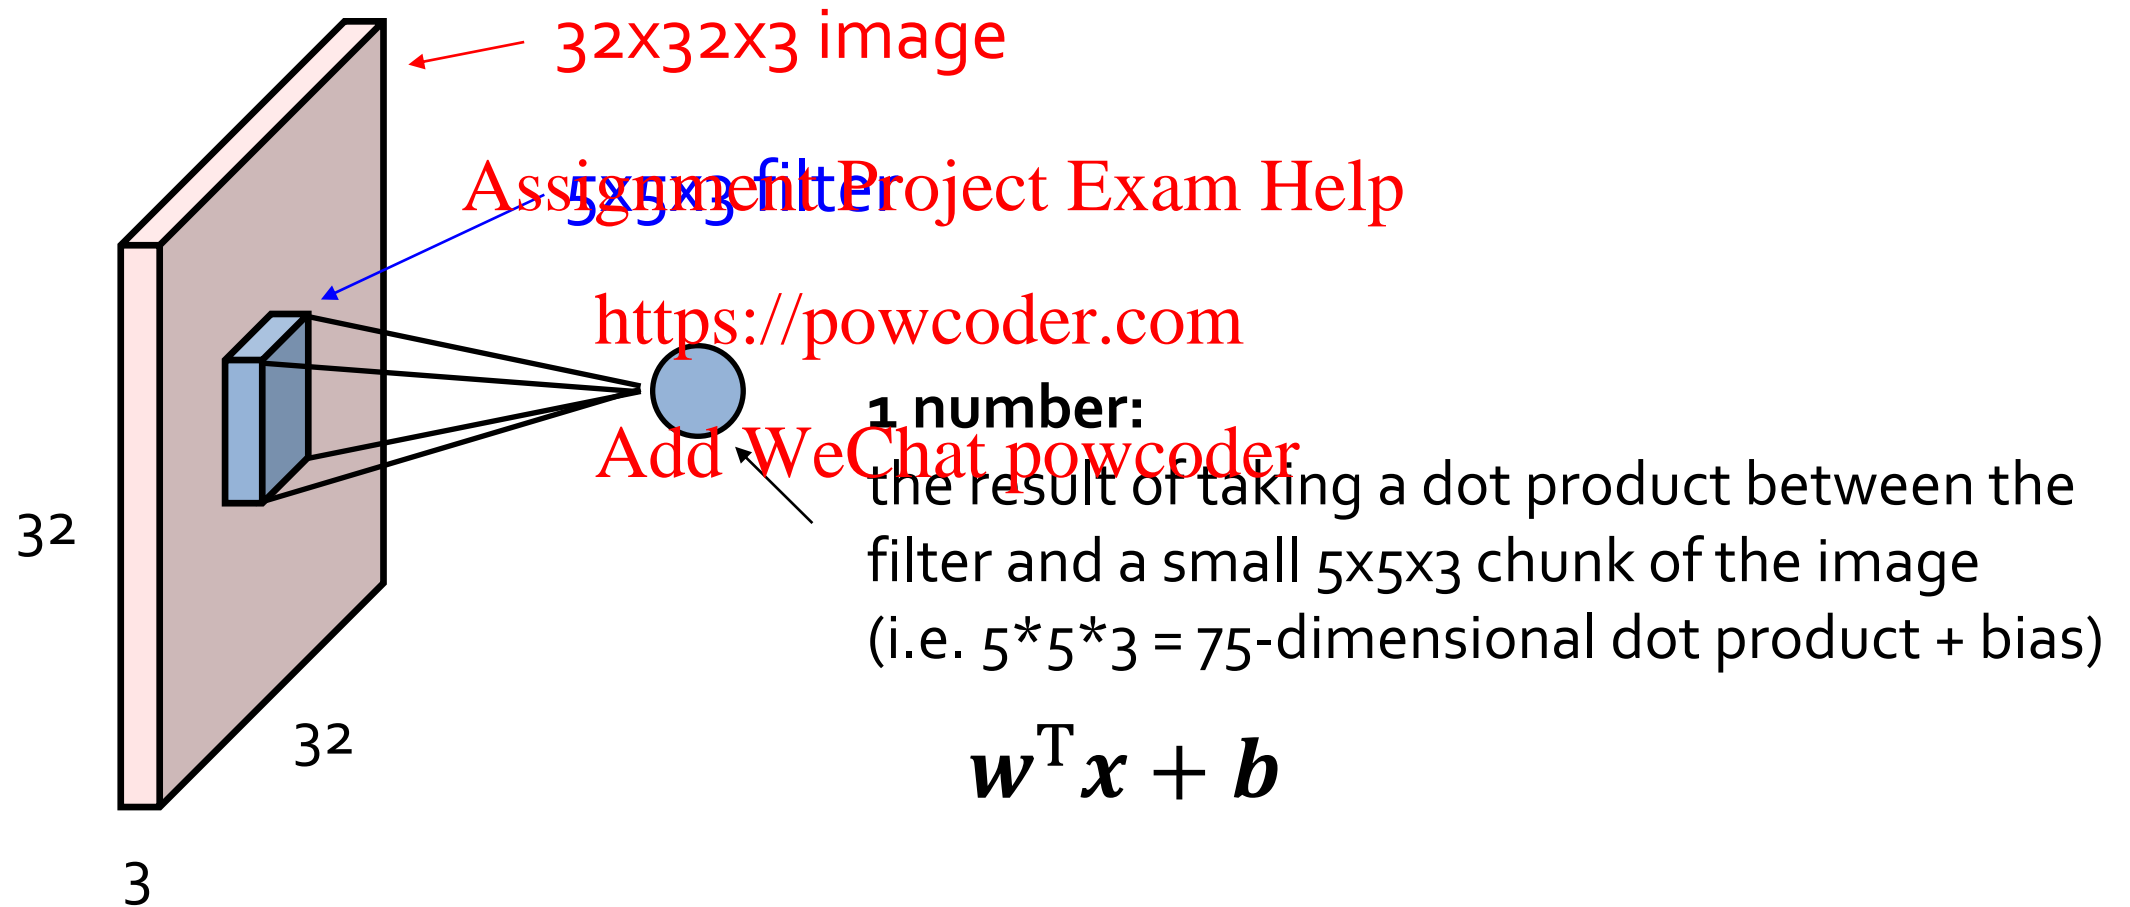
Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”

Convolutional Layer


To preserve spatial structure, use an original 2D image



Convolutional Layer




Convolutional Layer



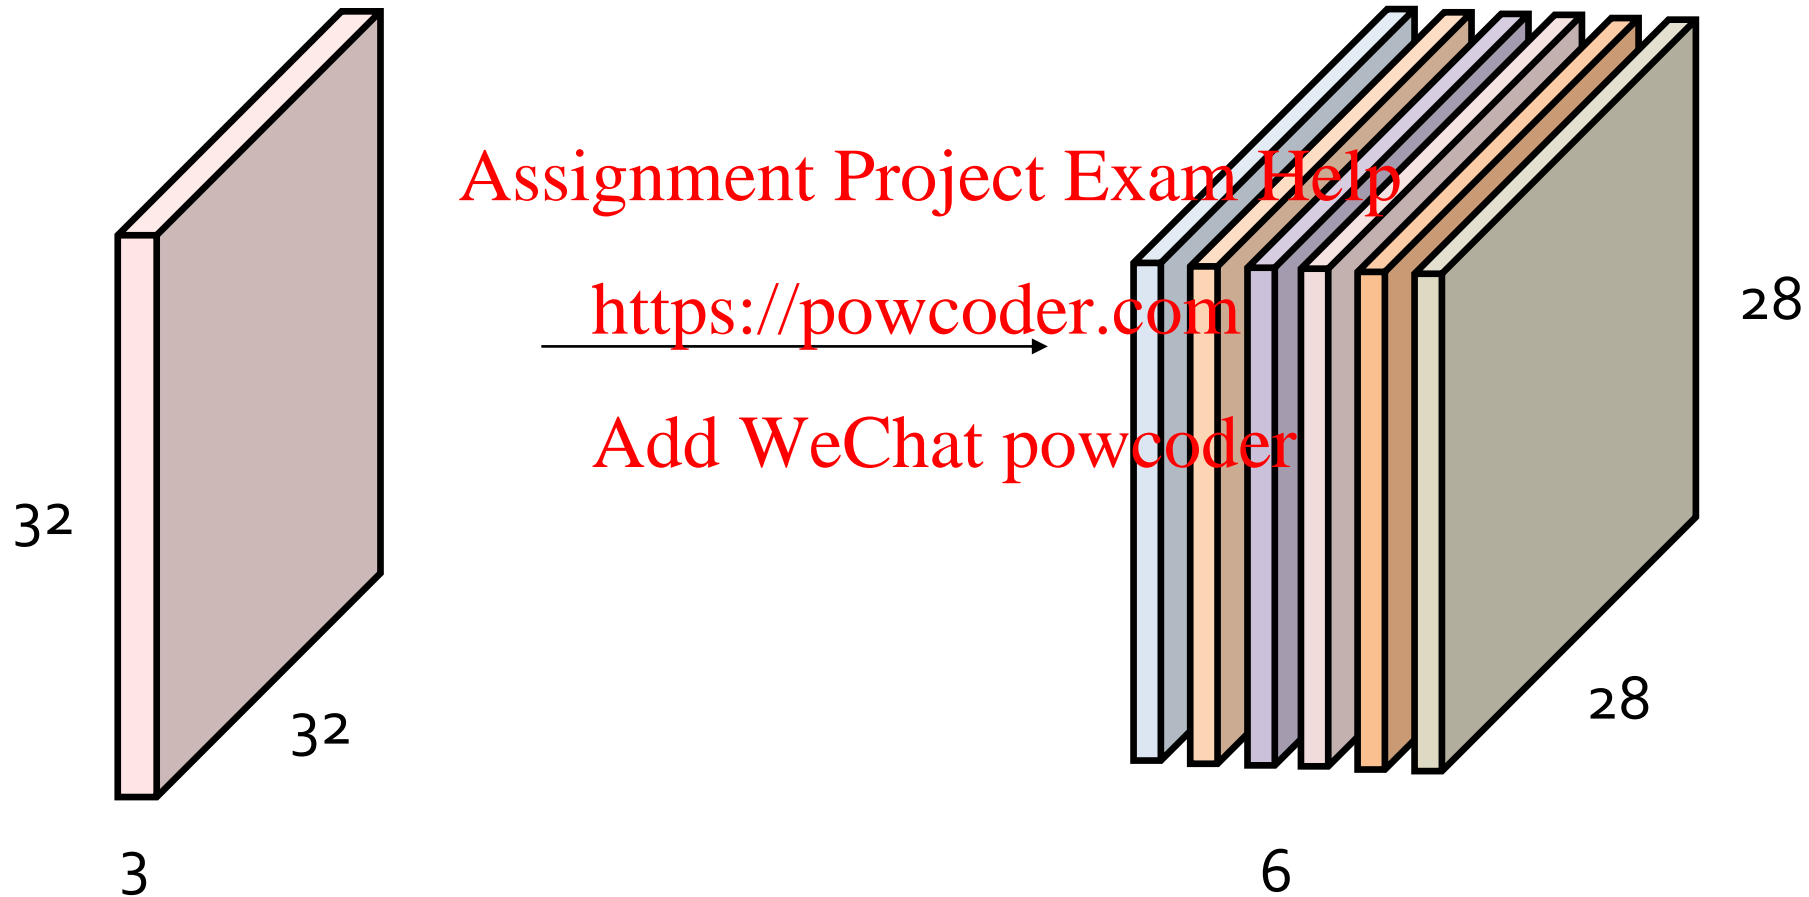
Convolutional Layer

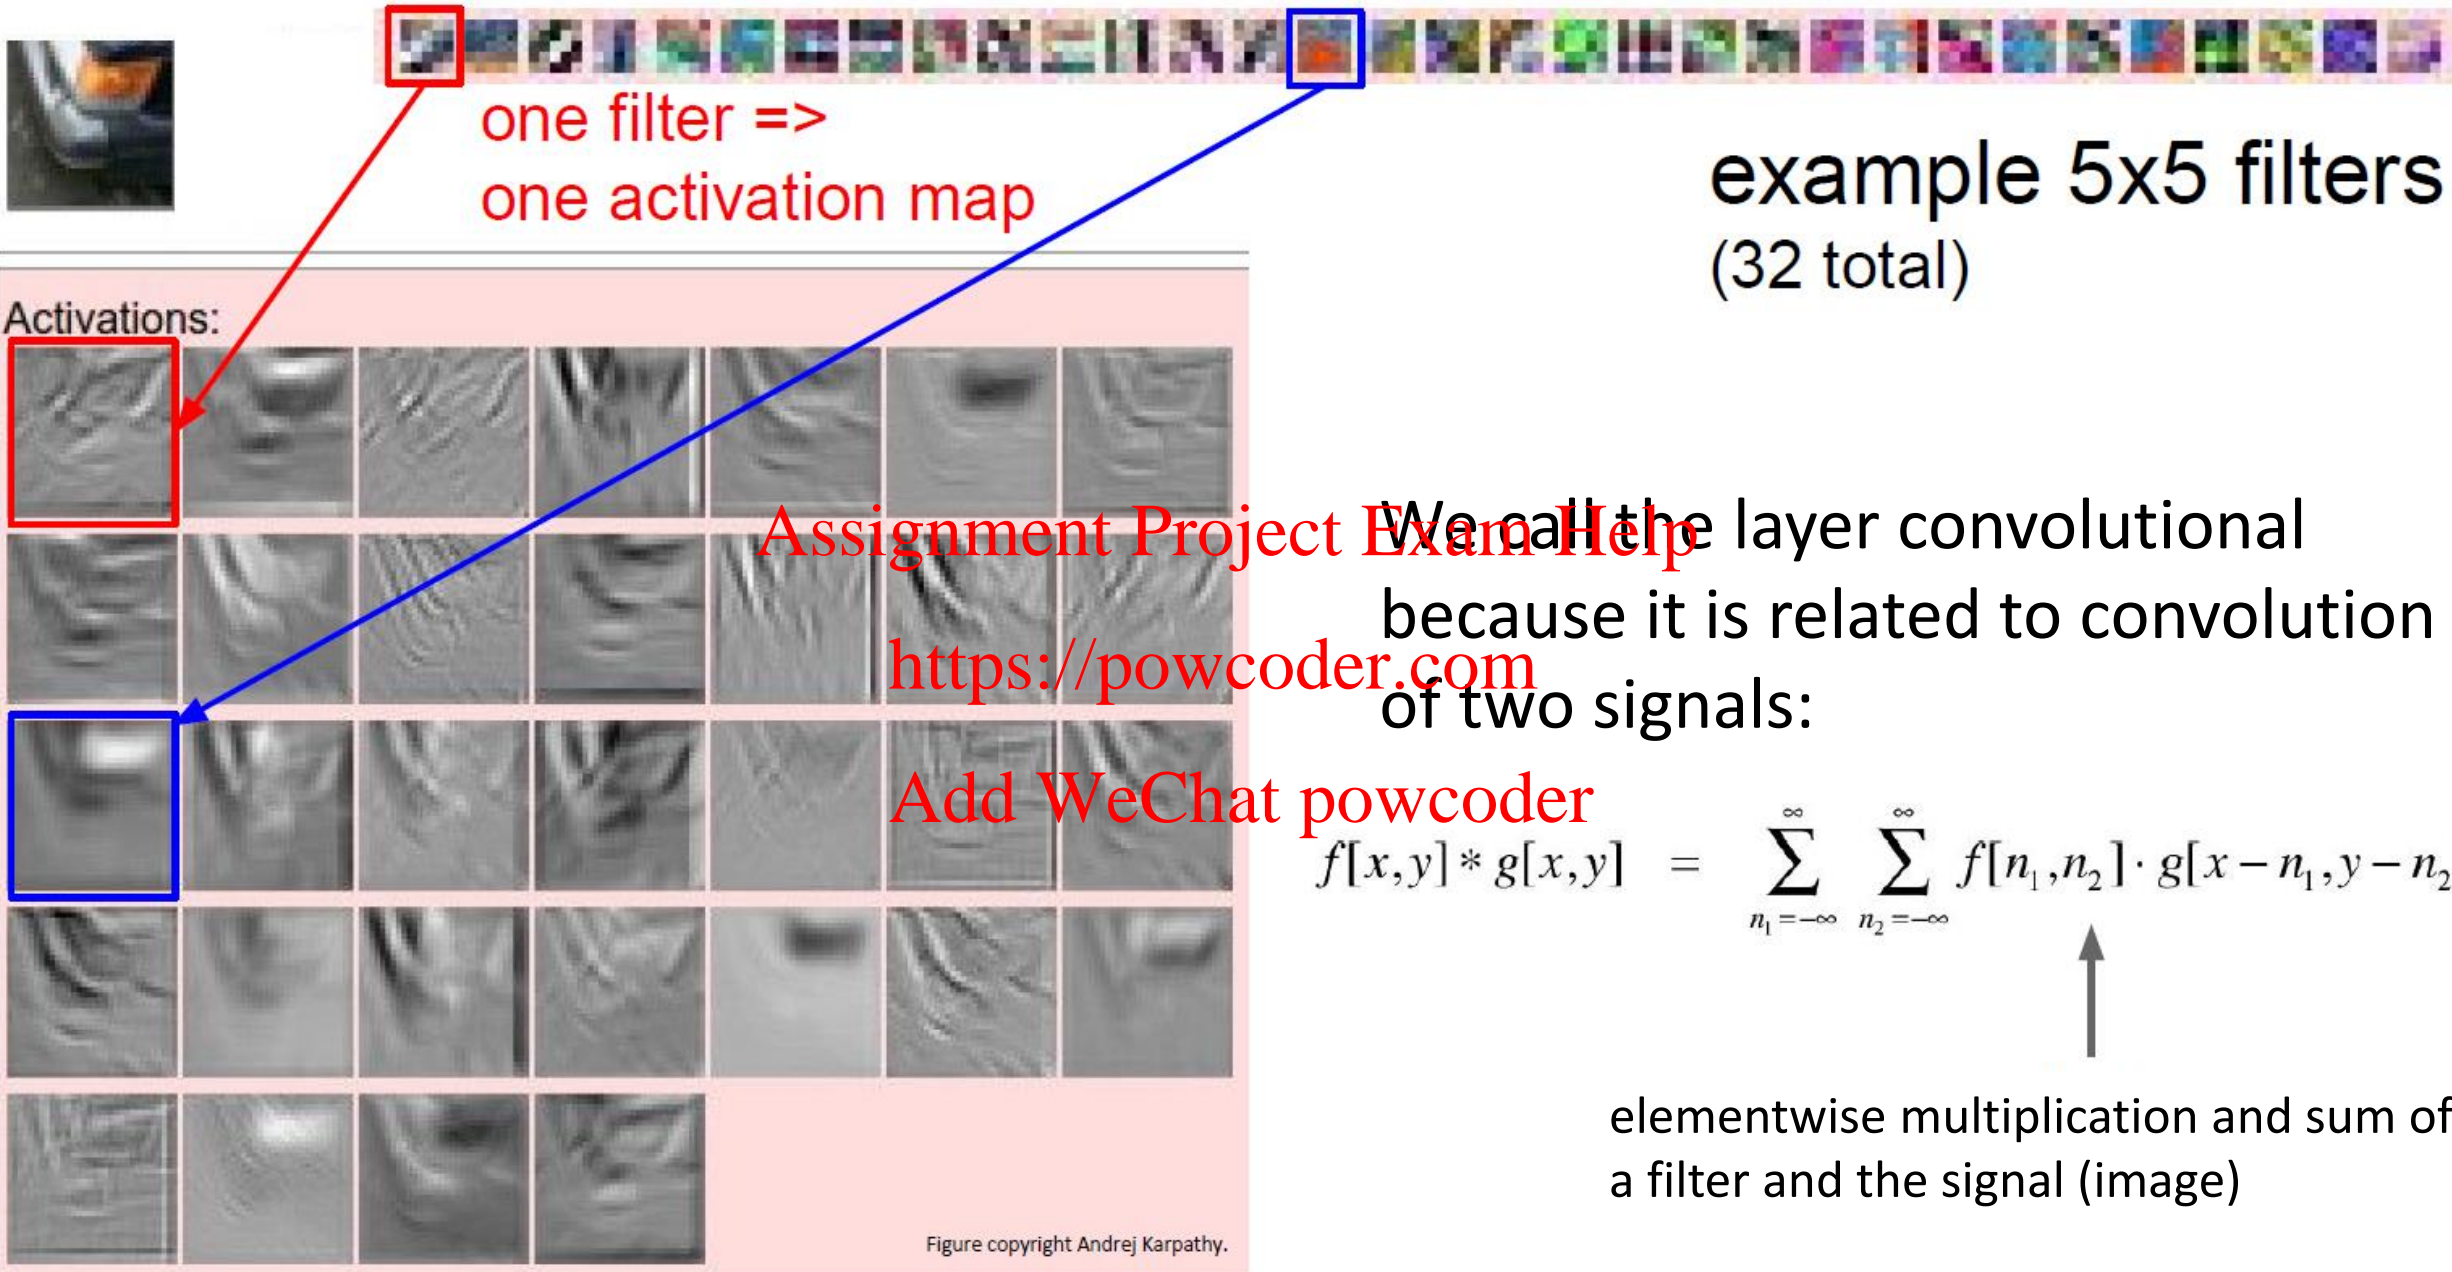
consider a second $5 \times 5 \times 3$ (orange) filter



Convolutional Layer

If we had **six** 5x5x3 filters, we'll get **six** separate activation maps:





one filter =>
one activation map

example 5x5 filters
(32 total)

We call the layer convolutional
because it is related to convolution
of two signals:

<https://powcoder.com>

Add WeChat powcoder

$$f[x,y] * g[x,y] = \sum_{n_1=-\infty}^{\infty} \sum_{n_2=-\infty}^{\infty} f[n_1,n_2] \cdot g[x-n_1,y-n_2]$$

↑
elementwise multiplication and sum of
a filter and the signal (image)

Figure copyright Andrej Karpathy.

Convolutional Layer

- The number of parameters in convolutional layer

Input	Weight	Output
$H_1 \times W_1 \times C_1$	C_2 filters of $F_h \times F_v \times C_1$	$H_1 \times W_1 \times C_2$

Assignment Project Exam Help



The number of weights: $C_2 \times (F_h \times F_v \times C_1)$

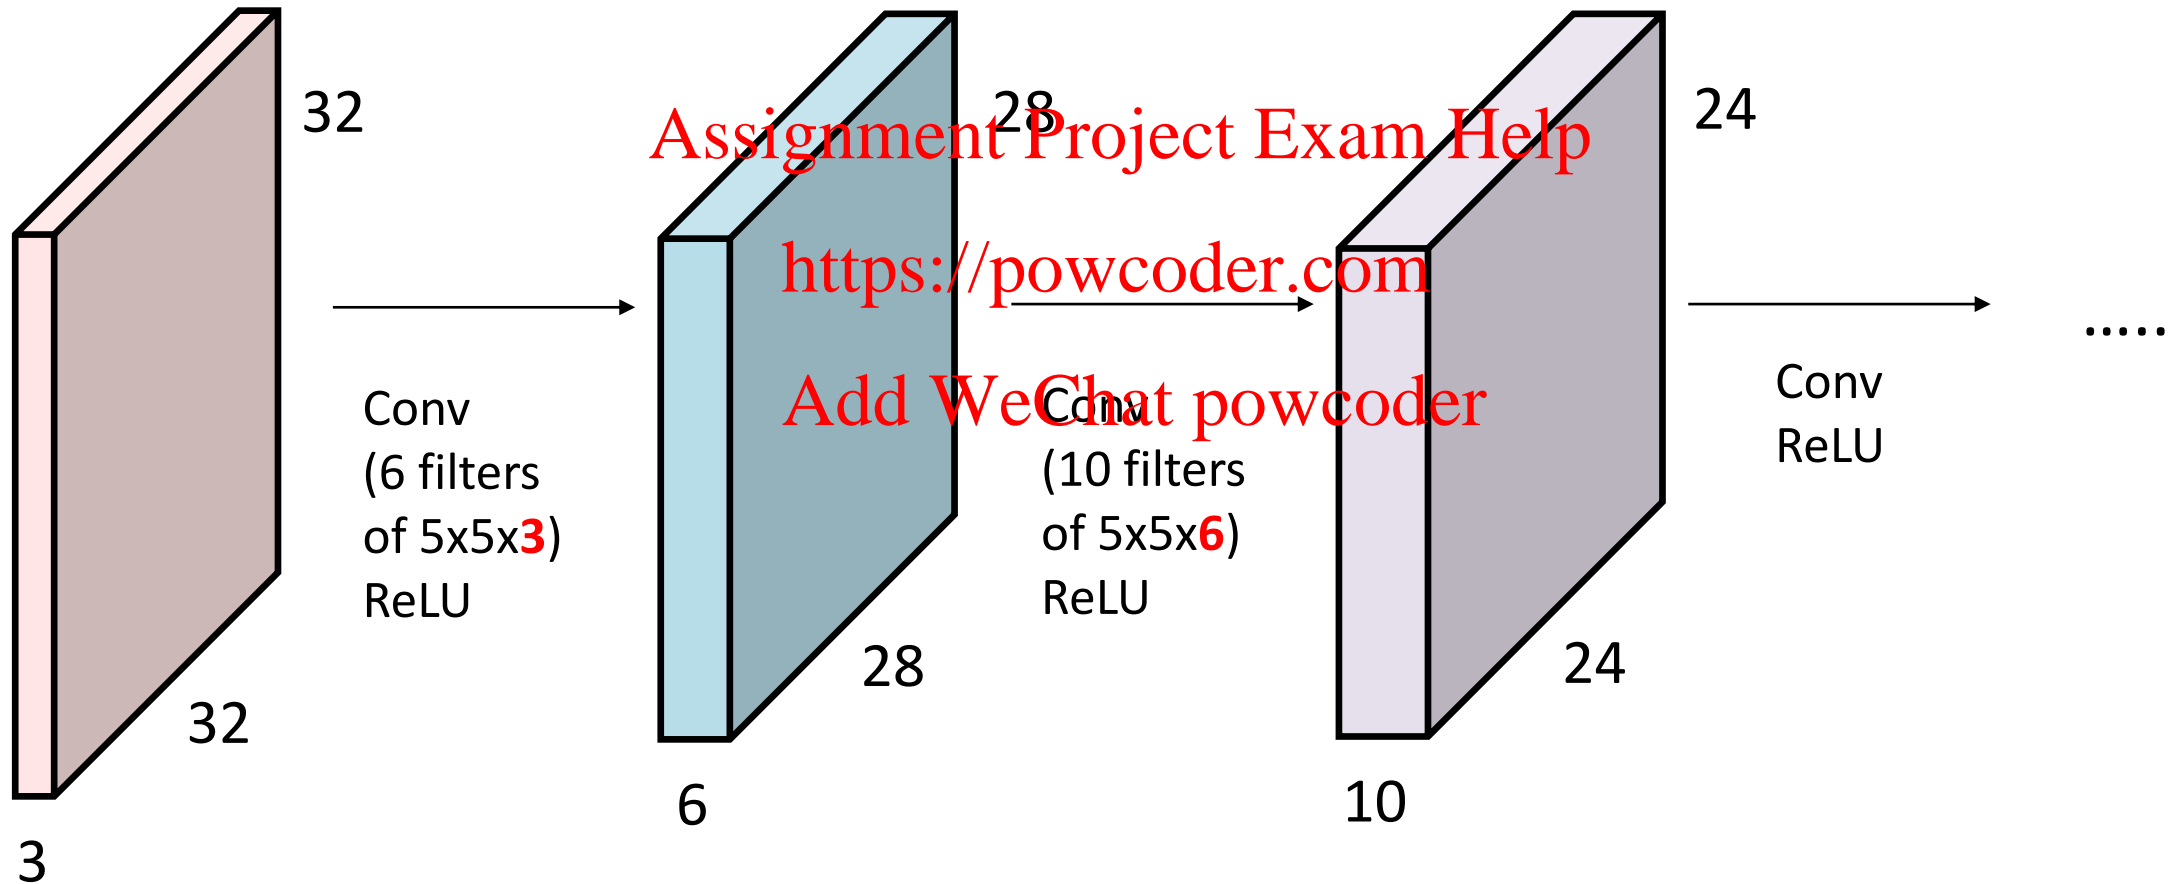
The number of bias: C_2

<https://powcoder.com>

Add WeChat powcoder

ConvNet

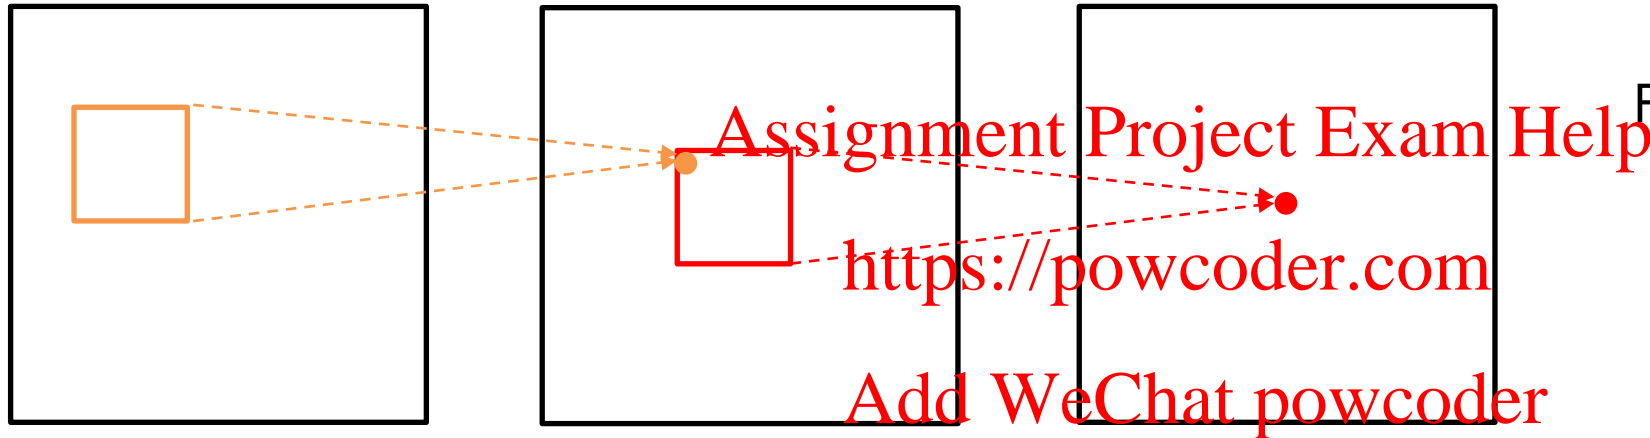
- A sequence of Convolutional Layers, interspersed with activation functions



ConvNet

- **Receptive field**

- The region of the input space that affects a particular unit of the network



From the convolution property,

$$y = x * h_1$$

$$z = y * h_2$$

$$\rightarrow z = x * h_1 * h_2 = x * h$$

$$h = h_1 * h_2$$

$5 \times 5 \times 1$ filter

$5 \times 5 \times 1$ filter


Effective receptive field at : $9 \times 9 (= 5 + 5 - 1)$

Convolutional Filter Size

Three 3x3 Conv layers

vs.

Single 7x7 Conv layer




Assume that zero-padding is applied to preserve a spatial resolution

- Receptive fields are equal. For three 3x3 Conv layers, $3+3-1+3-1 = 7$.
- # of Conv parameters: $3 \times C \times (3 \times 3 \times C) = 27C^2$ vs. $C \times (7 \times 7 \times C) = 49C^2$
- The three stacks of CONV layers produce more *expressive* activation maps

Spatial Dimension at Convolution Layer

7



7x7 input (spatially)

assume 3x3 filter applied **with stride 1**
-> 5x5 output

Assignment Project Exam Help

<https://powcoder.com>

7 assume 3x3 filter applied **with stride 2**
-> 3x3 output


Add WeChat powcoder

assume 3x3 filter applied **with stride 3**
-> **doesn't fit!**

cannot apply 3x3 filter on 7x7 input with stride 3.

***Stride**: is the number of pixels shifts over the input matrix

Spatial Dimension at Convolution Layer



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Output size:

$$(N - F) / \text{stride} + 1$$

e.g. $N = 7, F = 3$:

$$\text{stride } 1 \Rightarrow (7 - 3) / 1 + 1 = 5$$

$$\text{stride } 2 \Rightarrow (7 - 3) / 2 + 1 = 3$$

$$\text{stride } 3 \Rightarrow (7 - 3) / 3 + 1 = 2.33$$

Spatial Dimension at Convolution Layer

0	0	0	0	0	0			
0								
0								
0								
0								

In practice: Common to zero pad the border

e.g. input 7x7

3x3 filter, applied with **stride 1**

pad with **1 pixel** border

<https://powcoder.com>
-> **7x7 output!**

Add WeChat powcoder

in general, common to see CONV layers with
stride 1, filters of size $F \times F$, and zero-padding with
 $(F-1)/2$. (will preserve size spatially)

e.g. $F = 3 \Rightarrow$ zero pad with 1

$F = 5 \Rightarrow$ zero pad with 2

$F = 7 \Rightarrow$ zero pad with 3

Spatial Dimension at Convolution Layer

Input volume: $32 \times 32 \times 3$

$10 \times 5 \times 5 \times 3$ filters with stride 1, pad 2


<https://powcoder.com>

Output volume size:

$(32 + 2 * 2 - 5) / 1 + 1 = 32$ spatially, so

$32 \times 32 \times 10$

Add WeChat powcoder




Spatial Dimension at Convolution Layer

Input volume: $32 \times 32 \times 3$
 $10 \ 5 \times 5 \times 3$ filters with stride 1, pad 2

<https://powcoder.com>

Number of parameters in this layer?
each filter has $5 * 5 * 3 + 1 = 76$ params (+1 for bias)

-> $76 * 10 = 760$



Spatial Dimension at Convolution Layer

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
 - Number of filters K ,
 - their spatial extent F ,
 - the stride S ,
 - the amount of zero padding P
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F + 2P)/S + 1$
 - $H_2 = (H_1 - F + 2P)/S + 1$ (i.e. width and height are computed equally by symmetry)
 - $D_2 = K$
- With parameter sharing, it introduces $F \cdot F \cdot D_1$ weights per filter, for a total of $(F \cdot F \cdot D_1) \cdot K$ weights and K biases.
- In the output volume, the d -th depth slice (of size $W_2 \times H_2$) is the result of performing a valid convolution of the d -th filter over the input volume with a stride of S , and then offset by d -th bias.

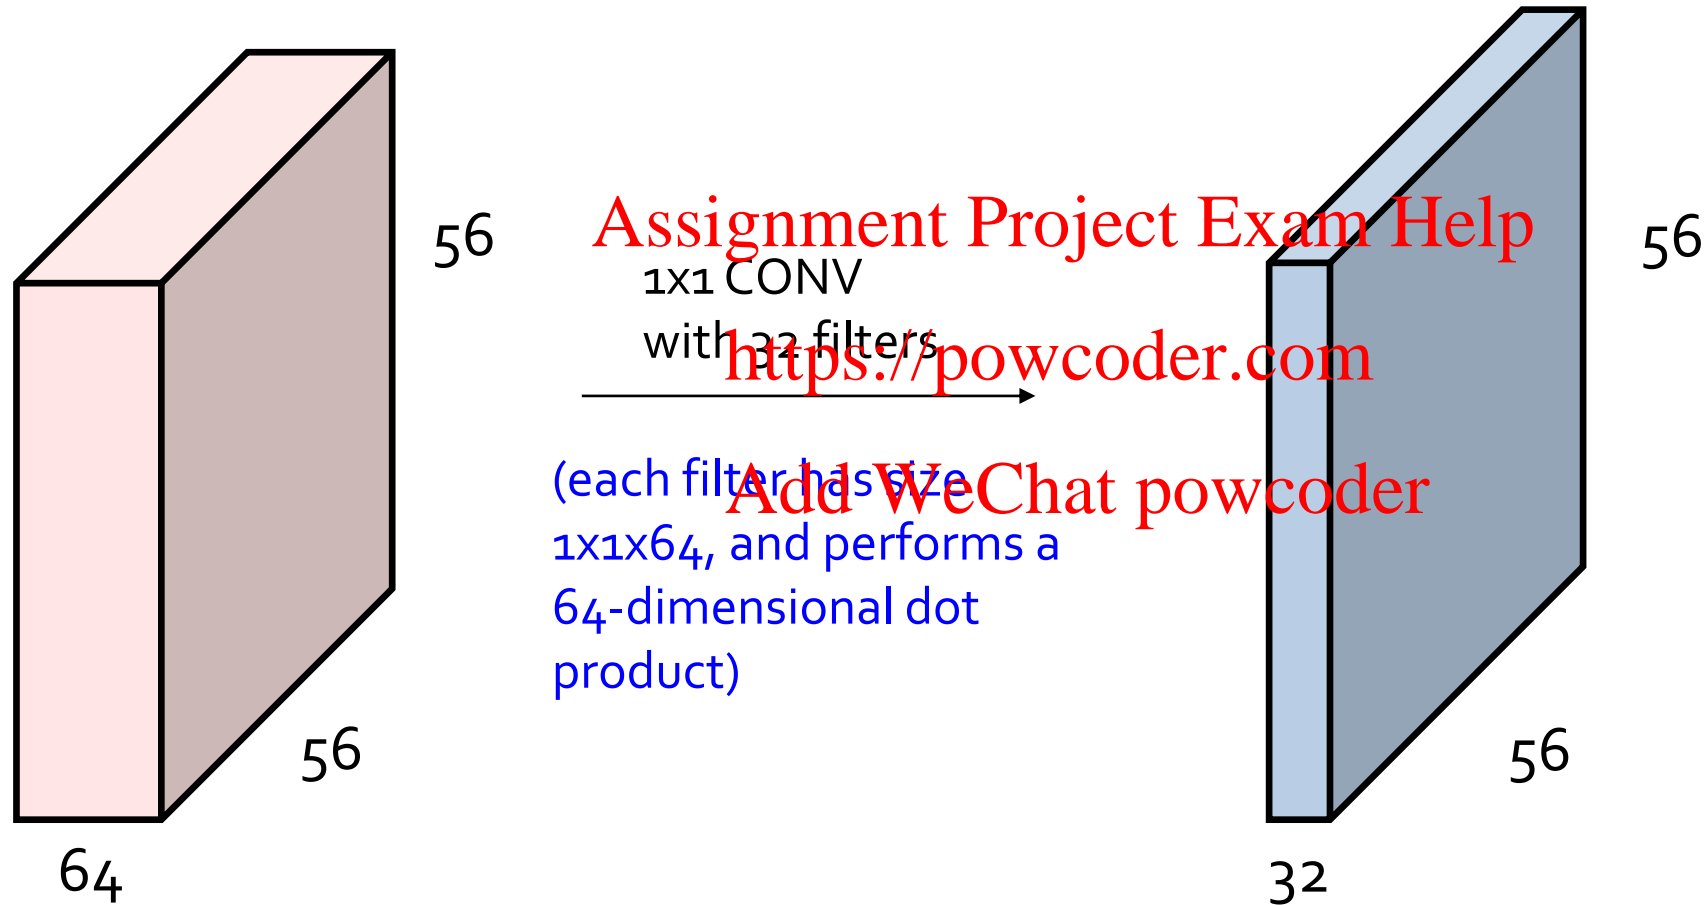
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

1x1 Convolution

1x1 convolution layers make perfect sense



Assignment Project Exam Help

Convolutional Layer

<https://powcoder.com>

Implementation and Backpropagation

Add WeChat powcoder

Implementation as Matrix Multiplication

- Convolution: dot products between the filters and local regions of the input
- Conv layer: the forward pass of a convolutional layer as one big matrix multiply

- Example of feed-forward process

1. Convert the input into X_{col} by taking a block of $11 \times 11 \times 3 (=363)$ pixels in the input for $55 \times 55 (=3025)$ times

X_{col} : $[363 \times 3025]$

2. Reshape the conv filter into W_{row} : $[96 \times 363]$

Reshape the conv bias (96x1 vector) into b_{col} : $[96 \times 3025]$ by stacking it for 3025 times

3. Perform matrix multiplication $O = W_{row} * X_{col} + b_{col}$

4. Reshape O : $[96 \times 3025]$ into $[55 \times 55 \times 96]$

Input: $[227 \times 227 \times 3]$
Conv filter: 96 filters of $[11 \times 11 \times 3]$
Conv bias: 96x1 vector
Stride: 4
Padding: 0

Output: $(227-11)/4+1 = 55$
-> $[55 \times 55 \times 96]$

$$O = W_{row} * X_{col} + b_{col}$$

96×3025

96×363

363×3025

96×3025

Backpropagation of Convolution Layer

Input: [227x227x3]
 Conv filter: 96 filters of [11x11x3]
 Stride: 4
 Padding: 0

 Output: (227-11)/4+1 = 55
 -> [55x55x96]

Convolution layer shares \mathbf{W} for all neurons of current activation map.
 For each neuron,

$$\frac{\partial L}{\partial \mathbf{x}_p} = \mathbf{W}^T \frac{\partial L}{\partial \mathbf{s}_p}$$

$$\frac{\partial L}{\partial \mathbf{W}} = \sum_p \frac{\partial L}{\partial \mathbf{s}_p} \mathbf{x}_p^T$$

$$\frac{\partial L}{\partial \mathbf{b}} = \sum_p \frac{\partial L}{\partial \mathbf{s}_p}$$


$p = 1, \dots, 3025 (= 55 \times 55)$

Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

Backpropagation of $\mathbf{s} = \mathbf{W}\mathbf{x} + \mathbf{b}$



$$\frac{\partial L}{\partial \mathbf{x}} = \frac{\partial \mathbf{s}}{\partial \mathbf{x}} \frac{\partial L}{\partial \mathbf{s}} = \mathbf{W}^T \frac{\partial L}{\partial \mathbf{s}}$$

$$\frac{\partial L}{\partial \mathbf{W}} = \frac{\partial L}{\partial \mathbf{s}} \mathbf{x}^T \quad \frac{\partial L}{\partial \mathbf{b}} = \frac{\partial L}{\partial \mathbf{s}}$$



Backpropagation of Convolution Layer

Input: [227x227x3]
 Conv filter: 96 filters of [11x11x3]
 Stride: 4
 Padding: 0

 Output: (227-11)/4+1 = 55
 -> [55x55x96]

Convolution layer shares \mathbf{W} for all neurons of current activation map.
 For all neurons,

$$\frac{\partial L}{\partial \mathbf{x}} = \mathbf{W}^T \frac{\partial L}{\partial \mathbf{s}}$$

$$\frac{\partial L}{\partial \mathbf{W}} = \frac{\partial L}{\partial \mathbf{s}} \mathbf{x}^T$$


$$\frac{\partial L}{\partial \mathbf{b}} = \frac{\partial L}{\partial \mathbf{s}} \mathbf{1}^T \quad \text{1x3025 vector}$$

Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

Backpropagation of $\mathbf{s} = \mathbf{W}\mathbf{x} + \mathbf{b}$



$$\frac{\partial L}{\partial \mathbf{x}} = \frac{\partial \mathbf{s}}{\partial \mathbf{x}} \frac{\partial L}{\partial \mathbf{s}} = \mathbf{W}^T \frac{\partial L}{\partial \mathbf{s}}$$

$$\frac{\partial L}{\partial \mathbf{W}} = \frac{\partial L}{\partial \mathbf{s}} \mathbf{x}^T \quad \frac{\partial L}{\partial \mathbf{b}} = \frac{\partial L}{\partial \mathbf{s}}$$



Backpropagation of Convolution Layer

$$\frac{\partial L}{\partial \mathbf{x}} = \mathbf{W}^T \frac{\partial L}{\partial \mathbf{s}}$$

1. Perform $\frac{\partial L}{\partial \mathbf{x}} = \mathbf{W}^T \frac{\partial L}{\partial \mathbf{s}}$
2. Reshape $\frac{\partial L}{\partial \mathbf{x}}$ (363x3025) into 3025 gradients of 11x11x3
3. Overlay the reshaped gradient into 3D matrix [227x227x3] in which overlapped gradients are accumulated.

Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

$$\frac{\partial L}{\partial \mathbf{W}} = \frac{\partial L}{\partial \mathbf{s}} \mathbf{x}^T$$

$$\frac{\partial L}{\partial \mathbf{b}} = \frac{\partial L}{\partial \mathbf{s}} \mathbf{1}^T$$

1. Convert the input into \mathbf{x} by taking a block of 11x11x3 (=363) pixels in the input for 55x55 (=3025) times. \mathbf{x} : [363x3025]

2. Perform $\frac{\partial L}{\partial \mathbf{W}} = \frac{\partial L}{\partial \mathbf{s}} \mathbf{x}^T$ and $\frac{\partial L}{\partial \mathbf{b}} = \frac{\partial L}{\partial \mathbf{s}} \mathbf{1}^T$



Fully Connected Layer

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Fully Connected Layer

32x32x3 image -> stretch to 3072 x 1

Assignment Project Exam Help



<https://powcoder.com>

Add WeChat powcoder

Each neuron looks at the full input volume

1 number:

the result of taking a dot product between a row of W and the input (a 3072-dimensional dot product)

Assignment Project Exam Help

<https://powcoder.com>
Pooling Layer


Add WeChat powcoder

Pooling layer

- makes the representations smaller and more manageable
- operates over each activation map independently

Pooling layer

- Max pooling
- Average pooling (rarely used)
- L2 norm pooling (rarely used)



MAX POOLING

Single depth slice

x ↑

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

→ y

max pool with 2x2 filters
and stride 2

<https://powcoder.com>

Add WeChat powcoder

6	8
3	4

EBU7240

Computer Vision

Assignment Project Exam Help

- <https://powcoder.com> -

Add WeChat powcoder

Semester 1, 2021

Changjae Oh

CNN Architectures

- **Case Studies**

- AlexNet
- VGG
- GoogLeNet
- ResNet

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Case Study: AlexNet

[Krizhevsky et al. 2012]

Architecture:

CONV1

MAX POOL1

NORM1

CONV2

MAX POOL2

NORM2

CONV3

CONV4


CONV5

Max POOL3

FC6

FC7

FC8




Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Case Study: AlexNet

[Krizhevsky et al. 2012]



Assignment Project Exam Help

Input: 227x227x3 images

First layer (CONV1): 96 11x11 filters applied at stride 4

=>

The output volume size: $(227-11)/4+1 = 55$

Output volume [55x55x96]

Total number of parameters in this layer

Parameters: $(11*11*3)*96 = 35K$


Bias: 96

<https://powcoder.com>

Add WeChat powcoder

Case Study: AlexNet

[Krizhevsky et al. 2012]



Assignment Project Exam Help

Input: 227x227x3 images
After CONV1: 55x55x96

<https://powcoder.com>

Add WeChat powcoder

Second layer (POOL1): 3x3 filters applied at stride 2

The output volume size: $(55-3)/2+1 = 27$


Output volume: 27x27x96

The number of parameters in this layer

Parameters: 0!

Case Study: AlexNet

[Krizhevsky et al. 2012]



Assignment Project Exam Help

Input: 227x227x3 images

After CONV1: 55x55x96

After POOL1: 27x27x96


...

<https://powcoder.com>

Add WeChat powcoder

Case Study: AlexNet

[Krizhevsky et al. 2012]



Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1

[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons

[4096] **FC7**: 4096 neurons

[1000] **FC8**: 1000 neurons (class scores)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Details/Retrospectives:

- first use of ReLU

- used Norm layers (not common anymore)

- heavy data augmentation

- dropout 0.5

- batch size 128


- SGD Momentum 0.9

- Learning rate 1e-2, reduced by 10 manually when val accuracy plateaus

- L2 weight decay 5e-4

- 7 CNN ensemble: 18.2% -> 15.4%

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



Case Study: VGGNet

[Simonyan and Zisserman, 2014]

Small filters, Deeper networks

Only 3x3 CONV stride 1, pad 1, and 2x2 MAX POOL stride 2

8 layers (AlexNet)

-> 16 - 19 layers (VGGNet)

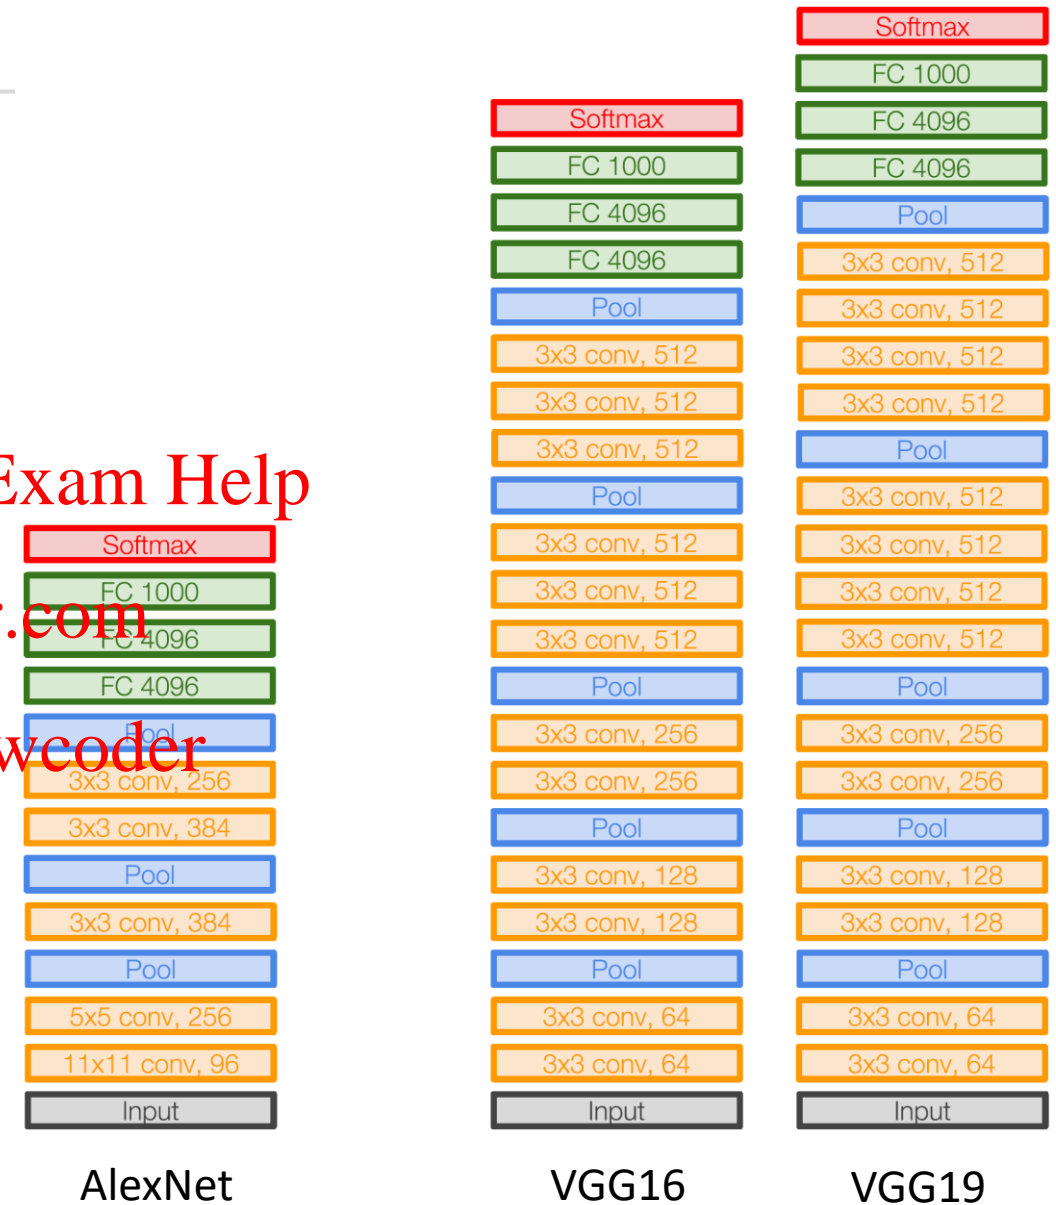
11.7% top 5 error in ILSVRC'13 (ZFNet)

-> 7.3% top 5 error in ILSVRC'14

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Case Study: VGGNet

[Simonyan and Zisserman, 2014]

Q: Why use smaller filters? (3x3 conv)

Stack of three 3x3 conv (stride 1) layers
has same effective receptive field as
one 7x7 conv layer

But deeper, more non-linearity

And fewer parameters: $3 * (32C^2)$ vs.
 $72C^2$ for C channels per layer


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



AlexNet




VGG16



VGG19

(not counting biases)

INPUT: [224x224x3] **memory:** 224*224*3=150K **params:** 0
 CONV3-64: [224x224x64] **memory:** 224*224*64=3.2M **params:** (3*3*3)*64 = 1,728
 CONV3-64: [224x224x64] **memory:** 224*224*64=3.2M **params:** (3*3*64)*64 = 36,864
 POOL2: [112x112x64] **memory:** 112*112*64=800K **params:** 0
 CONV3-128: [112x112x128] **memory:** 112*112*128=1.6M **params:** (3*3*64)*128 = 73,728
 CONV3-128: [112x112x128] **memory:** 112*112*128=1.6M **params:** (3*3*128)*128 = 147,456
 POOL2: [56x56x128] **memory:** 56*56*128=400K **params:** 0
 CONV3-256: [56x56x256] **memory:** 56*56*256=800K **params:** (3*3*128)*256 = 294,912
 CONV3-256: [56x56x256] **memory:** 56*56*256=800K **params:** (3*3*256)*256 = 589,824
 CONV3-256: [56x56x256] **memory:** 56*56*256=800K **params:** (3*3*256)*256 = 589,824
 POOL2: [28x28x256] **memory:** 28*28*256=200K **params:** 0
 CONV3-512: [28x28x512] **memory:** 28*28*512=400K **params:** (3*3*256)*512 = 1,179,648
 CONV3-512: [28x28x512] **memory:** 28*28*512=400K **params:** (3*3*512)*512 = 2,359,296
 CONV3-512: [28x28x512] **memory:** 28*28*512=400K **params:** (3*3*512)*512 = 2,359,296
 POOL2: [14x14x512] **memory:** 14*14*512=100K **params:** 0
 CONV3-512: [14x14x512] **memory:** 14*14*512=100K **params:** (3*3*512)*512 = 2,359,296
 CONV3-512: [14x14x512] **memory:** 14*14*512=100K **params:** (3*3*512)*512 = 2,359,296
 CONV3-512: [14x14x512] **memory:** 14*14*512=100K **params:** (3*3*512)*512 = 2,359,296
 POOL2: [7x7x512] **memory:** 7*7*512=25K **params:** 0
 FC: [1x1x4096] **memory:** 4096 **params:** 7*7*512*4096 = 102,760,448
 FC: [1x1x4096] **memory:** 4096 **params:** 4096*4096 = 16,777,216
 FC: [1x1x1000] **memory:** 1000 **params:** 4096*1000 = 4,096,000



VGG16

TOTAL memory: 15.2M * 4 bytes \approx 61MB / image (for a forward pass)

TOTAL params: 138M parameters

(not counting biases)

INPUT: [224x224x3] **memory:** $224*224*3=150\text{K}$ **params:** 0

CONV3-64: [224x224x64] **memory:** $224*224*64=3.2\text{M}$ **params:** $(3*3*3)*64 = 1,728$

CONV3-64: [224x224x64] **memory:** $224*224*64=3.2\text{M}$ **params:** $(3*3*64)*64 = 36,864$

POOL2: [112x112x64] **memory:** $112*112*64=800\text{K}$ **params:** 0

CONV3-128: [112x112x128] **memory:** $112*112*128=1.6\text{M}$ **params:** $(3*3*64)*128 = 73,728$

CONV3-128: [112x112x128] **memory:** $112*112*128=1.6\text{M}$ **params:** $(3*3*128)*128 = 147,456$

POOL2: [56x56x128] **memory:** $56*56*128=400\text{K}$ **params:** 0

CONV3-256: [56x56x256] **memory:** $56*56*256=800\text{K}$ **params:** $(3*3*128)*256 = 294,912$

CONV3-256: [56x56x256] **memory:** $56*56*256=800\text{K}$ **params:** $(3*3*256)*256 = 589,824$

CONV3-256: [56x56x256] **memory:** $56*56*256=800\text{K}$ **params:** $(3*3*256)*256 = 589,824$

POOL2: [28x28x256] **memory:** $28*28*256=200\text{K}$ **params:** 0

CONV3-512: [28x28x512] **memory:** $28*28*512=400\text{K}$ **params:** $(3*3*256)*512 = 1,179,648$

CONV3-512: [28x28x512] **memory:** $28*28*512=400\text{K}$ **params:** $(3*3*512)*512 = 2,359,296$

CONV3-512: [28x28x512] **memory:** $28*28*512=400\text{K}$ **params:** $(3*3*512)*512 = 2,359,296$

POOL2: [14x14x512] **memory:** $14*14*512=100\text{K}$ **params:** 0

CONV3-512: [14x14x512] **memory:** $14*14*512=100\text{K}$ **params:** $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] **memory:** $14*14*512=100\text{K}$ **params:** $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] **memory:** $14*14*512=100\text{K}$ **params:** $(3*3*512)*512 = 2,359,296$

POOL2: [7x7x512] **memory:** $7*7*512=25\text{K}$ **params:** 0

FC: [1x1x4096] **memory:** 4096 **params:** $7*7*512*4096 = 102,760,448$

FC: [1x1x4096] **memory:** 4096 **params:** $4096*4096 = 16,777,216$

FC: [1x1x1000] **memory:** 1000 **params:** $4096*1000 = 4,096,000$

Most memory is in
early CONV

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Most params are
in late FC

TOTAL memory: $15.2\text{M} * 4 \text{ bytes} \approx 61\text{MB}$ / image (for a forward pass)

TOTAL params: 138M parameters

Case Study: VGGNet


[Simonyan and Zisserman, 2014]

Details:

- ILSVRC'14 2nd in classification, 1st in localization
- Similar training procedure as Krizhevsky, 2012
- No Local Response Normalisation (LRN)
- Use VGG16 or VGG19 (VGG19 only slightly better, more memory)
- Use ensembles for best results
- FC7 features generalize well to other tasks



AlexNet



VGG16



VGG19

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Case Study: GoogLeNet

[Szegedy et al., 2014]


Deeper and wider networks,
with computational efficiency

- 22 layers
- Efficient “Inception” module
- No FC layers
- Only 5 million parameters!
12x less than AlexNet
- ILSVRC’14 classification winner
(6.7% top 5 error)


Assignment Project Exam Help

<https://powecoder.com>

Add WeChat powecoder



Inception module



Case Study: GoogLeNet


[Szegedy et al., 2014]

“Inception module”: design a good local network topology (network within a network) and then stack these modules on top of each other


Assignment Project Exam Help

<https://powecoder.com>

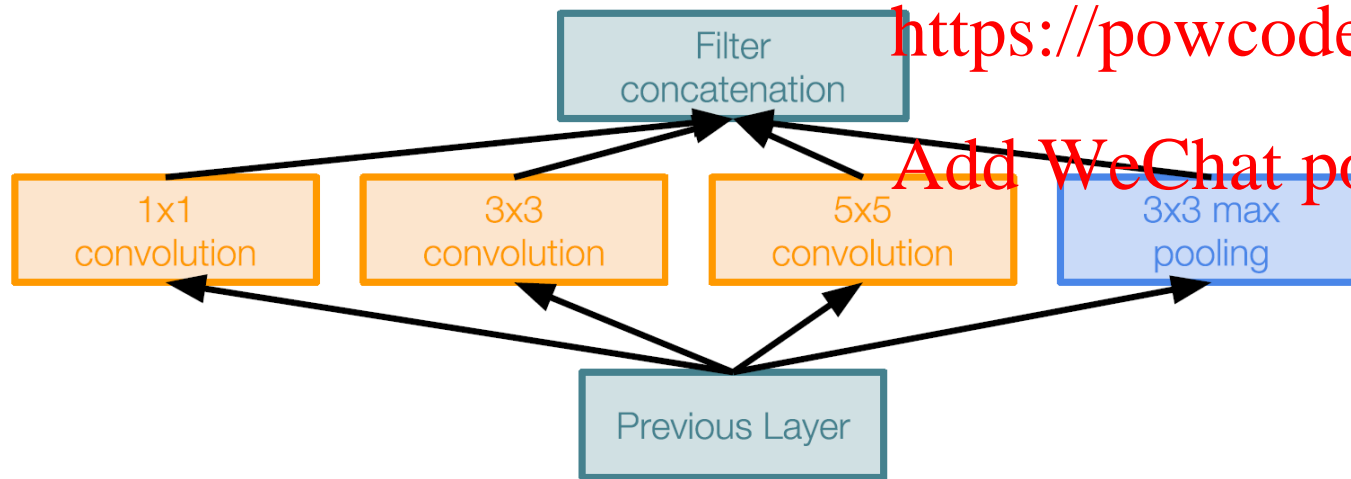
Add WeChat powecoder



Inception module



Case Study: GoogLeNet



Naive Inception module

Apply parallel filter operations on the input from previous layer:

- Multiple receptive field sizes for convolution (1x1, 3x3, 5x5)
- Pooling operation (3x3)

Concatenate all filter outputs together depth-wise

Q: What is the problem with this?
[Hint: Computational complexity]

Case Study: GoogLeNet

Example:

Q: What are the output size of filter concatenation?

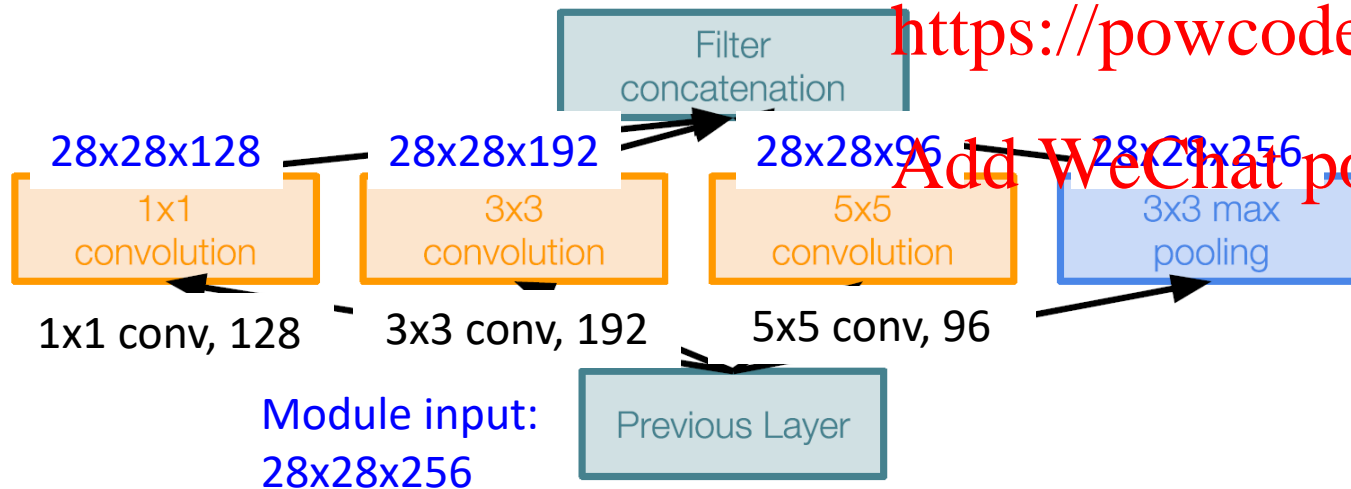
(Assume that a proper size of zero-padding is used)

Q: What is the problem with this?
[Hint: Computational complexity]

Assignment Project Exam Help
 $28 \times 28 \times (128 + 192 + 96 + 256) = 28 \times 28 \times 672$

<https://powcoder.com>

Add WeChat powcoder



Naive Inception module

Case Study: GoogLeNet


Example: Q: What are the output size of filter concatenation?
(Assume that a proper size of zero-padding is used)

Q: What is the problem with this?
[Hint: Computational complexity]

Assignment Project Exam Help
 $28 \times 28 \times (128 + 192 + 96 + 256) = 28 \times 28 \times 672$

<https://powcoder.com>

Add WeChat powcoder



Naive Inception module

Conv Ops:

[1x1 conv, 128] $28 \times 28 \times 128 \times 1 \times 1 \times 256$

[3x3 conv, 192] $28 \times 28 \times 192 \times 3 \times 3 \times 256$

[5x5 conv, 96] $28 \times 28 \times 96 \times 5 \times 5 \times 256$

Total: 854M ops

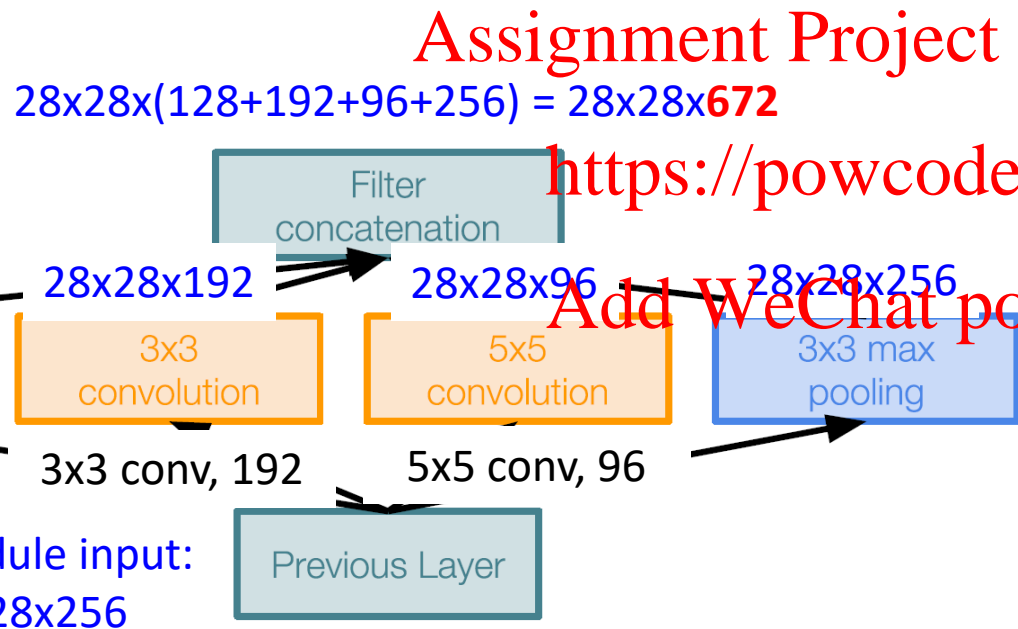
Very expensive compute

Pooling layer also preserves feature depth.
-> Total depth after concatenation can only grow at every layer!

Case Study: GoogLeNet

Example: Q: What are the output size of filter concatenation?
(Assume that a proper size of zero-padding is used)

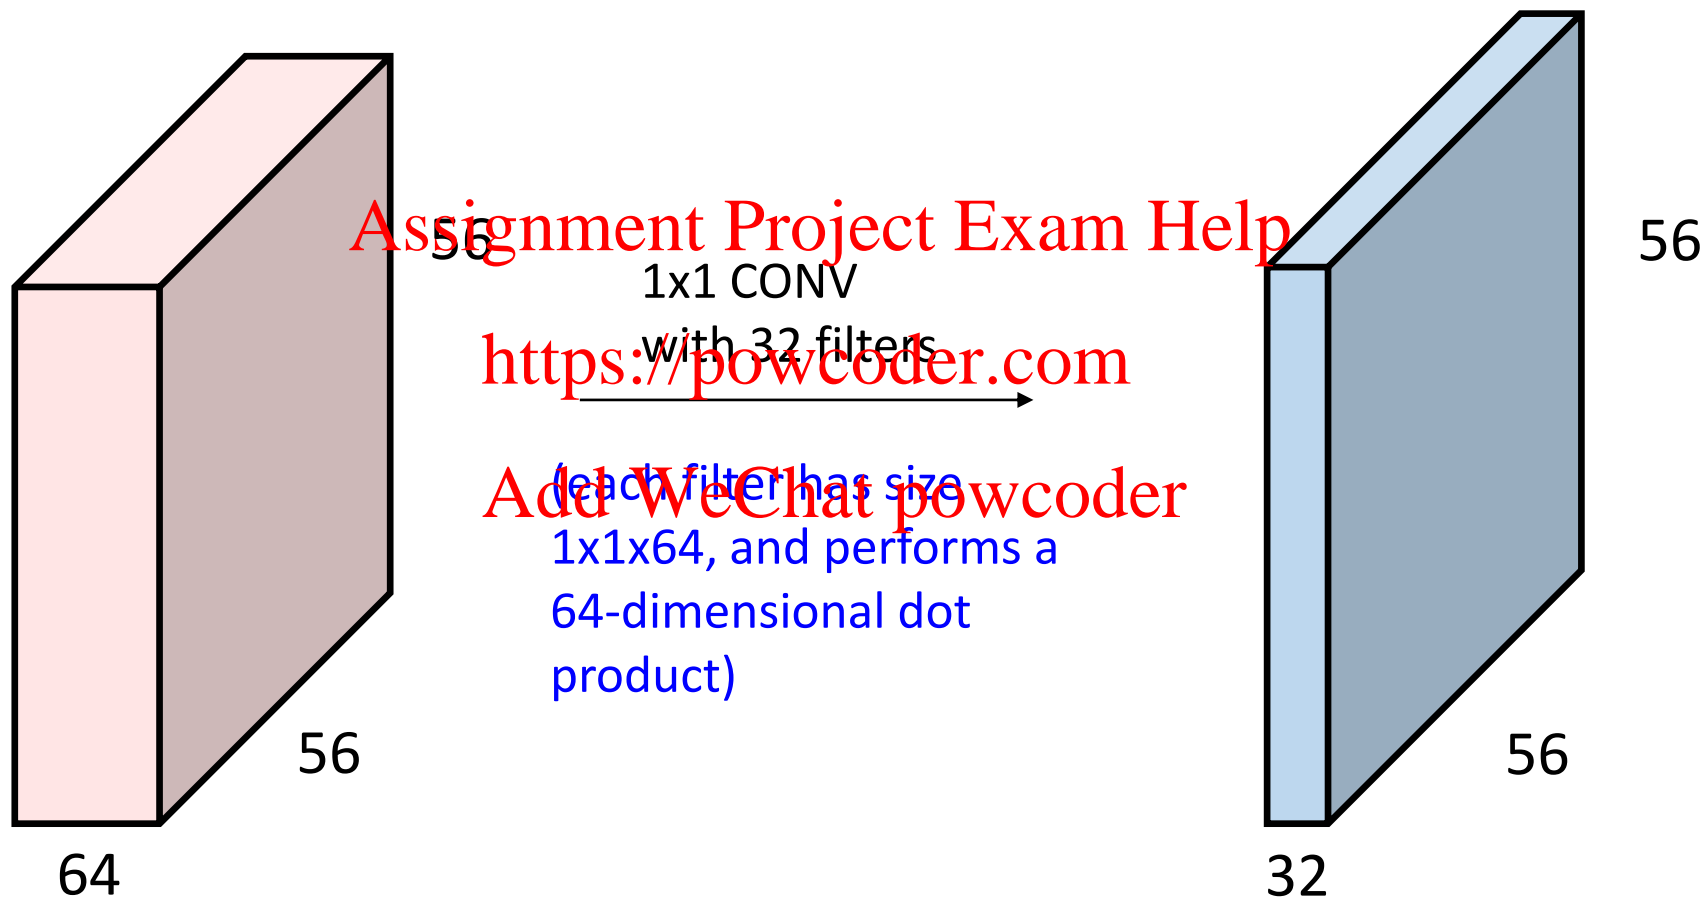
Q: What is the problem with this?
[Hint: Computational complexity]



Naive Inception module

Solution:
“bottleneck” layers that
use 1x1 convolutions to reduce
feature depth

Reminder: 1x1 Convolution

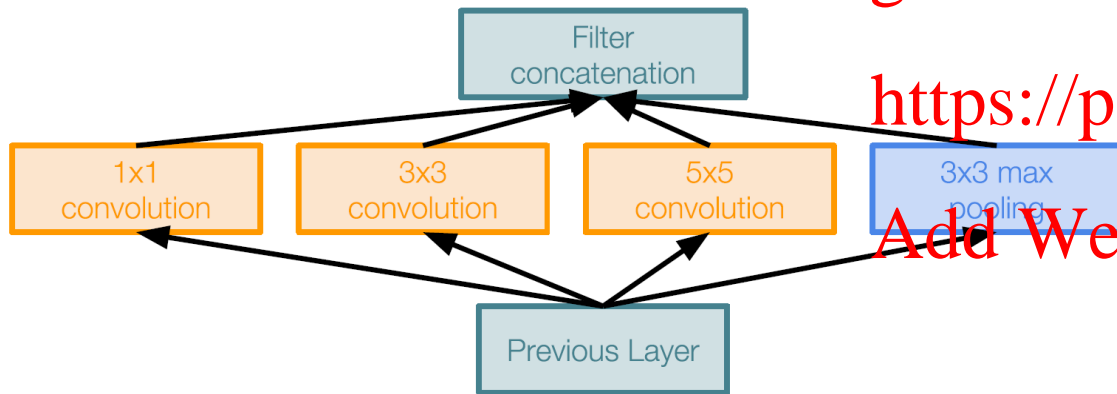


Case Study: GoogLeNet

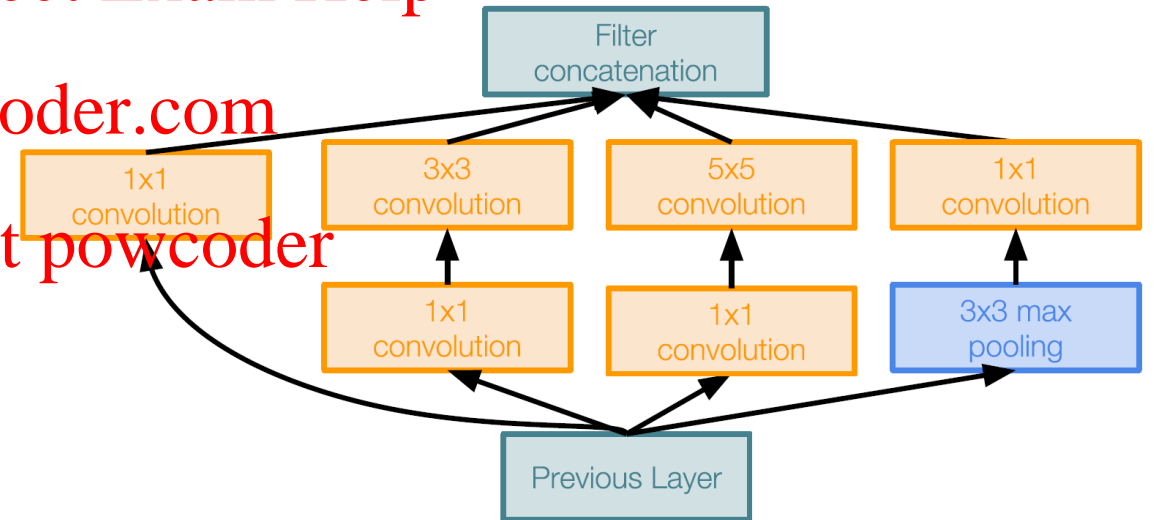
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

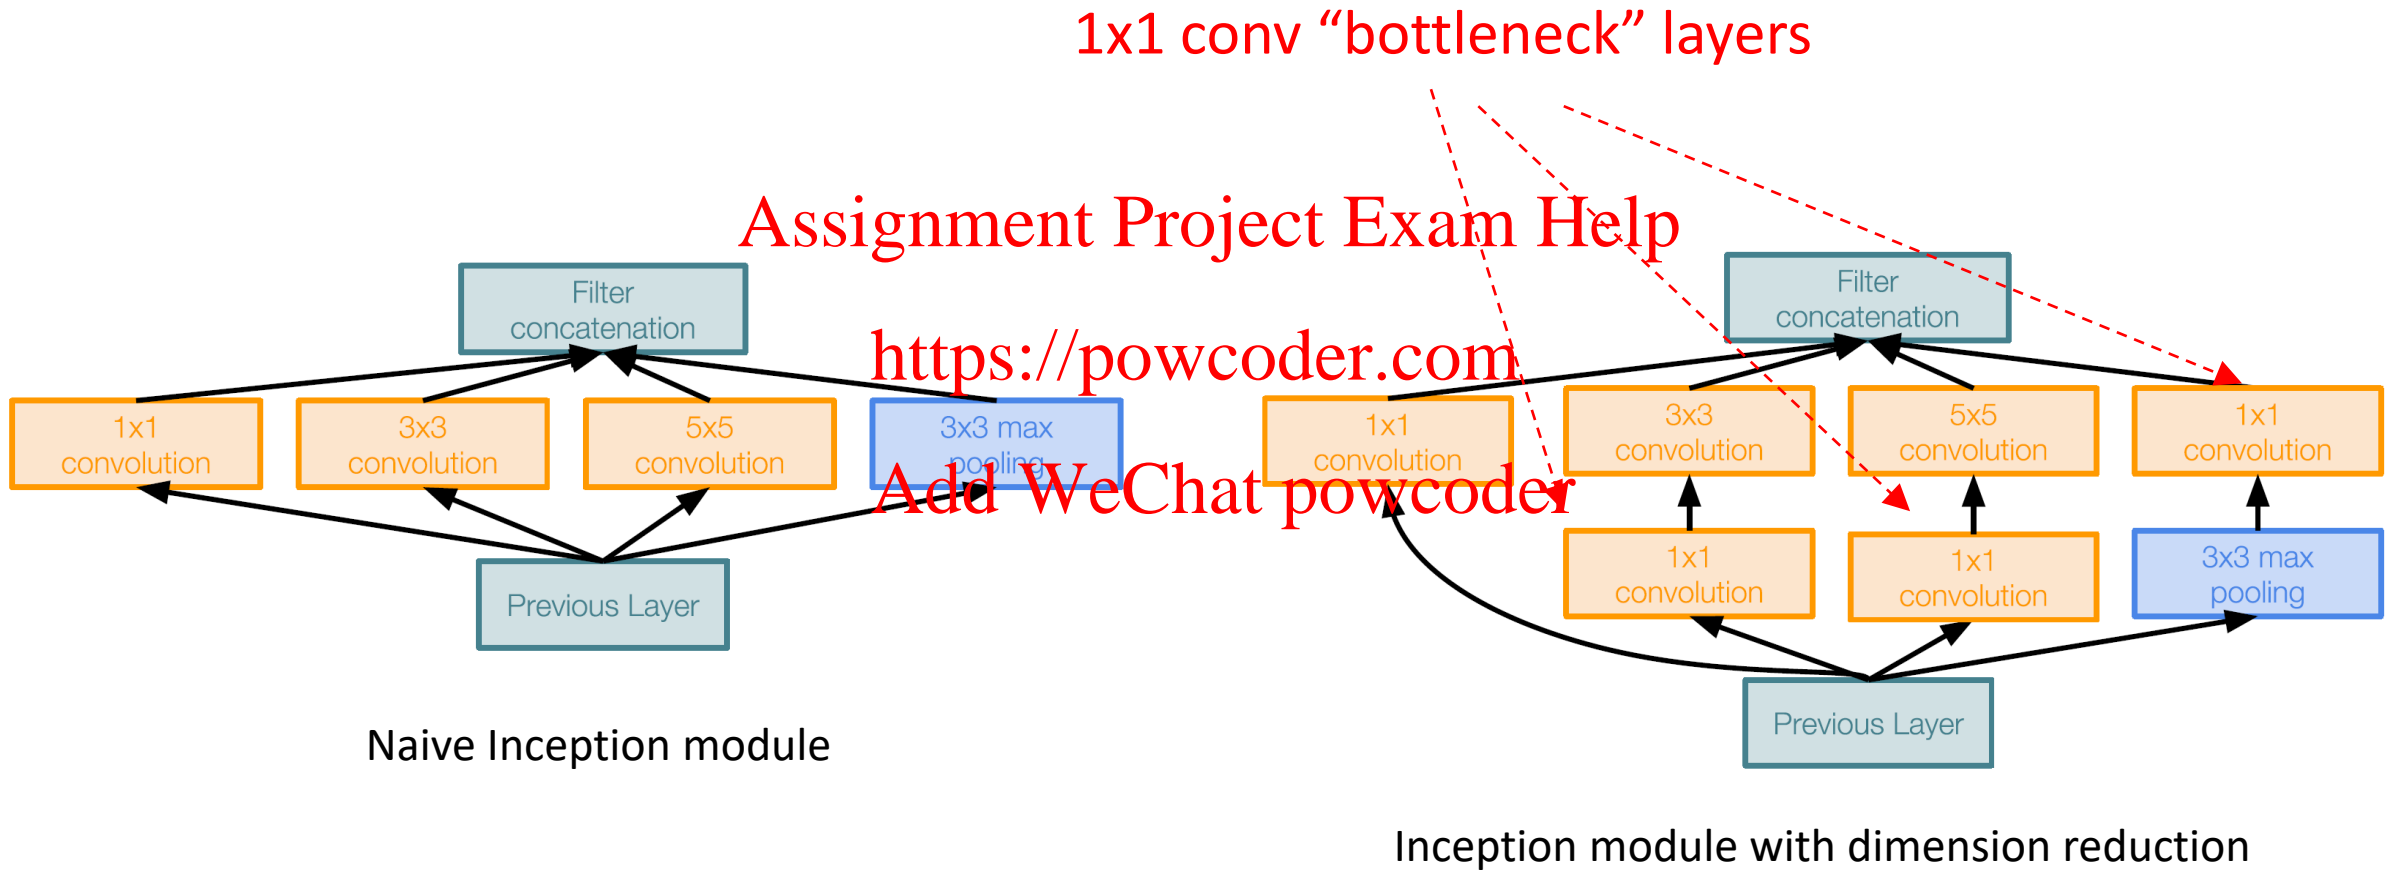


Naive Inception module

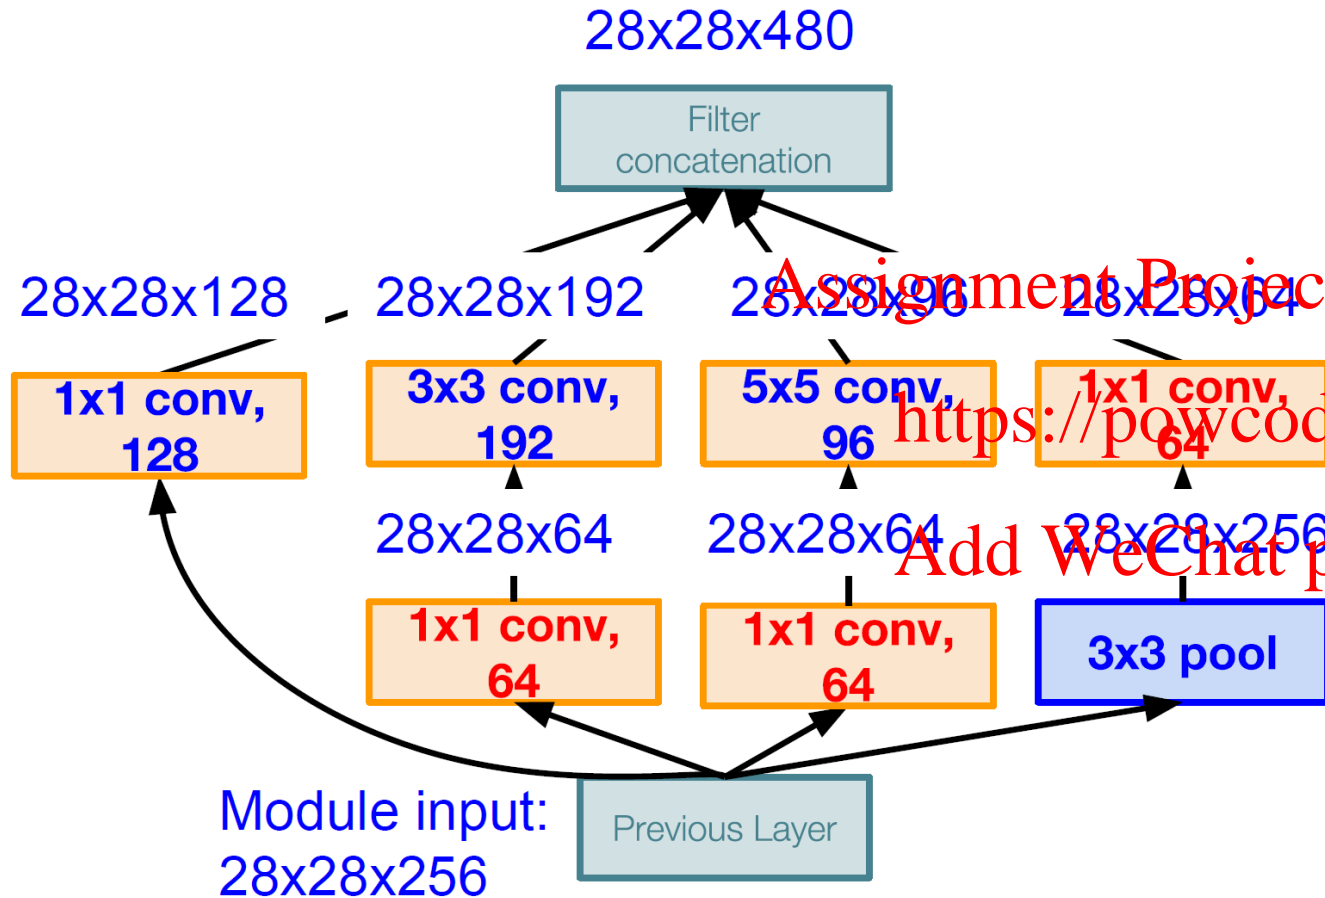


Inception module with dimension reduction

Case Study: GoogLeNet



Case Study: GoogLeNet



Inception module with dimension reduction

Using same parallel layers as naive example, and adding “1x1 conv, 64 filter” bottlenecks:

Conv Ops:

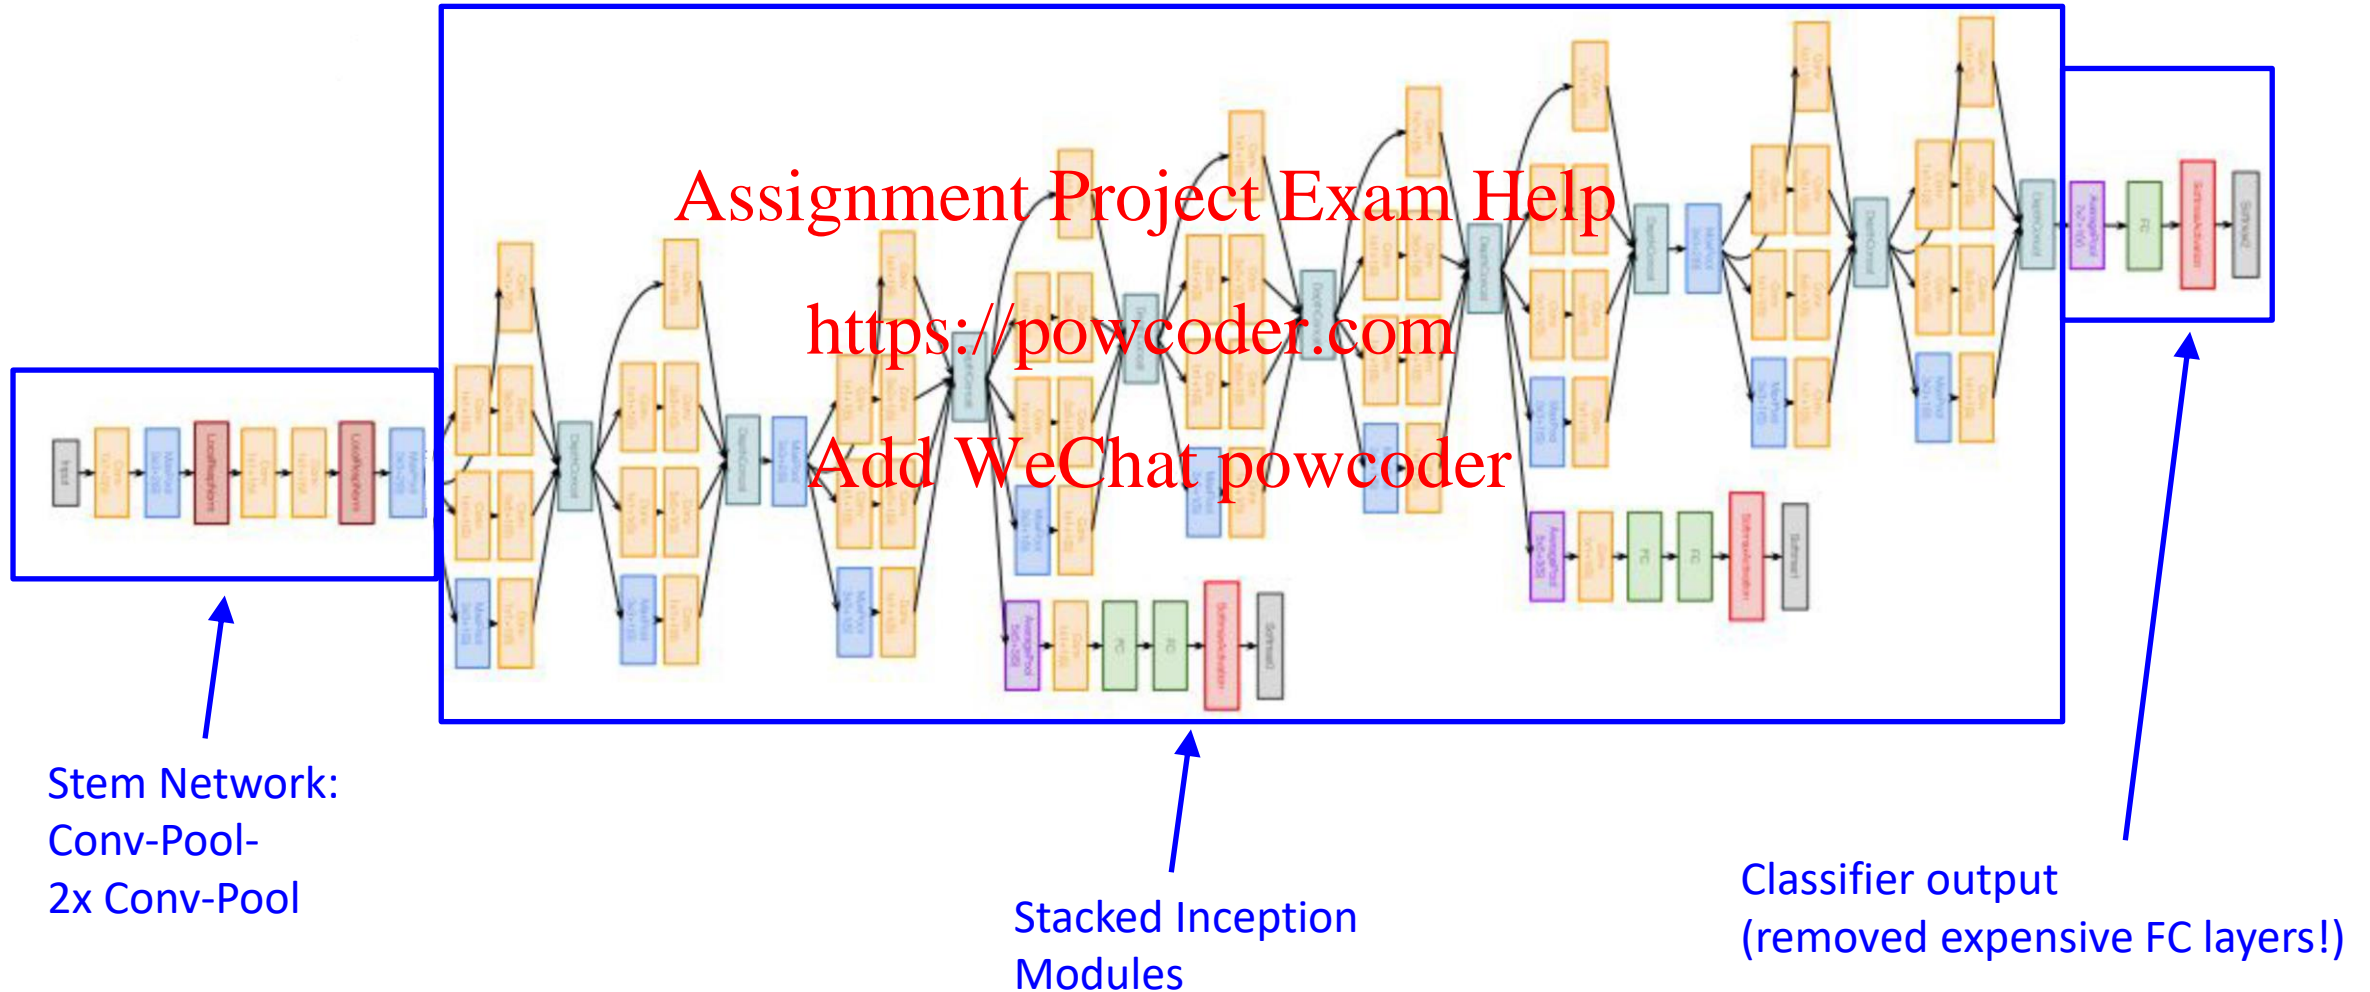
[1x1 conv, 64] 28x28x64x1x1x256
[1x1 conv, 64] 28x28x64x1x1x256
[1x1 conv, 128] 28x28x128x1x1x256
[3x3 conv, 192] 28x28x192x3x3x64
[5x5 conv, 96] 28x28x96x5x5x64
[1x1 conv, 64] 28x28x64x1x1x256

Total: 358M ops

Compared to 854M ops for naive version
Bottleneck can also reduce depth after pooling layer

Case Study: GoogLeNet

Full GoogLeNet architecture



Case Study: GoogLeNet

[Szegedy et al., 2014]


Deeper networks, with computational efficiency

- 22 layers
- Efficient “Inception” module
- No FC layers
- Only 5 million parameters!
12x less than AlexNet
- ILSVRC’14 classification winner
(6.7% top 5 error)


Assignment Project Exam Help

<https://powecoder.com>


Add WeChat powecoder



Inception module



ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



Case Study: ResNet

[He et al., 2015]


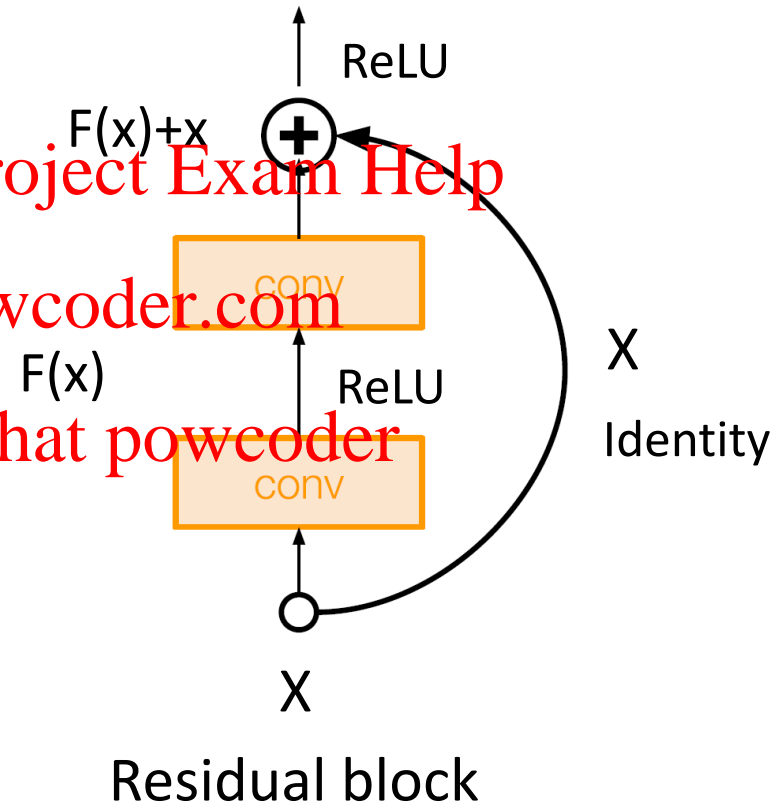
Very deep networks using residual connections

- 152-layer model for ImageNet
- ILSVRC'15 classification winner (3.57% top 5 error)
- Swept all classification and detection competitions in ILSVRC'15 and COCO'15!

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Case Study: ResNet



What happens when we continue stacking deeper layers on a “plain” convolutional neural network?

56-layer model performs worse on both training and test error

-> The deeper model performs worse, but it's not caused by overfitting!

<https://powcoder.com>

Add WeChat powcoder



Case Study: ResNet

Hypothesis: the problem is an *optimization*, not *the model itself*,
-> deeper models are harder to optimize

Assignment Project Exam Help

The deeper model should be able to perform at least as well as the shallower model.

<https://powcoder.com>

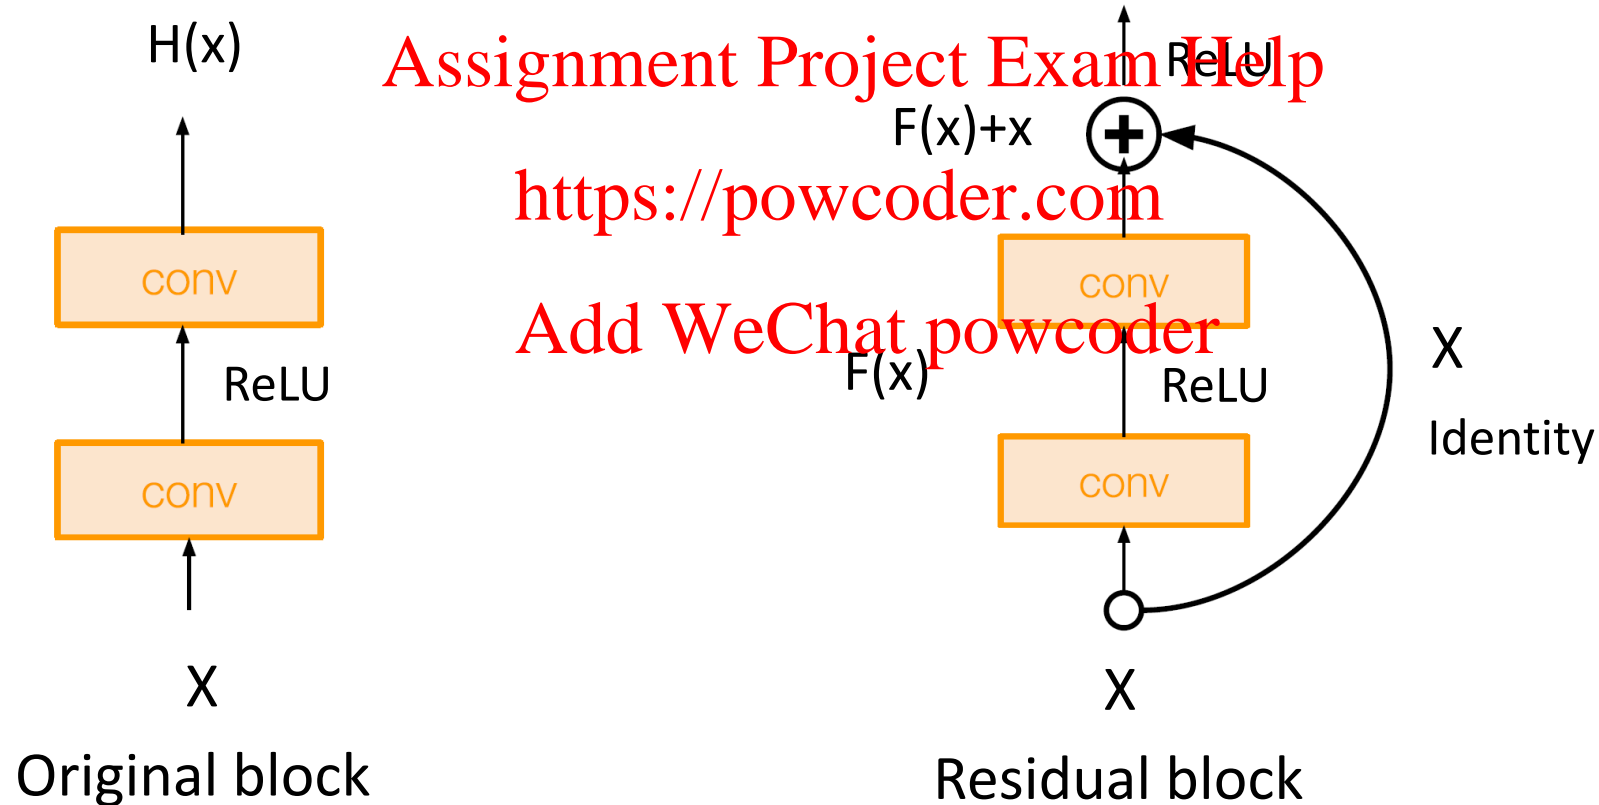
-> Solution by construction is copying the learned layers from the shallower model and setting additional layers to identity mapping.

Add WeChat powcoder

Case Study: ResNet

Solution

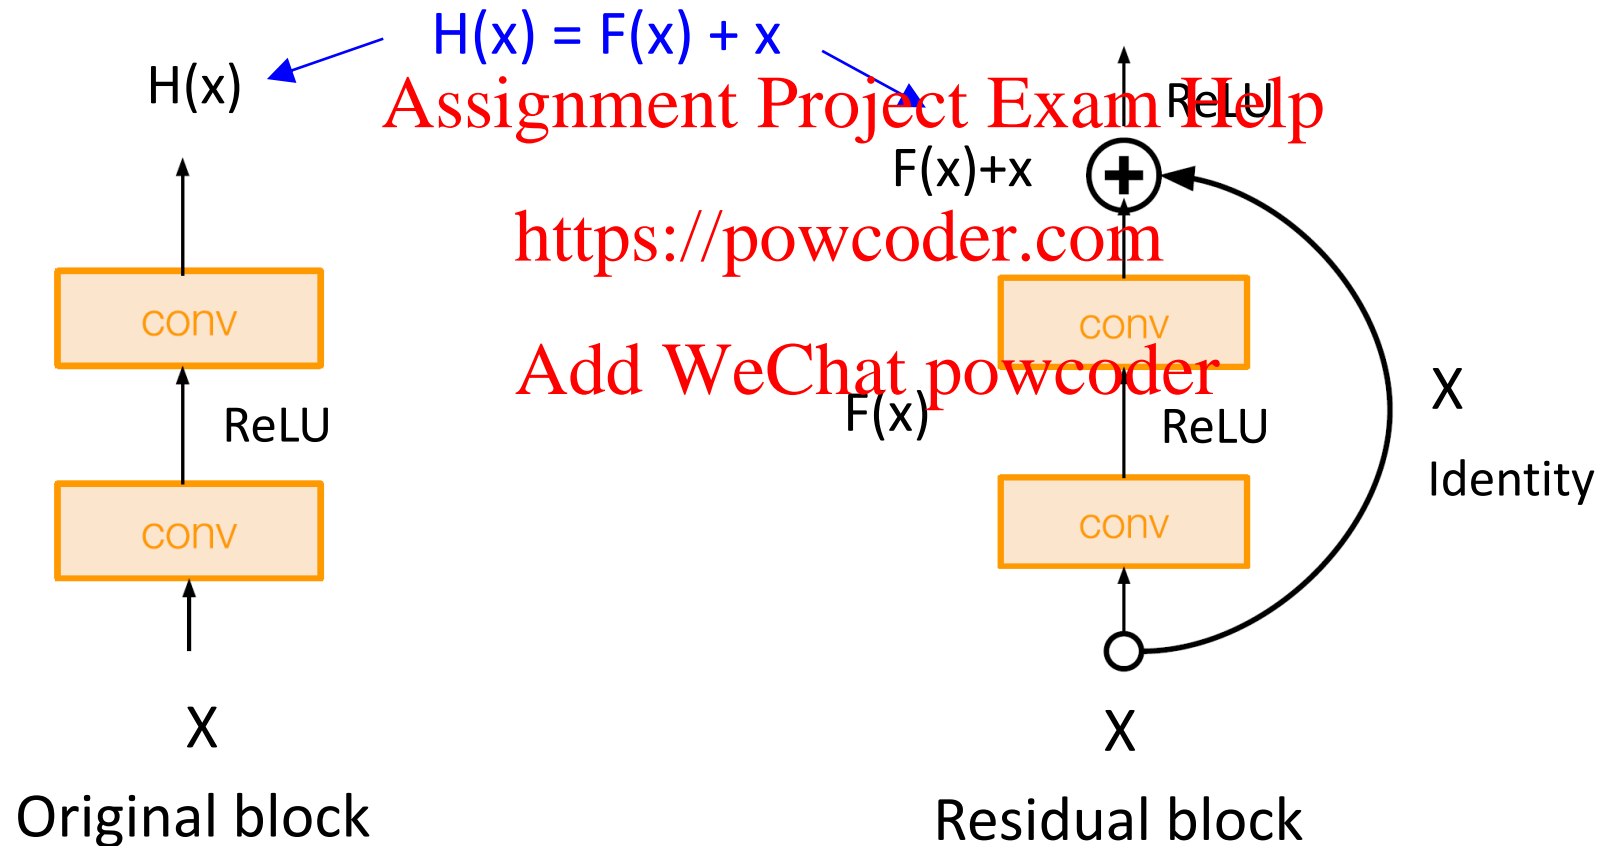
Use network layers to fit a residual mapping instead of directly trying to fit a desired underlying mapping



Case Study: ResNet

Solution

Use network layers to fit a residual mapping instead of directly trying to fit a desired underlying mapping



Use layers to
fit residual
 $F(x) = H(x) - x$
instead of
 $H(x)$ directly

Case Study: ResNet

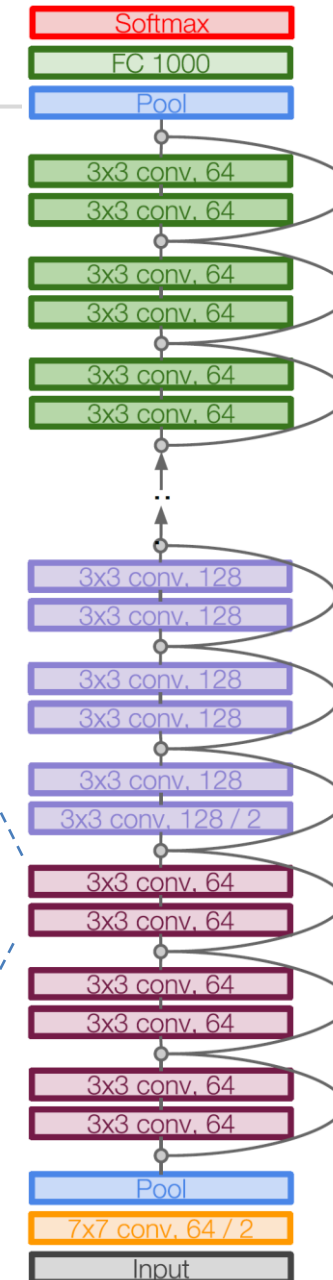
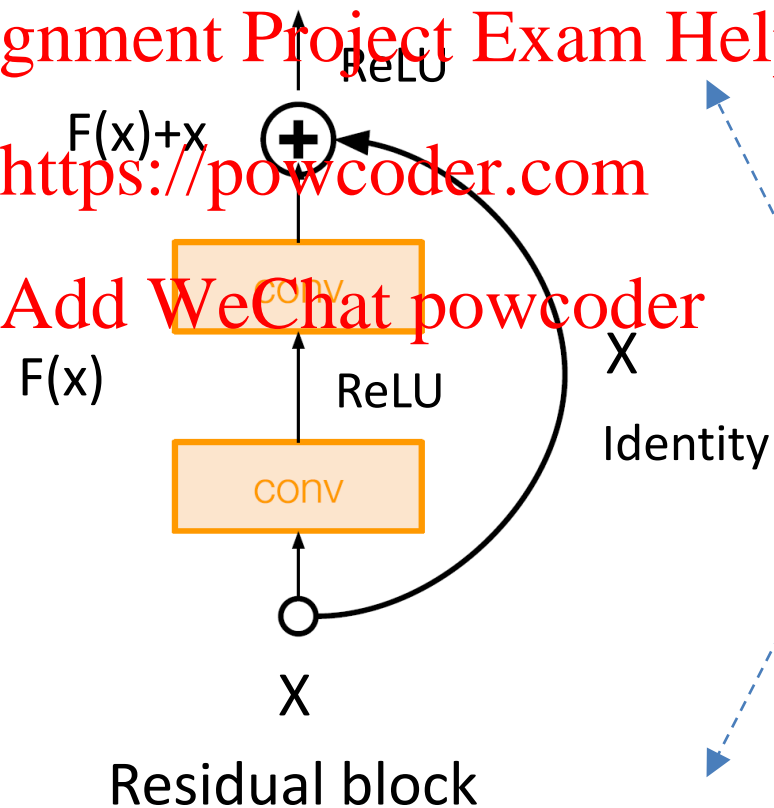
Full ResNet architecture:

- Stack residual blocks
- Every residual block has two 3x3 conv layers

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Case Study: ResNet

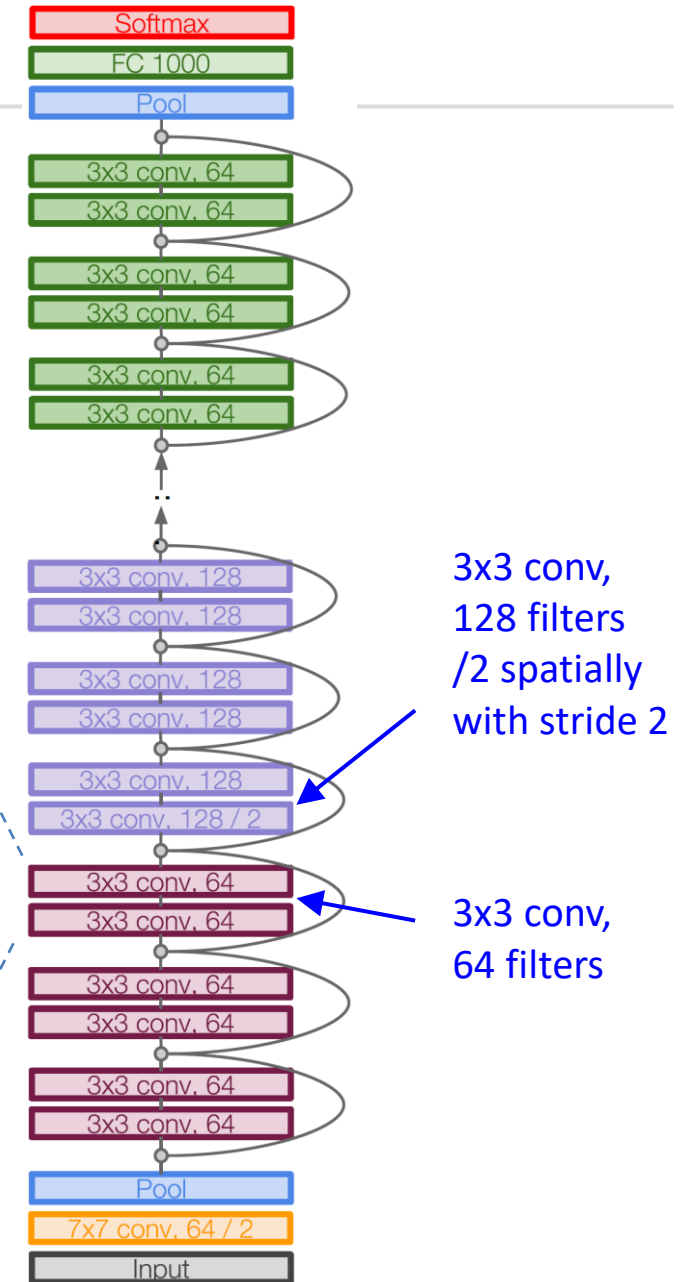
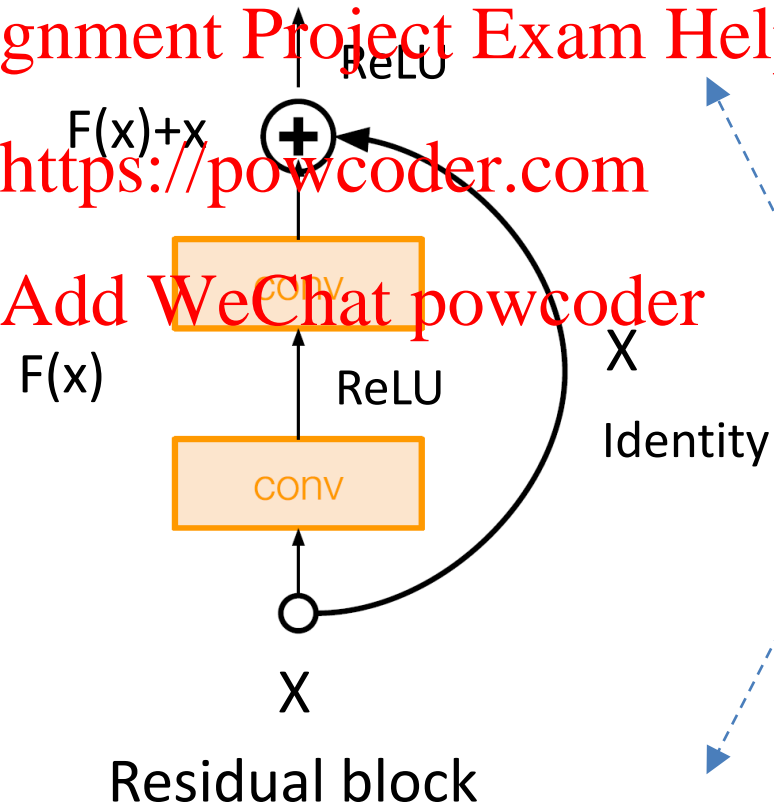
Full ResNet architecture:

- Stack residual blocks
- Every residual block has two 3x3 conv layers
- Periodically, double # of filters and downsample spatially using stride 2 (/2 in each dimension)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Case Study: ResNet

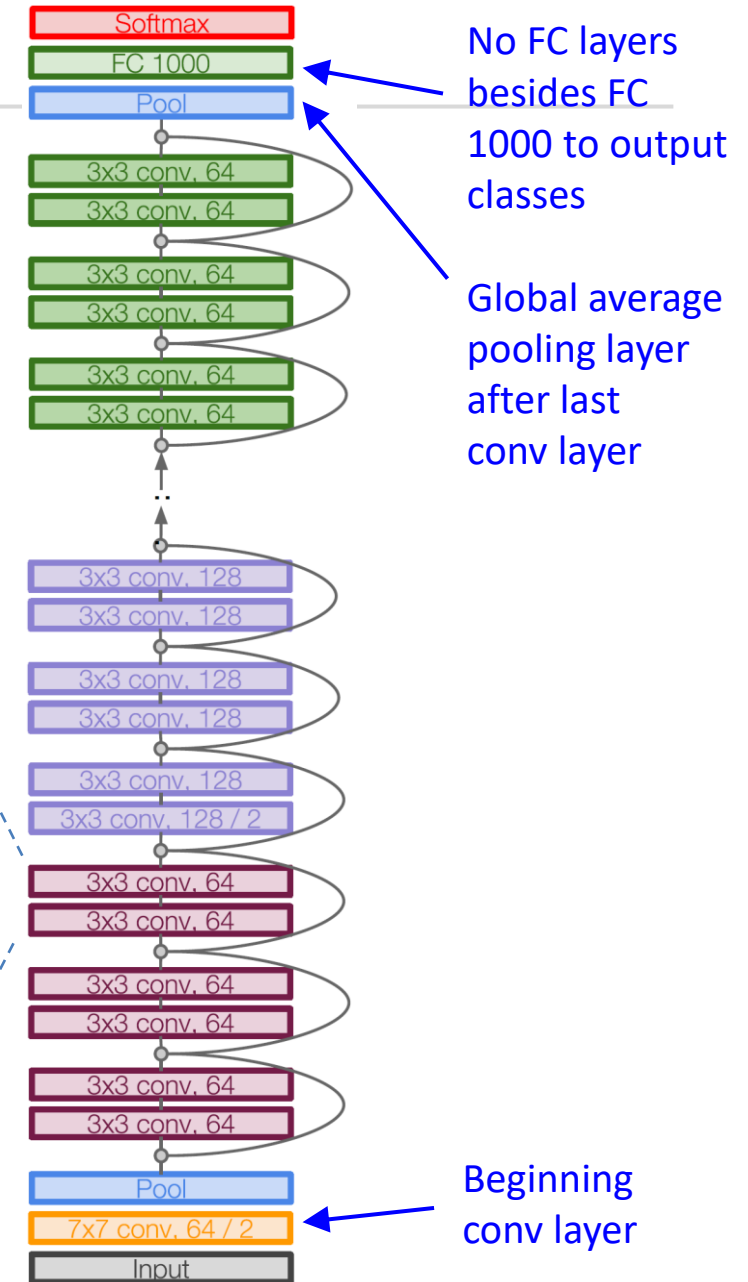
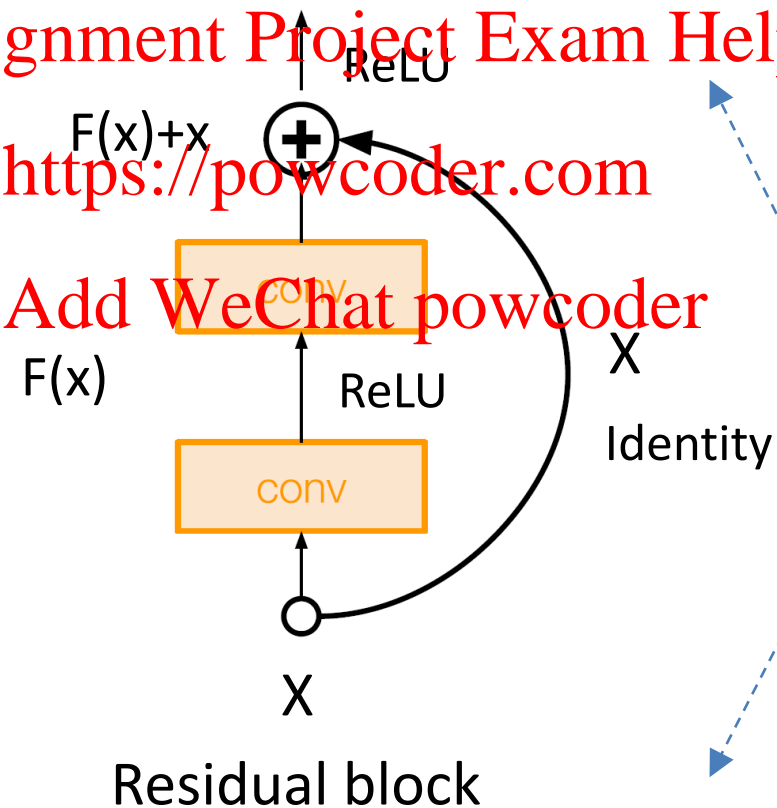
Full ResNet architecture:

- Stack residual blocks
- Every residual block has two 3x3 conv layers
- Periodically, double # of filters and downsample spatially using stride 2 (/2 in each dimension)
- Additional conv layer at the beginning
- No FC layers at the end (only FC 1000 to output classes)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Case Study: ResNet


Total depths of 18, 34, 50, 101, or 152 layers
(ResNet-18, ResNet-34, ResNet-50, ResNet-101, ResNet-152)

Assignment Project Exam Help

For deeper networks
(ResNet-50+), use “bottleneck”
layer to improve efficiency
(similar to GoogLeNet)

<https://powcoder.com>

Add WeChat powcoder



Case Study: ResNet

For deeper networks
(ResNet-50+), use “bottleneck”
layer to improve efficiency
(similar to GoogLeNet)


Assignment Project Exam Help

<https://powcoder.com>
3x3 conv operates over
only 64 feature maps

Add WeChat powcoder

1x1 conv, 64 filters
to project to 28x28x64

1x1 conv, 256 filters projects
back to 256 feature maps
(28x28x256)



Case Study: ResNet

- Experimental Results

- Able to train very deep networks without degrading (152 layers on ImageNet, 1202 on Cifar)
- Deeper networks now achieve lowering training error as expected
- Swept 1st place in all ILSVRC and COCO 2015 competitions

MSRA @ ILSVRC & COCO 2015 Competitions

- **1st places in all five main tracks**

- ImageNet Classification: “Ultra-deep” (quote Yann) **152-layer** nets
- ImageNet Detection: **16%** better than 2nd
- ImageNet Localization: **27%** better than 2nd
- COCO Detection: **11%** better than 2nd
- COCO Segmentation: **12%** better than 2nd

ILSVRC 2015 classification winner
(3.6% top 5 error)
better than “human performance”!
(Russakovsky 2014)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder